


# INF113: Linker and Syscalls

Kirill Simonov  
29.08.2025




# Feel the assembly

- Learn how old computers used to work
  - Example: PDP-11 from 1970s, watch **PDP11.mp4** in **Modules**
  - No OS, no processes
  - Setup the program directly in memory, and hit “run”
- Videogames\*




Human Resource Machine

\* not a paid advertisement




Turing Complete



TIS-100

# From assembly to binary



# Binaries

- Assembler converts mnemonic commands into machine codes
- Sets the file up in the ELF format, adding certain headers/metadata
  - Architecture/sizing information
  - Memory information
  - Sections and their offsets

but there's more...

## Executable and Linkable Format (**ELF**)

a.k.a. the standard binary file format on Unix-like systems

[https://en.wikipedia.org/wiki/Executable\\_and\\_Linkable\\_Format](https://en.wikipedia.org/wiki/Executable_and_Linkable_Format)

- Linker then adjusts metadata to make a standalone executable
  - Also “connects” with other binaries—will see later today!

# How to read ELFish

- hexdump outputs a binary file byte-by-byte in hex

executables are  $\geq 4\text{ kB}$ ...

```
$ cat hello.txt
Hello, world!
$ hexdump -C hello.txt
00000000  48 65 6c 6c 6f 2c 20 77  6f 72 6c 64 21 0a
                  |Hello, world!.|
```

0000000e

- objdump parses the ELF structure

```
$ objdump -d program
b:      file format elf64-x86-64
Disassembly of section .text:
0000000000401000 <_start>:
 401000: 48 c7 c7 23 00 00 00    mov     $0x23,%rdi
 401007: 48 c7 c0 3c 00 00 00    mov     $0x3c,%rax
 40100e: 0f 05                 syscall
```

actual code only!

# Linker

- The object file has “unset” addresses

```
$ nm program.o  
0000000000000000 T _start  
$ objdump -s program.o  
b.o:      file format elf64-x86-64  
Contents of section .text:  
0000 48c7c723 ...
```

- The linker adjusts the address where instructions start

```
$ nm program  
...  
0000000000401000 T _start  
$ objdump -s program  
b:      file format elf64-x86-64  
Contents of section .text:  
401000 48c7c724 ...
```

0x401000 is the default address  
for Instruction Counter to start


# Linking multiple files

```
.code64
.section .text
.global _start

_start:
    mov $19, %rax
    mov $16, %rbx

    call _add

    mov %rbx, %rdi
    mov $60, %rax
    syscall
```



```
.code64
.section .text
.global _add

_add:
    add %rax, %rbx
    ret
```

# Linking addresses



```
$ nm program.o
              U _add
0000000000000000 T _start
```

```
$ nm add.o
0000000000000000 T _add
```

each object file has an  
entry point at zero

```
$ nm program
000000000040101f T _add
0000000000402000 T __bss_start
0000000000402000 T _edata
0000000000402000 T _end
0000000000401000 T _start
```

linker arranges all records

# Linking calls



```
$ objdump -d program.o
```

```
0000000000000000 <_start>:  
 0: 48 c7 c0 13 00 00 00  
 7: 48 c7 c3 10 00 00 00  
 e: e8 00 00 00 00  
13: 48 89 df  
16: 48 c7 c0 3c 00 00 00  
1d: 0f 05
```

$$0x13 + 0x0c = 0x1f$$

old IC                    new IC

```
$ objdump -d program
```

```
0000000000401000 <_start>:  
401000: 48 c7 c0 13 00 00 00 mov $0x13,%rax  
401007: 48 c7 c3 10 00 00 00 mov $0x10,%rbx  
40100e: e8 0c 00 00 00 00 call 40101f <_add>  
401013: 48 89 df mov %rbx,%rdi  
401016: 48 c7 c0 3c 00 00 00 mov $0x3c,%rax  
40101d: 0f 05 syscall  
000000000040101f <_add>:  
40101f: 48 01 c3 add %rax,%rbx  
401022: c3 ret
```

- Linker puts the correct address for each function call

# System calls

- The programs are run in a “sandbox” environment
  - No I/O
  - No external memory
  - No setting exit code
  - ...
- To perform these actions, the program makes a syscall:
  1. Write what exactly needs to be done
  - 2. Pass the control back to the OS
  3. OS checks that instructions left by the program are ok
  4. OS completes the action and passes control back

```
.code64  
.section .text  
.global _start  
  
_start:  
    mov $19, %rax  
    mov $16, %rbx  

```

Here: the value 60 in %rax is interpreted as the exit command  
OS expects to find the exit code in %rdi

# How to find syscalls

- Whenever syscall is invoked, the OS checks the operation code in %rax, e.g.:
  - 0: read
  - 1: write
  - 60: exit
- Look for codes
  - Locally in `/usr/include/x86_64-linux-gnu/asm/unistd_64.h`
  - Or in the source repo
    - [https://github.com/torvalds/linux/blob/master/arch/x86/entry/syscalls/syscall\\_64.tbl](https://github.com/torvalds/linux/blob/master/arch/x86/entry/syscalls/syscall_64.tbl)
  - Or in the guides, here also which argument goes to which register:
    - [https://blog.rchapman.org/posts/Linux\\_System\\_Call\\_Table\\_for\\_x86\\_64/](https://blog.rchapman.org/posts/Linux_System_Call_Table_for_x86_64/)
- OS will expect arguments in certain order:

register mapping for system call invocation using <code>syscall</code>								
system call number	1 <sup>st</sup> parameter	2 <sup>nd</sup> parameter	3 <sup>rd</sup> parameter	4 <sup>th</sup> parameter	5 <sup>th</sup> parameter	6 <sup>th</sup> parameter	result	
	<code>rax</code>	<code>rdi</code>	<code>rsi</code>	<code>rdx</code>	<code>r10</code>	<code>r8</code>	<code>r9</code>	<code>rax</code>

[https://en.wikibooks.org/wiki/X86\\_Assembly/Interfacing\\_with\\_Linux](https://en.wikibooks.org/wiki/X86_Assembly/Interfacing_with_Linux)

# Hello, world!

- write has code 1, and expects arguments:
  1. stream in %rdi,
  2. address of the string in %rsi,
  3. length of the string in %rdx
- Pre-defined values have a separate section, here in read-only (.rodata)
- We can use assembly expressions to get the length

```
.code64
.section .rodata
msg:
    .ascii "Hello, world!\n"
    .set msglen, (. - msg)
.section .text
.global _start
_start:
    mov $1, %rax
    mov $1, %rdi
    lea msg, %rsi
    mov $msglen, %rdx
    syscall

    mov $60, %rax
    xor %rdi, %rdi
    syscall
```

# System calls in C

- System calls have the same principle for C programs
- System calls themselves are implemented as C functions
  - Look up in the source repo  
[https://github.com/torvalds/linux/blob/master/arch/x86/entry/syscalls/syscall\\_64.tbl](https://github.com/torvalds/linux/blob/master/arch/x86/entry/syscalls/syscall_64.tbl)
  - Also local manuals:

```
$ man 2 write
SYNOPSIS
    #include <unistd.h>

    ssize_t write(int fd, const void *buf [.count], size_t count);
```

- Compiling with gcc both generates the object file, and calls the linker
  - The result is always linked with libc, hence calling library function works