

# INF113: Scheduling

Kirill Simonov  
10.09.2025

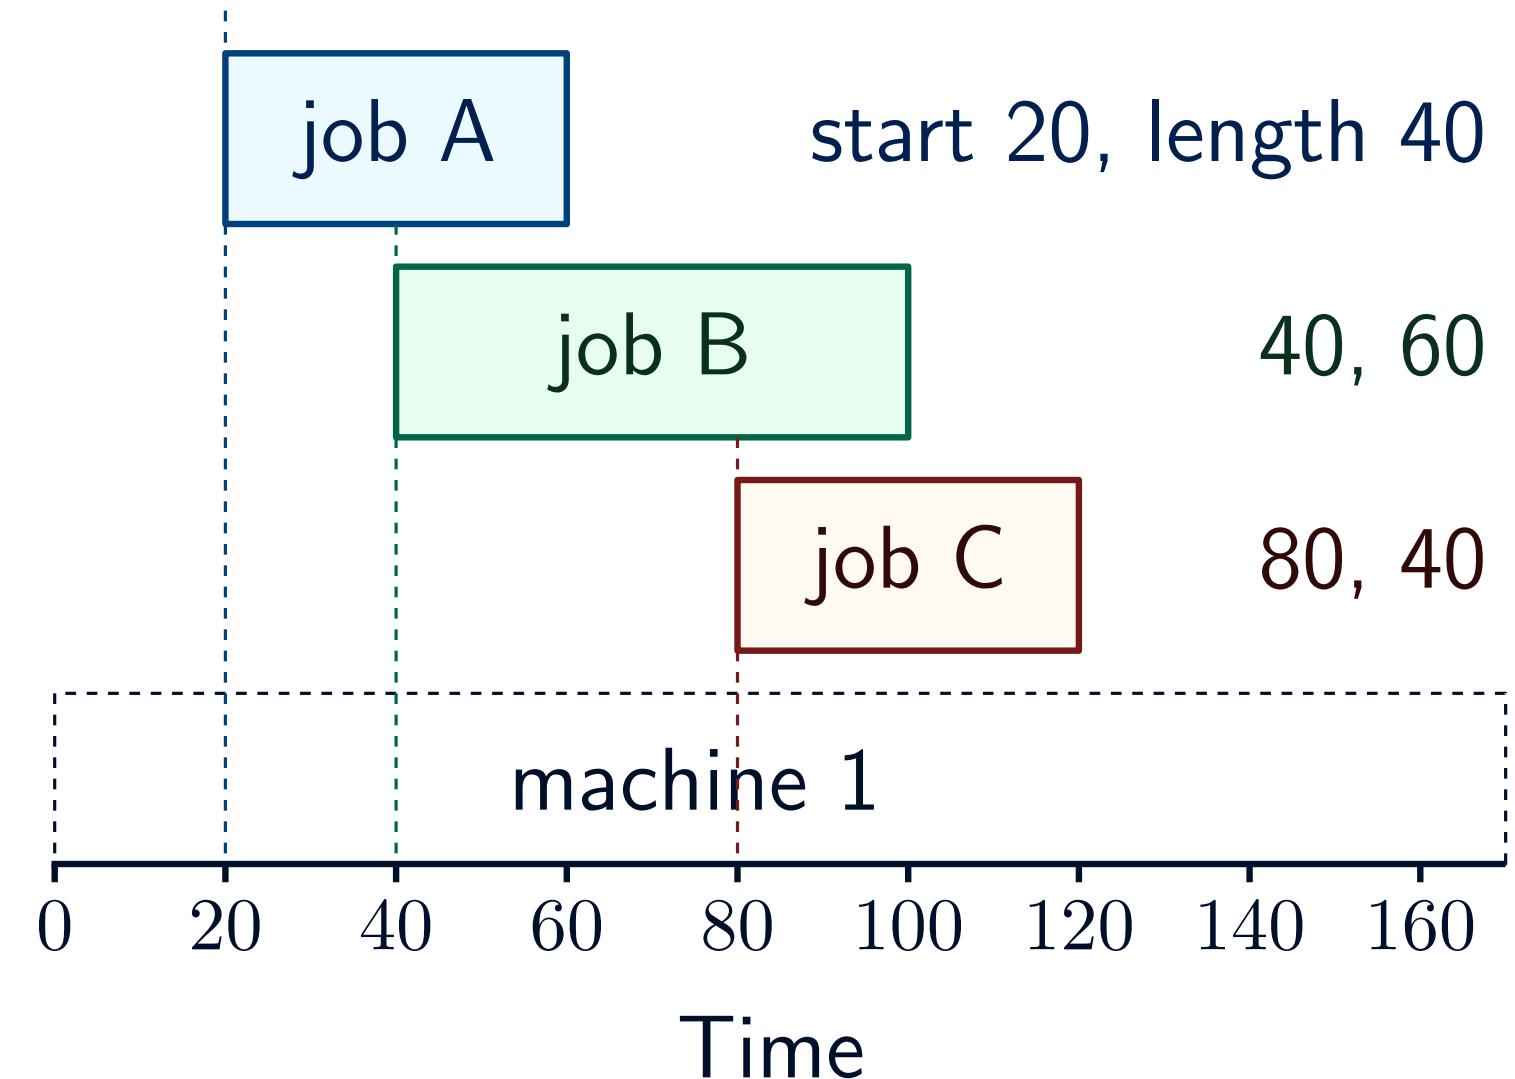


# Questions for today

- How to model scheduling?
- How to measure that a schedule is “good” ?
- How to schedule fast and well without prior knowledge?

# Scheduling in general

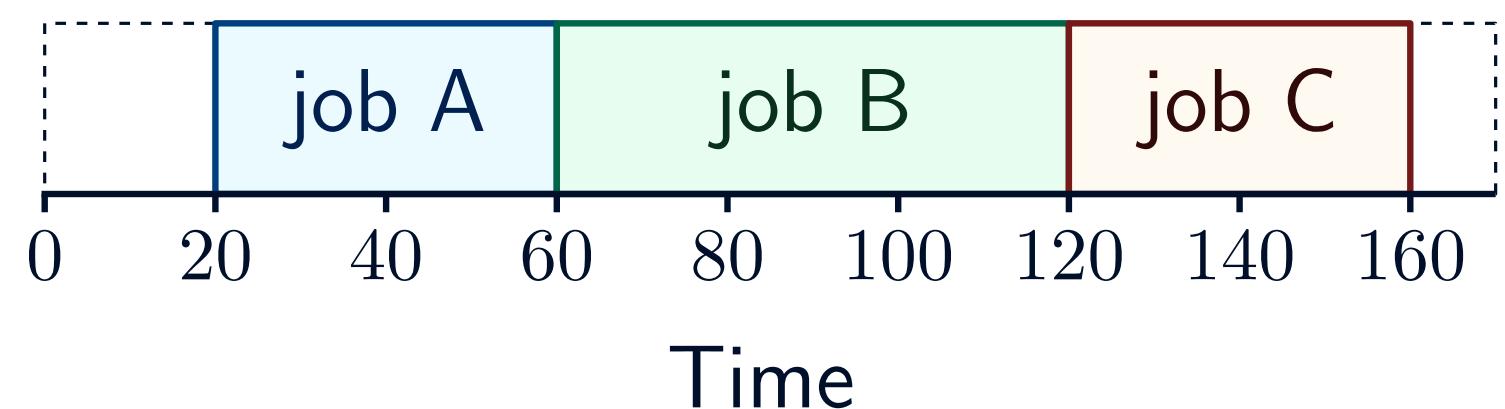
- A broad class of optimization problems can be viewed as scheduling
- **Input:** a collection of jobs, and a number of machines
- Each job is defined by
  - length
  - arrival time
- Each machine serves jobs, one at a time
- **Task:** Come up with an optimal schedule



# Scheduling in general

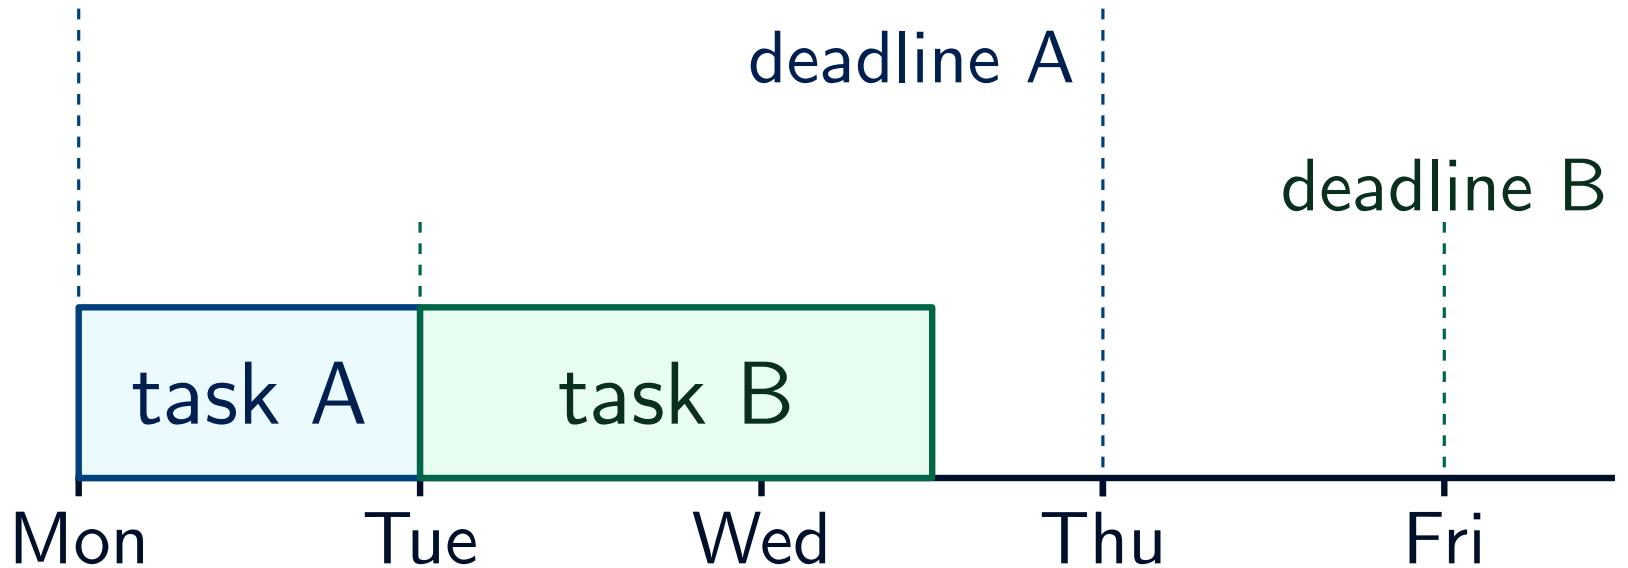
- A broad class of optimization problems can be viewed as scheduling
- **Input:** a collection of jobs, and a number of machines
- Each job is defined by
  - length
  - arrival time
- Each machine serves jobs, one at a time
- **Task:** Come up with an optimal schedule

what is optimal?..



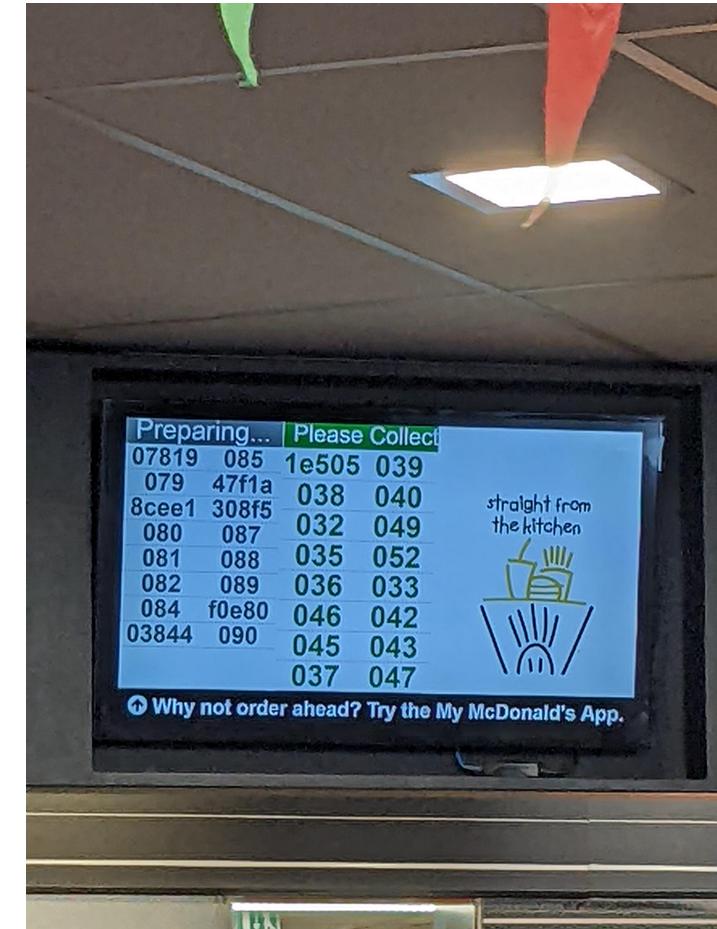
# Examples of scheduling

- Completing homework
  - Food/drink orders



- Arranging classes

	Monday 8/9	Tuesday 9/9	Wednesday 10/9	Thursday 11/9	Friday 12/9
8			E. Ullestad: null 08:00 - 10:00 A66 Aud A	INF170 : 08:15 - 10:00 A. Hemmati A66 Aud A	
9					
10	INF100 : 10:15 - 12:00 T.J.F. Strømme A66 Aud A	INF170 : 10:15 - 12:00 A. Hemmati A66 Aud A	INF113 : 10:15 - 12:00 K. Simonov A66 Aud A	STAT101 : 10:15 - 12:00 B. Støve A66 Aud A	PHYS101 : 09:15 - 11:00 M.M. Greve A66 Aud A
11					M.M. Greve: null
12	STAT101 : 12:15 - 14:00 A66 Aud A	S. Gisnås: null 12:00 - 15:30 A66 Aud A	EXPHIL-MNSEM : 12:15 - 14:00 T.T. Huvenes A66 Aud A		
13					
14	STAT110 : 14:15 - 16:00 H.J. Skaug A66 Aud A			M.M. Greve: null 14:15 - 16:00 A66 Aud A	
15					
16					

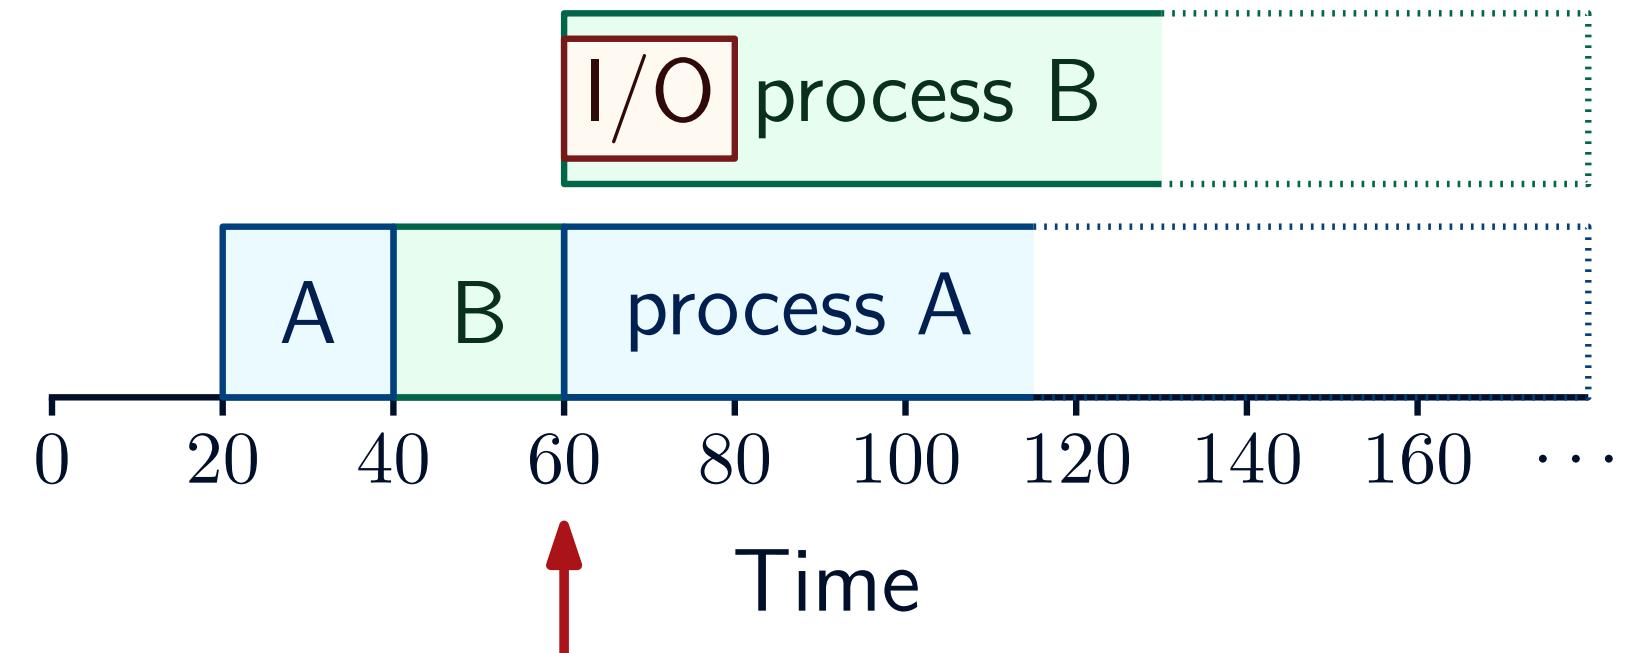


- Ambulance dispatch, air traffic, TikTok queue...

# CPU Scheduling

- One machine—the CPU
- Jobs—processes—are only revealed at arrival time
- The length of a job is unknown
- Jobs may be interrupted
- Jobs may need to wait on I/O

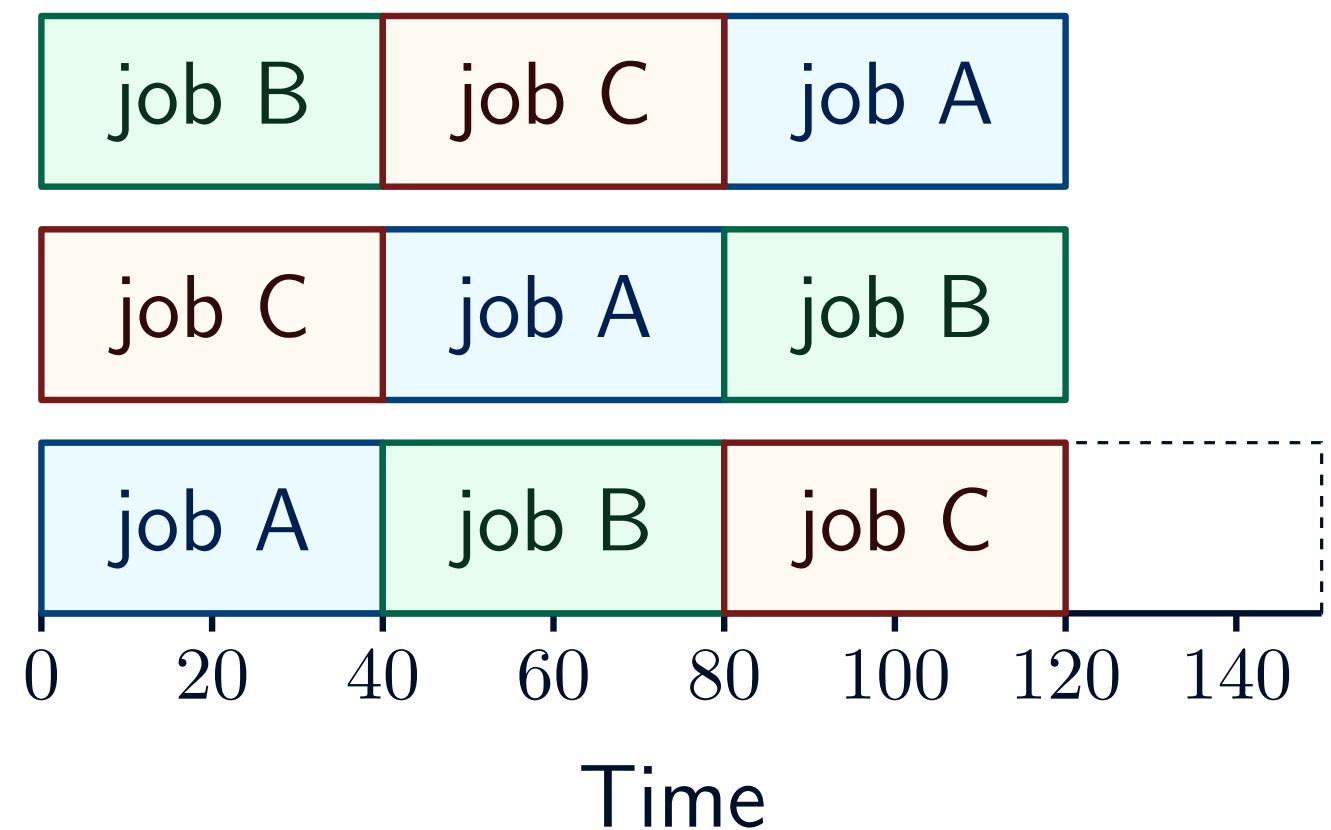
Scheduler should be  
**fast and robust**



# Starting simple

- For now, assume that
  1. All jobs arrive at the same time
  2. All jobs have the same length
  3. The jobs cannot be interrupted
  4. The jobs do not wait on I/O
  5. The length of a job is known in advance

At this point, maybe just run in any order?



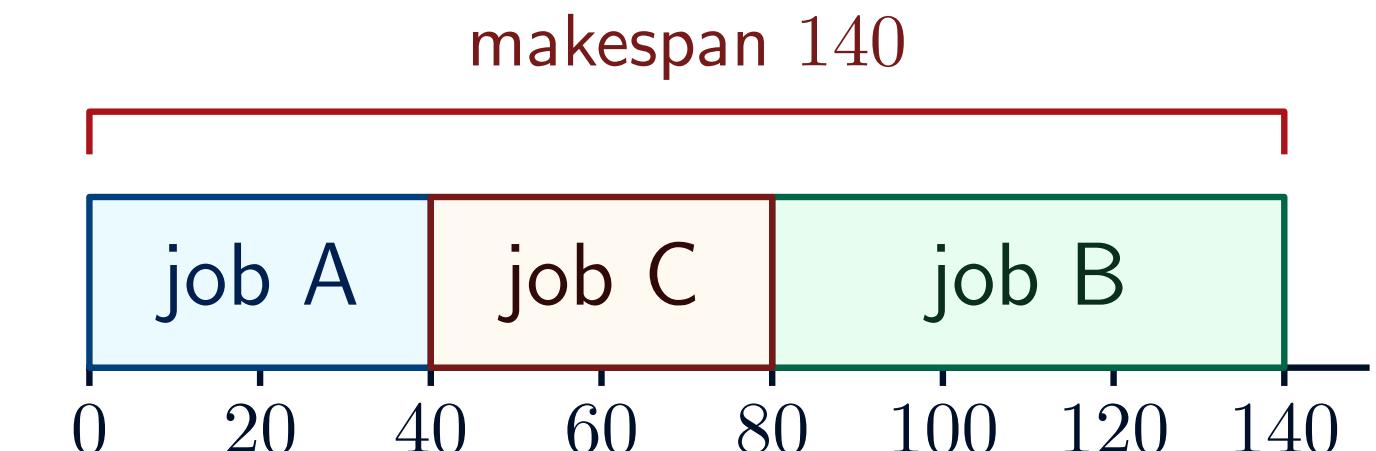
# Measuring performance

- Convenient to have a numerical value for the efficiency
- **Attempt:** Let **makespan** be the completion time of the last job
  - Effectively, aiming for less gaps
  - Almost useless for a single machine
- **Better:** Let **turnaround** be the completion time minus arrival time of the job

$$T_{\text{completion}} - T_{\text{arrival}}$$

- We will measure the **average turnaround**

$$\frac{\sum T_{\text{completion}} - T_{\text{arrival}}}{|\text{jobs}|}$$



turnaround 40    80    140

$$\text{average turnaround} = \frac{40+80+140}{3} \approx 86.67$$

# First In, First Out

- **Algorithm:** Run the job that arrived the earliest
- **Theorem:** FIFO is optimal for average turnaround when lengths are the same  
*no proof in this class*

FIFO: makespan 120  
av. turnaround  $\approx 46.67$

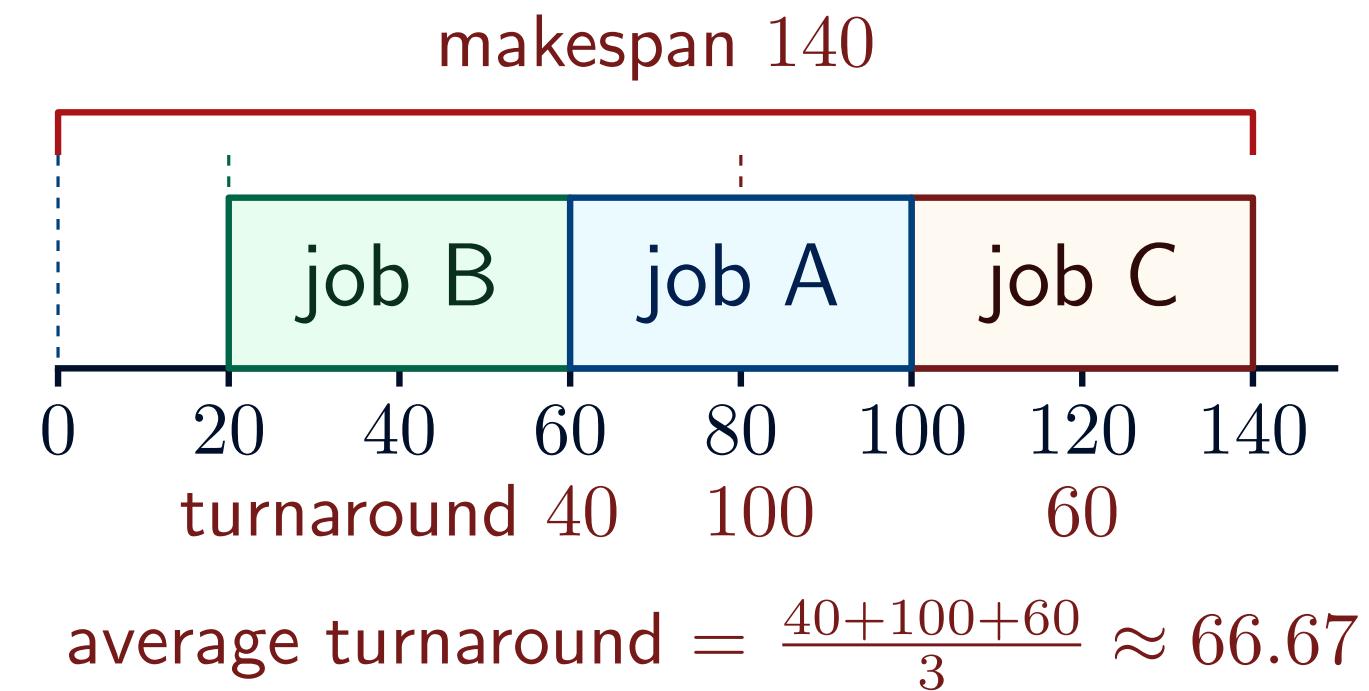


Dummy: makespan 140  
av. turnaround  $\approx 66.67$



Assume that

1. All jobs arrive at the same time
2. All jobs have the same length
3. The jobs cannot be interrupted
4. The jobs do not wait on I/O
5. The length of a job is known in advance



# Shortest Job First

- **Algorithm:** Run the shortest available job
- **Theorem:** SJF is optimal for average turnaround when starting time is the same  
*no proof in this class*

SJF: makespan 140  
av. turnaround  $\approx 86.67$



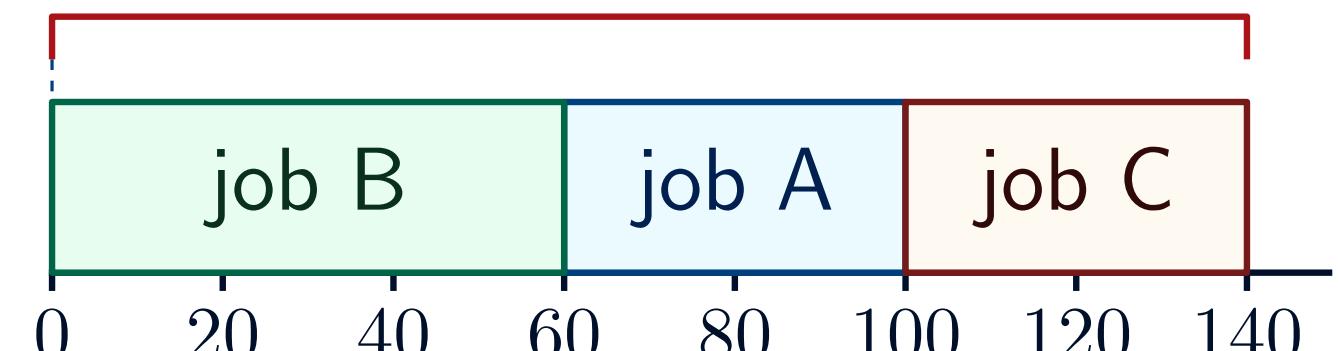
FIFO: makespan 140  
av. turnaround = 100



makespan 140

Assume that

1. All jobs arrive at the same time
2. ~~All jobs have the same length~~
3. The jobs cannot be interrupted
4. The jobs do not wait on I/O
5. The length of a job is known in advance



turnaround 60 100 140

$$\text{average turnaround} = \frac{60+100+140}{3} = 100$$

# SJF needs interrupts

- **Problem:** SJF does not use arrival times

SJF:

av. turnaround = 100



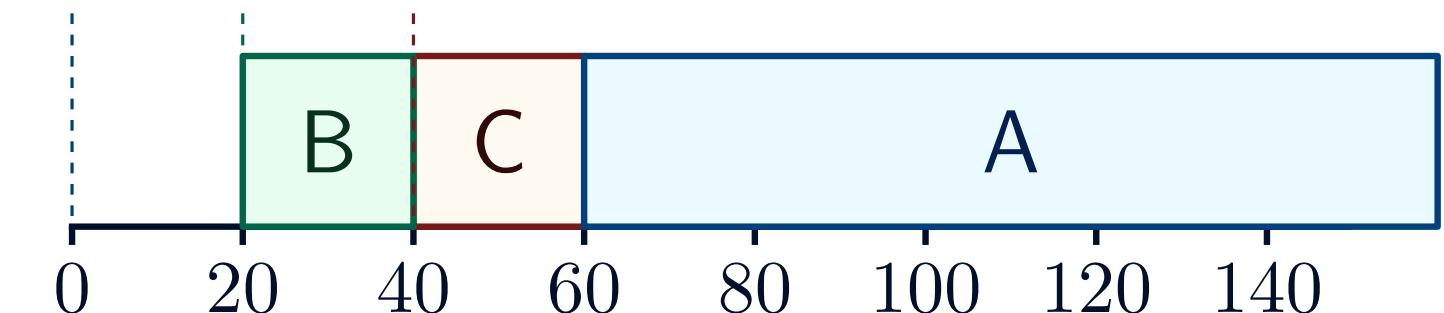
Optimal:

av. turnaround  $\approx 66.67$



Assume that

1. All jobs arrive at the same time
2. All jobs have the same length
3. The jobs cannot be interrupted
4. The jobs do not wait on I/O
5. The length of a job is known in advance



$$\text{average turnaround} = \frac{20+20+160}{3} \approx 66.67$$

# Shortest to Completion First

- **Algorithm:** Run the job with the least time **left**
- No need to wait for completion thanks to interrupts and context switches

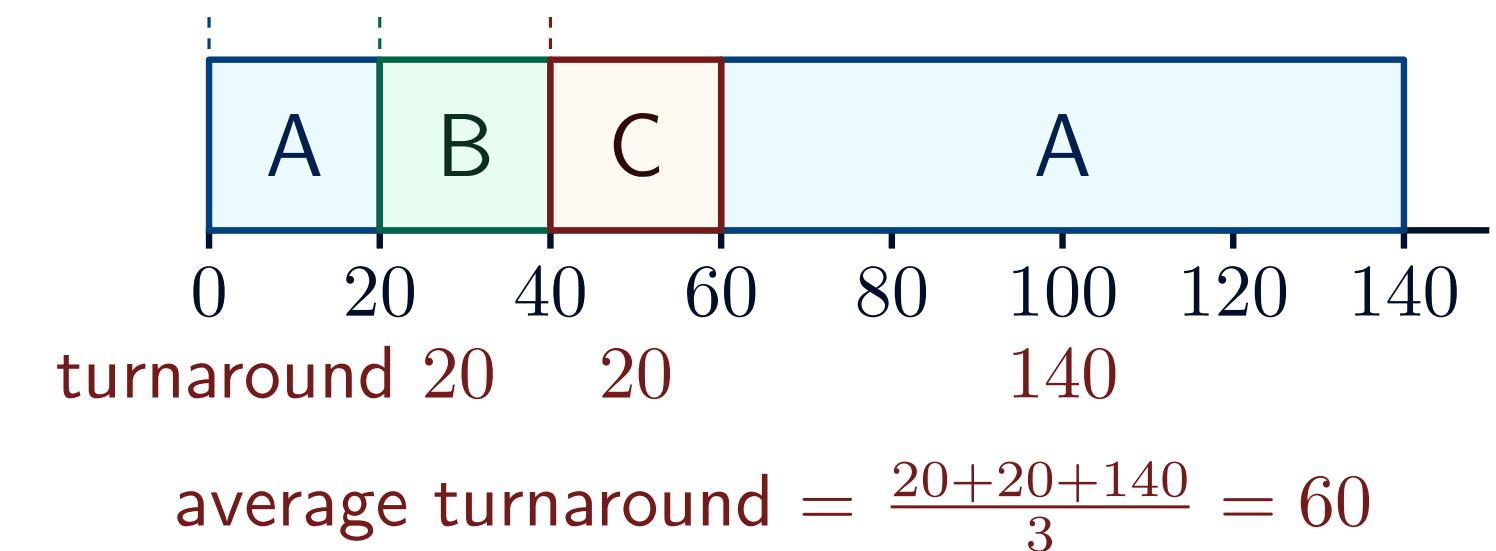
Optimal w/o interrupts:  
av. turnaround  $\approx 66.67$



STCF:  
av. turnaround = 60



1. All jobs arrive at the same time
2. All jobs have the same length
3. The jobs cannot be interrupted
4. The jobs do not wait on I/O
5. The length of a job is known in advance



# Response time

- STCF would be perfect for the early-day OS
- Nowadays, we do not only care about performance, but also interactivity
- **Response time** is the time the job has to wait:

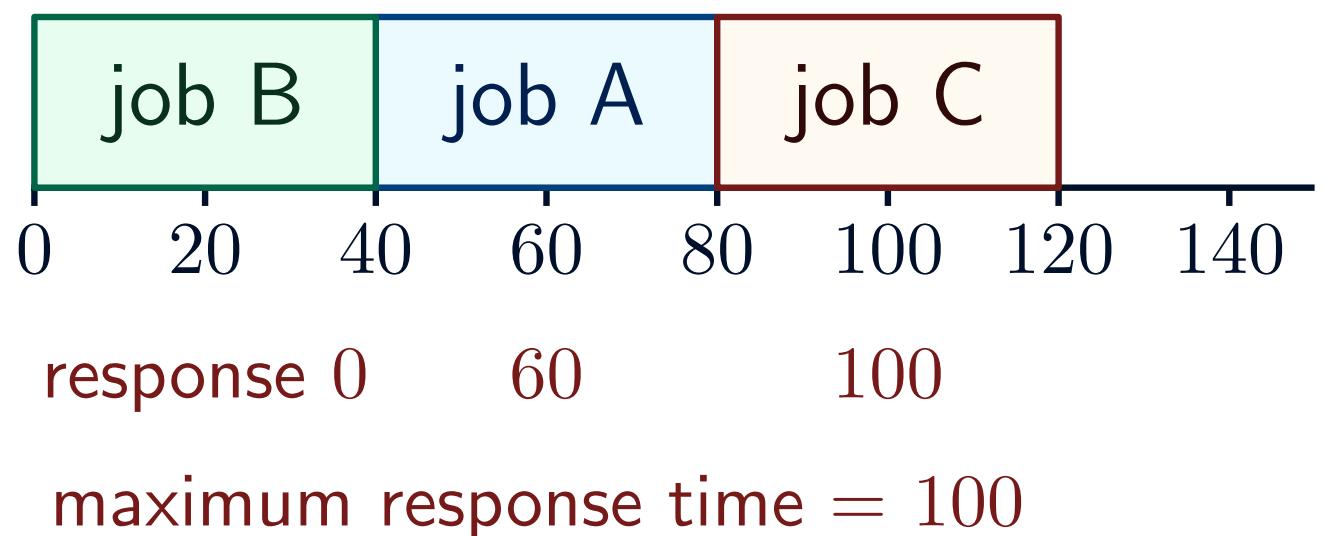
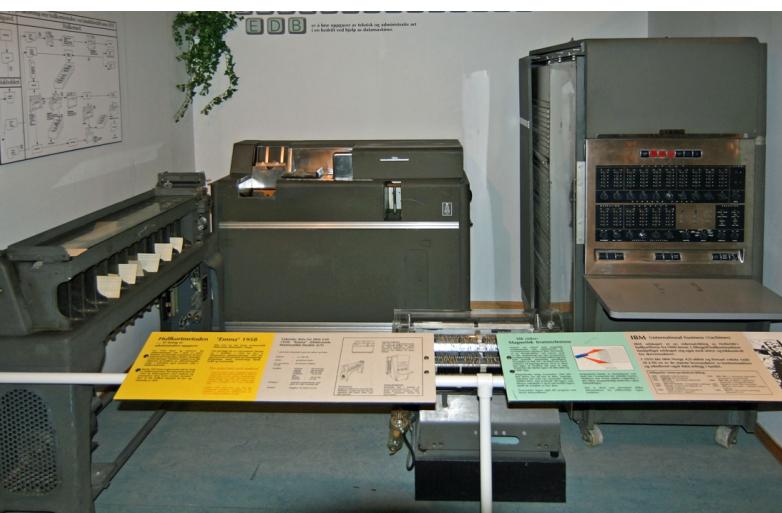
$$T_{\text{firstrun}} - T_{\text{arrival}}$$

- We would like any job to have low response time

$$T_{\text{firstrun}} - T_{\text{arrival}}$$

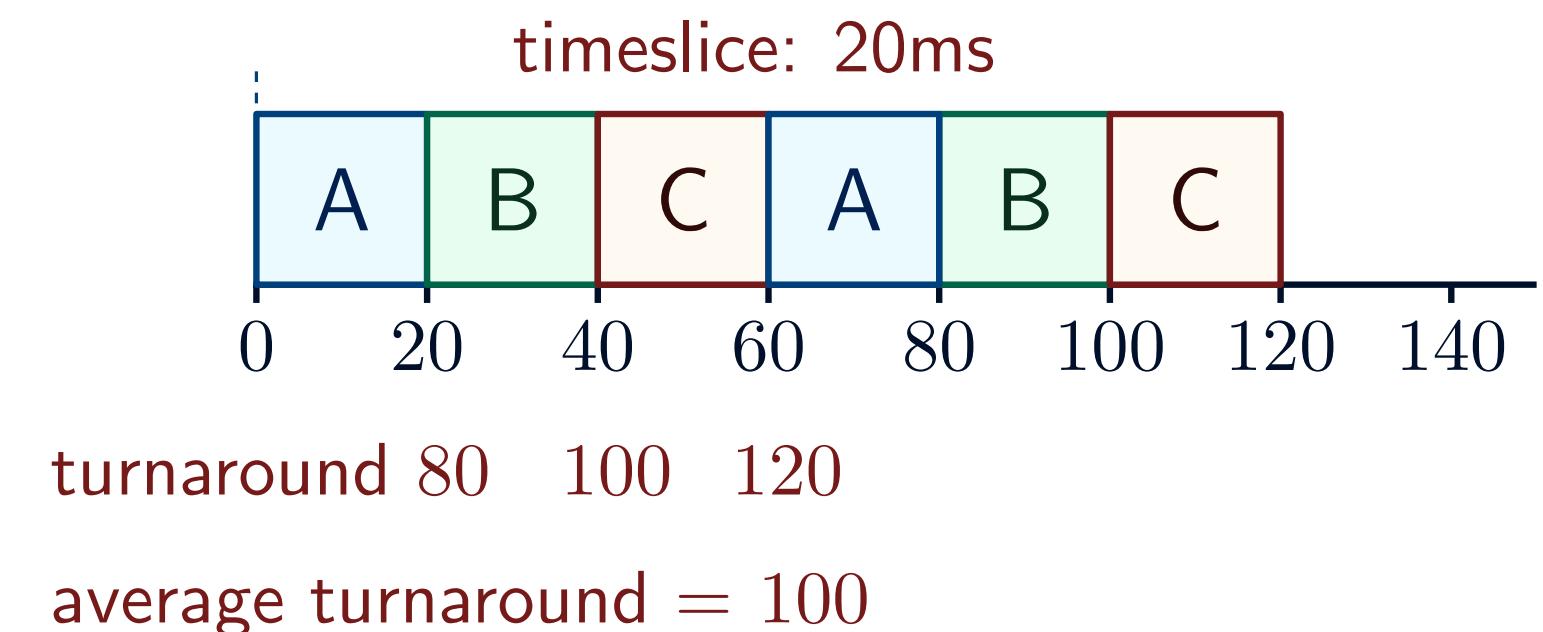
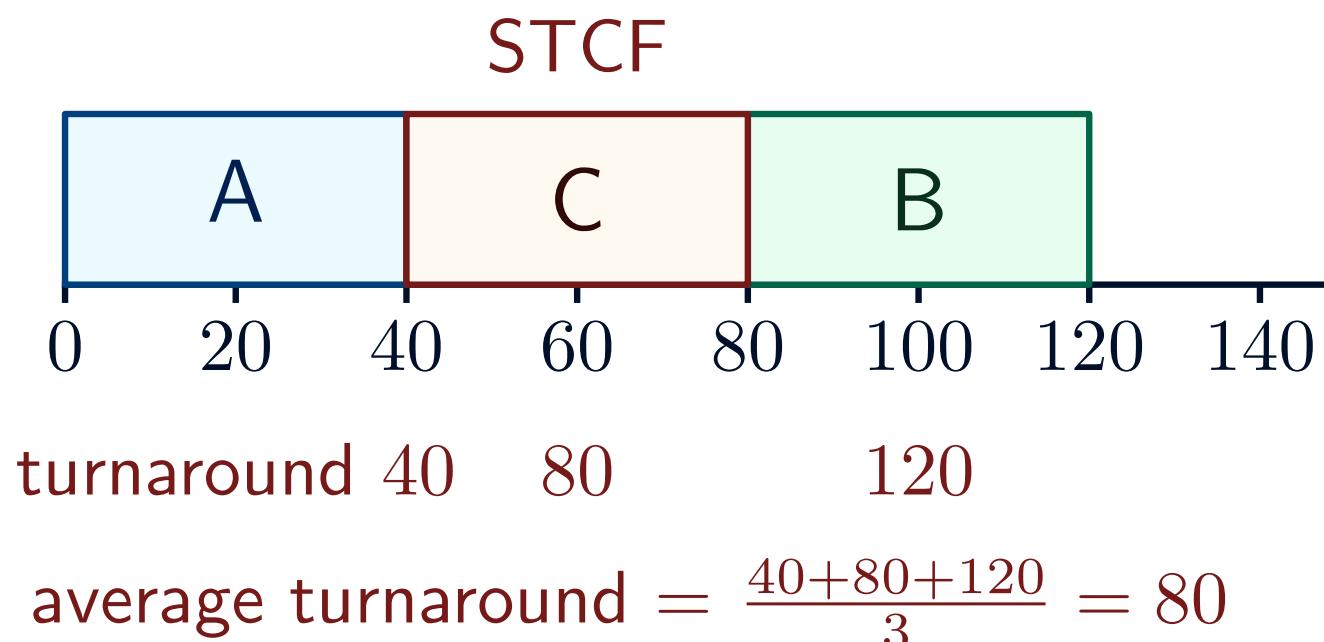
- STCF is not great for response time:

batch processing



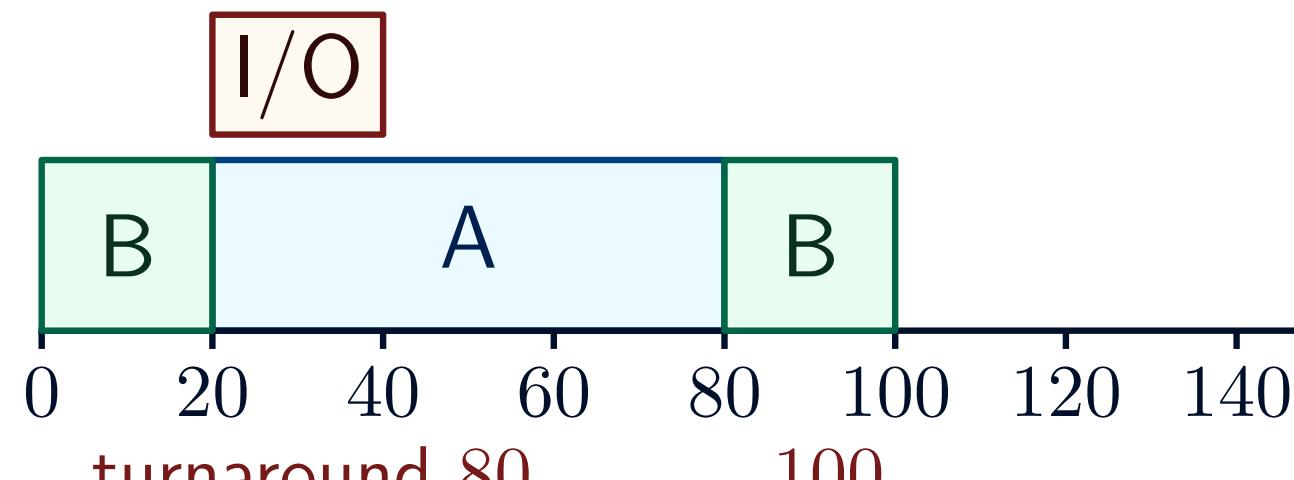
# Round Robin

- **Algorithm:** Run any job for a short timeslice, then put it in the back of the queue
- Need to balance the length of the slice
  - Too long—bad response time
  - Too short—context switching dominates
- Round Robin is not great for turnaround time

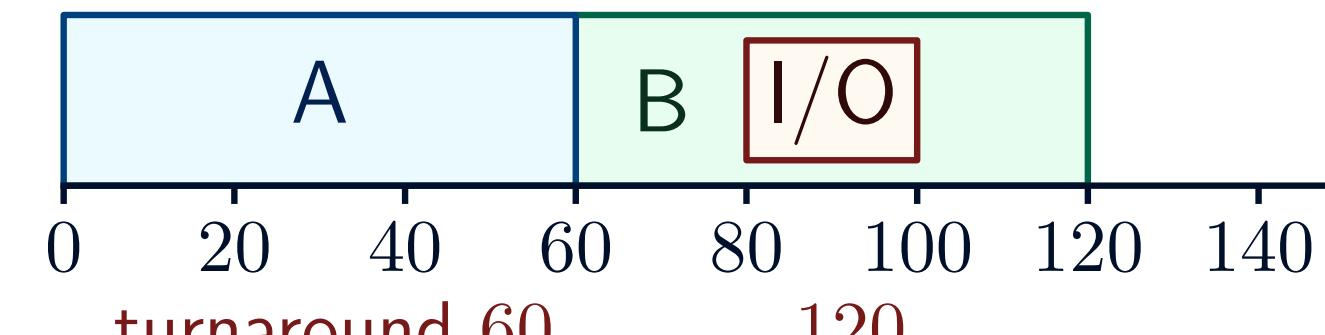


# Incorporating I/O

- While a job waits for I/O, run something else
- Putting jobs with I/O first is fine with turnaround time
- Jobs that do I/O a lot tend to use CPU less



average turnaround = 60

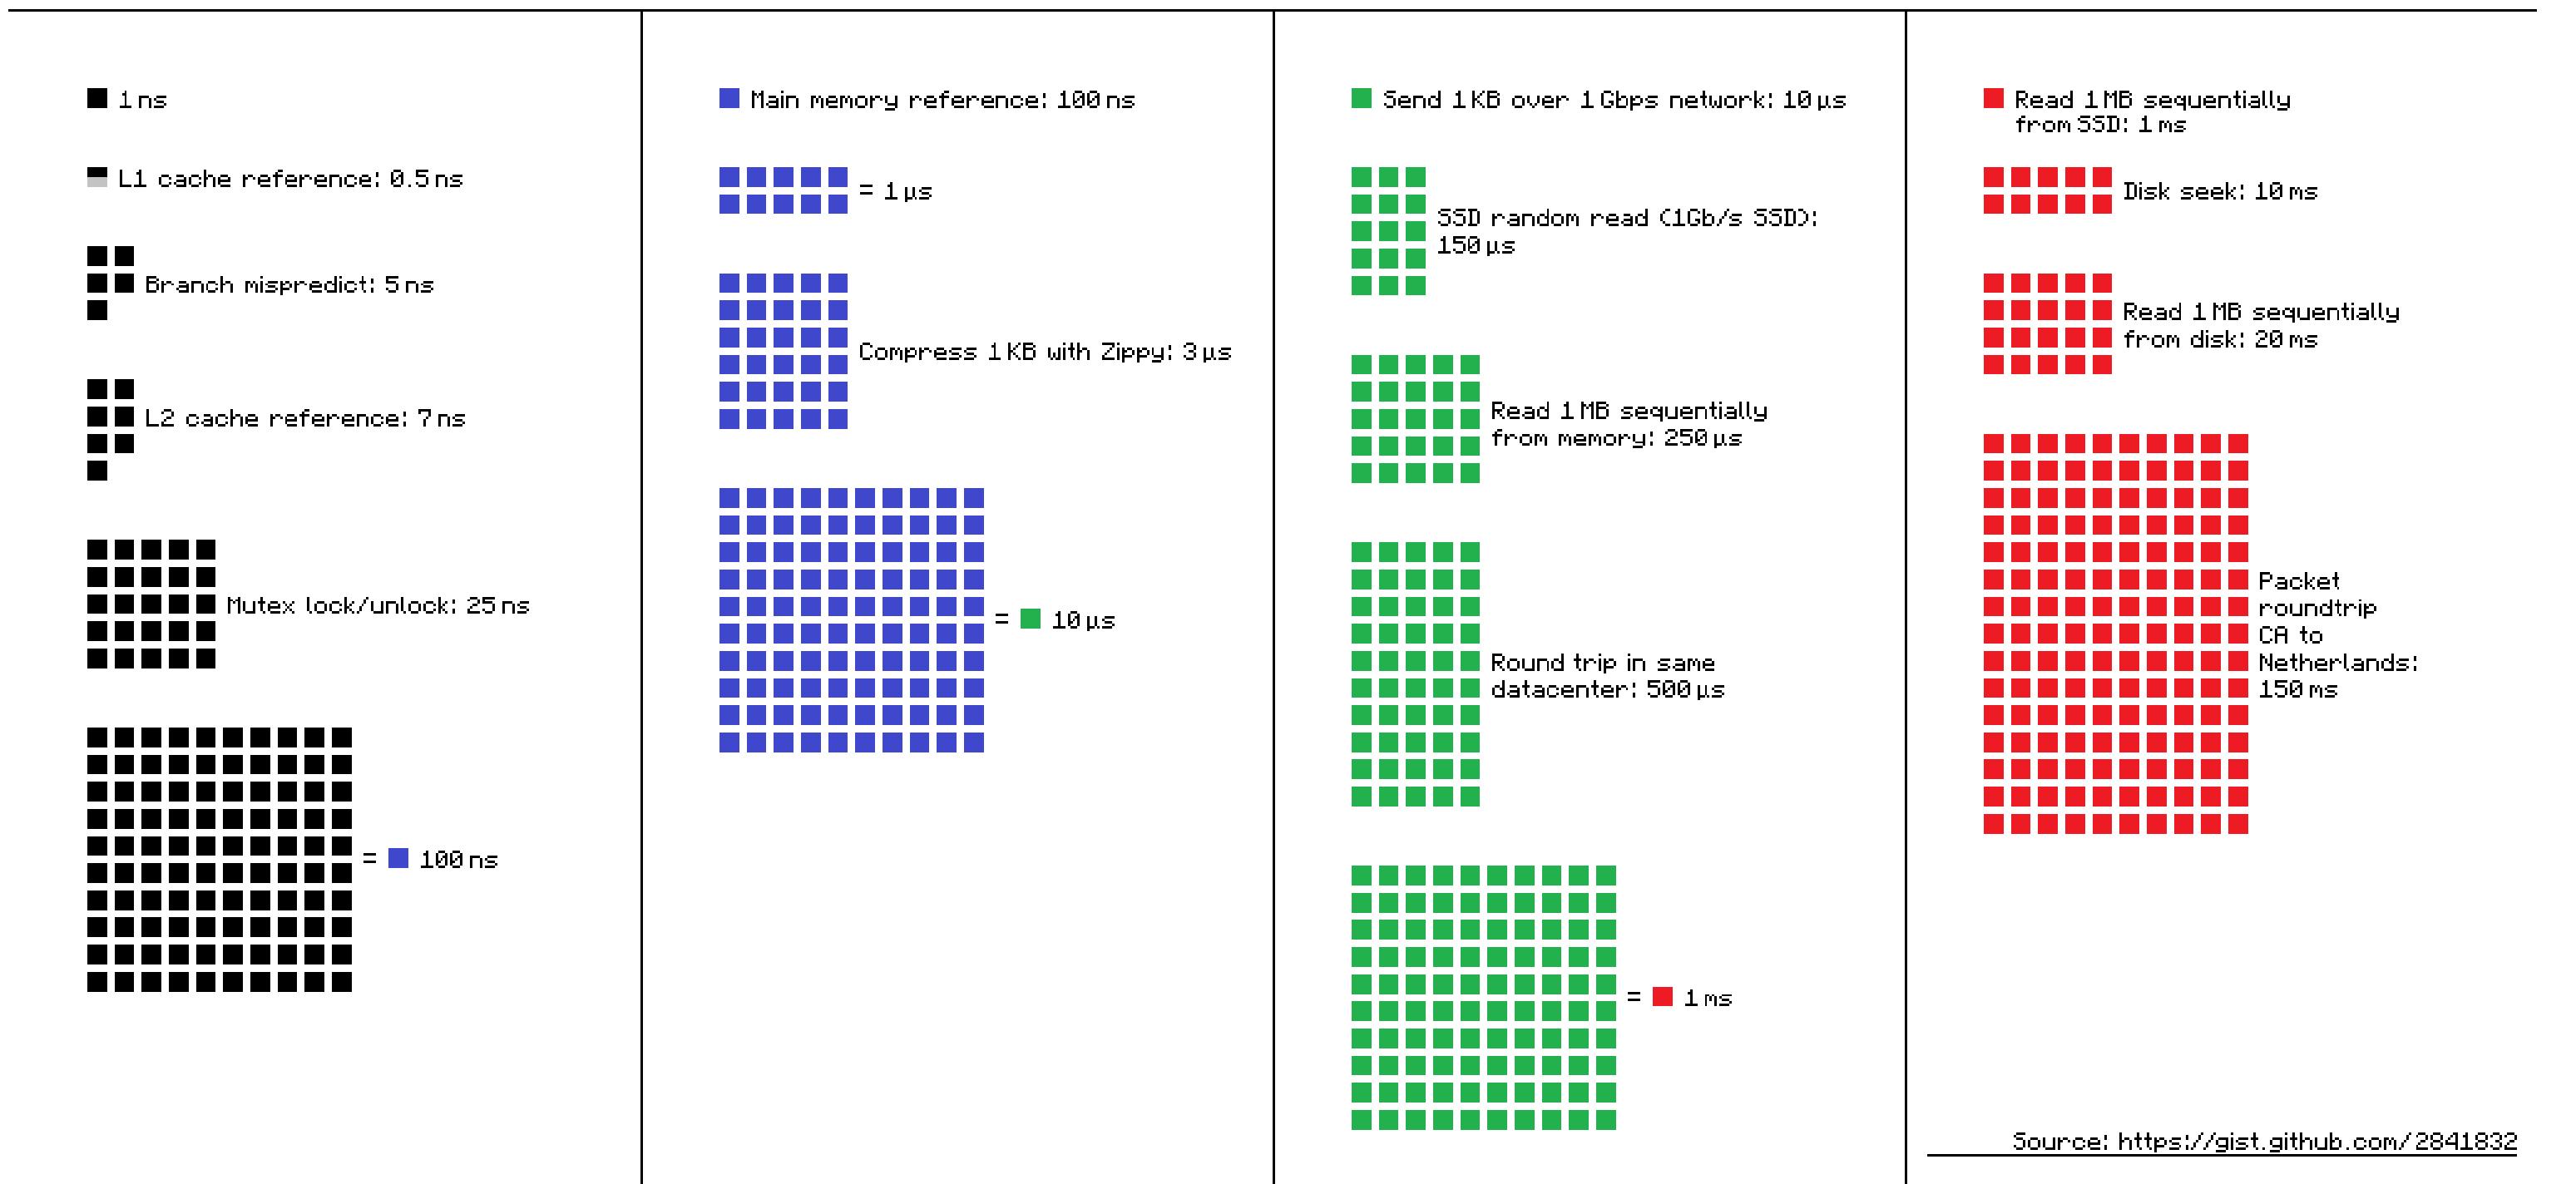


average turnaround = 60

same turnaround time, better response time

# Why wait on I/O

## Latency Numbers Every Programmer Should Know



# So far

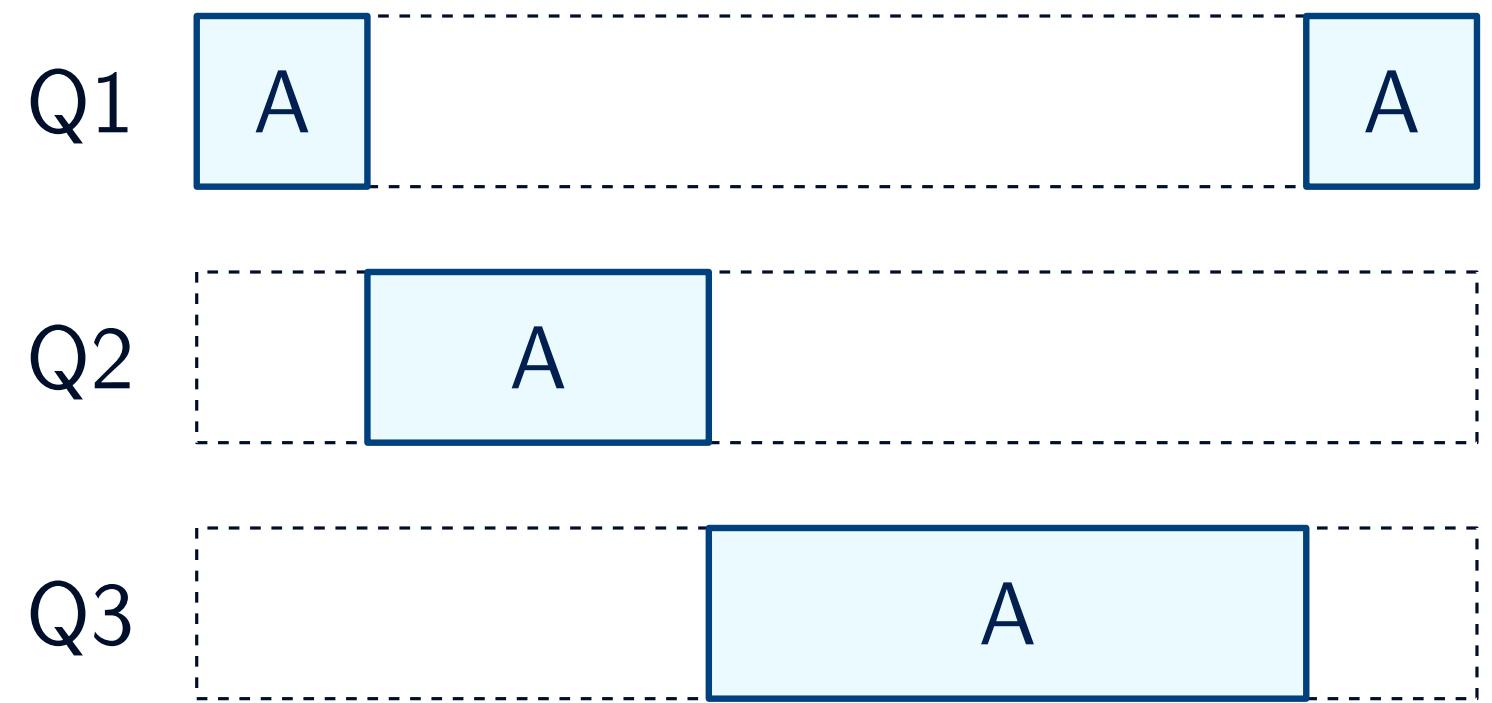
- Two measures of success: performance and interaction
  - Both need to be good for an all-purpose OS
  - Knowing running times in advance is a very strong assumption
  - Running low-compute I/O-focused jobs first is a good idea
- 
1. All jobs arrive at the same time
  2. All jobs have the same length
  3. The jobs cannot be interrupted
  4. The jobs do not wait on I/O
  5. The length of a job is known in advance

# Multi-Level Feedback Queue

- Proposed in 1962 by Corbato et al.
  - Good turnaround time while also good response time
  - “Learns” from history
  - Baseline approach for in Windows, FreeBSD, Solaris
- 
1. All jobs arrive at the same time
  2. All jobs have the same length
  3. The jobs cannot be interrupted
  4. The jobs do not wait on I/O
  5. The length of a job is known in advance

# The Queues in MLFQ

- There are several queues, modelling different priorities
- Run the processes from the highest non-empty queue in RR
- When a process comes, put it in the topmost queue
- If a process uses up its allotment, move it down
- Priority boost: after time  $S$ , move all jobs to topmost queue



# Before you go

- More details from today: Chapters 7, 8 in the OSTEP book
  - Try out the simulators there
- Mandatory Assignment 1
  - open Friday **September 12**
  - deadline Friday **September 26**
  - description and delivery on Mitt
  - you will do a few experiments in C
- Friday: more on MLFQ
- Friday: starting from fairness in scheduling Linux and its Completely Fair Scheduler (CFS)