

INF113: Introduction to Memory

Kirill Simonov
01.10.2025



Program address space

How does the memory of a process look like?


- First **static** data: code, global vars
- Then the **heap**: dynamic allocation

```
int* a = malloc(n * sizeof(int));
```

- Free space afterwards
- At the end, the **stack**: local variables

```
void rec(int d) {  
    ...  
}
```

- Both heap and stack will grow dynamically



Memory issues


- Memory leaks: explicitly free allocated memory

```
int* a = malloc(n * sizeof(int));  
...  
free(a);
```

- Even in higher-level languages:
 - Automatic deallocation via garbage collector
 - But an object will not be deleted while there is still a reference to it
- Dangling pointers: make sure the address is still available

~~```
int* add(int a, int b) {
 int sum = a + b;
 return ∑
}
```~~

- Addressing invalid memory might not produce any error!



# Valgrind

- Install `sudo apt install valgrind`
- Compile your C program with `gcc -g -o program program.c`
- Run within valgrind `valgrind ./program`
- Look for leaks and errors:

```
==95170== Invalid write of size 4
==95170== at 0x109199: main (merror2.c:7)
==95170== Address 0x4a87068 is 0 bytes after a block of size 40
==95170== at 0x4846828: malloc (in /usr/...)
==95170== by 0x10917E: main (merror2.c:5)


==95092== LEAK SUMMARY:
==95092== definitely lost: 40 bytes in 1 blocks

==94931== Conditional jump or move depends on uninitialised value(s)
```

# Memory virtualization

OS provides a virtual memory address space to each process

- **Abstraction:** The process can access virtual addresses as if it had direct control over the physical memory
- **Protection:** The process cannot access memory of other processes/OS
- **Efficiency:**
  - Transitioning from a virtual address to the physical address should be very fast
  - Avoid moving large chunks of memory
  - Avoid allocating much more memory to a process than it needs
  - ...




# Address translation

```
int main() {
 int x = 3000;
 x = x + 3;
}
```

```
128: movl 0x0(%ebx), %eax
132: addl $0x03, %eax
135: movl %eax, 0x0(%ebx)
```

- CPU cycle:
1. Fetch instruction at 128
  2. Execute: Move from 15678 to eax
  3. Fetch instruction at 132
  4. Execute: Add to eax
  5. Fetch instruction at 135
  6. Execute: Move eax to 15678

program address space



# Address translation

```
int main() {
 int x = 3000;

 x = x + 3;
}
```


physical memory:

The program is not  
actually at address 0!

```
128: movl 0x0(%ebx), %eax
132: addl $0x03, %eax
135: movl %eax, 0x0(%ebx)
```

CPU cycle:  
(intended)

1. Fetch instruction at **16128**
2. Execute: Move from **31678** to eax
3. Fetch instruction at **16132**
4. Execute: Add to eax
5. Fetch instruction at **16135**
6. Execute: Move eax to **31678**



# Address translation

```
int main() {
 int x = 3000;
 x = x + 3;
}
```


physical memory:

The program is not  
actually at address 0!

```
128: movl 0x0(%ebx), %eax
132: addl $0x03, %eax
135: movl %eax, 0x0(%ebx)
```

How does the program/CPU know how to adjust instructions,  
depending on where the program is loaded in memory?

3. Fetch instruction at **10132**
4. Execute: Add to eax
5. Fetch instruction at **16135**
6. Execute: Move eax to **31678**



# Attempt 1: Loader

- OS can run the **loader** once it knows the target memory location
- Loader will increment all addresses in the program code by the offset in the memory
- Pros:
  - No additional hardware requirements
  - Small overhead after loading
- Cons:
  - Complicated to implement
  - Costly preprocessing
  - **No protection!**

16KB  
17KB

Code


```
movl 0x0(%ebx), %eax
```



```
addl $0x1000, %ebx
movl 0x0(%ebx), %eax
subl $0x1000, %ebx
```

# Attempt 2: Context switch

- We already have **limited direct execution**: pass control to the OS when needed
- OS knows offsets of all address spaces in memory, so it can both
  - Adjust the address from virtual to physical
  - Check that the address is within the program's space
- Cons: context switch is **100x** slower than addressing memory...



# Base and bound


- **Assume:** Each process occupies a contiguous, dedicated part of the memory
- Translating addresses is fairly simple:

```
physical address = virtual address + base
```

- To check validity, just compare against the size:

```
if virtual address >= bound
 fail
```



- base and bound set in hardware  
⇒ automatic address translation and protection!



# Base and bound: Hardware

- Memory management unit (MMU) on the CPU has base and bound registers
- When OS loads a program, it sets the registers to the program address space bounds
- When a program requests address, MMU converts virtual address into physical and checks validity

```
if virtual address >= bound
 fail
physical address = virtual address + base
```




# Base and bound: Hardware

- Memory management unit (MMU) on the CPU has base and bound registers
- When OS loads a program, it sets the registers to the program address space bounds
- When a program requests address, MMU converts virtual address into physical and checks validity

```
if virtual address >= bound
 fail
physical address = virtual address + base
```

- Upon a context switch, base and bound need to be stored too
- Easy to relocate program: just update the registers




# Simple VM

- **Assume:** Each process occupies a same-sized contiguous part of the memory
- OS can keep a **free list**: list of free slots in the memory
  - New process: Take first slot, remove from free list
  - Process over: Add its slot to the free list

Free list:


80 → 96 → 112

Memory



# Simple VM

- **Assume:** Each process occupies a same-sized contiguous part of the memory
- OS can keep a **free list**: list of free slots in the memory
  - New process: Take first slot, remove from free list
  - Process over: Add its slot to the free list
- Issues so far:
  - Processes may have very different memory requirements
  - Lots of memory within process slots is wasted



# Summary

- Memory organization within a single process: code, heap, stack
- Using valgrind to detect memory issues
- The goals of memory virtualization
- Base and bound registers for address translation, and protection
- Simple VM based on free list
- **Next time:** Make it more efficient