

INF113: Segmentation and Free Space

Kirill Simonov
03.10.2025




Base and bound: Reminder

- **Mechanism:** Store boundaries of process space in registers
- Translating addresses:

```
physical address = virtual address + base
```


- Checking validity:

```
if virtual address >= bound  
    fail
```





Base and bound: Issues

- Issues with same-sized address spaces:
 - Each address space is small
 - Most processes will not use most of their free memory—**fragmentation**
- If we made address spaces different-sized but still contiguous?
- Each process should have a large chunk of free memory within
 - If e.g. heap needs more memory, whole process has to be relocated



Segmentation



Implementing segmentation


- Store the segment table on registers:

	base	bound
code	16KB	1KB
heap	22KB	1KB
stack	31KB	1KB

- To translate a virtual address, use the bounds for the respective segment

virtual address	physical address
0.5KB in heap	→ 22.5KB
0.5KB in code	→ 16.5KB
2KB in stack	→ fail

2KB in stack →




How to know the segment?

- Virtual addresses should still be simple
get value at 0x014f **vs** get value at (code, 0x014f)

- Possible solution: Use leading bits to describe the segment


	15 14 13 12 11 10 9 8 7 6 5 4 3 2 1 0	
segment	offset	
00: code	260	0x0104
01: heap	544	0x4220
11: stack	11	0x800b



How to know the segment?


- Virtual addresses should still be simple
get value at 0x014f **vs** get value at (code, 0x014f)

- Possible solution: Use leading bits to describe the segment



- MMU can quickly compute the physical address

```
segment = address >> 14;  
offset = address & 0x3fff;  
if (offset >= bounds[segment])  
    fail;  
phys_address = base[segment] + offset;
```





What about stack?

- Store direction of growth

	base	bound	dir
code	16KB	1KB	1
heap	22KB	1KB	1
stack	32KB	1KB	0

- Interpret offset as a negative number



Segment protection

- Specify which operations are allowed on the segments


	base	bound	dir	rights
code	16KB	1KB	1	r-x
heap	22KB	1KB	1	rw-
stack	32KB	1KB	0	rw-

- MMU additionally checks operation type

```
...
if (op & rights == 0
    || offset >= bounds[segment])
    fail;
...


```

- Read-only segments can be shared!




Why segmentation?

- No “reserved” gaps within individual address spaces
- More flexibility → filling gaps more easily
- Can relocate segments individually → less overhead
- Additional savings via sharing segments



Why segmentation?

- No “reserved” gaps within individual address spaces
- More flexibility → filling gaps more easily
- Can relocate segments individually → less overhead
- Additional savings via sharing segments
- Still, lots of room for fragmentation...
- Real solution next week: **Paging**



Segmentation fault


- The scariest thing a programmer can see...
- Simply means that the segmentation table is violated

	base	bound	dir	rights
code	16KB	1KB	1	r-x
heap	22KB	1KB	1	rw-
stack	32KB	1KB	0	rw-



Free space management

- How to manage the memory well?
 - Find a place for the next process/segment efficiently
 - Avoid fragmentation
- Easy when all requests have same size
 - Also holds for paging-based VM in OS




Free space management



- How to manage the memory well?
 - Find a place for the next process/segment efficiently
 - Avoid fragmentation
- Easy when all requests have same size
 - Also holds for paging-based VM in OS
- More fun when sizes vary
 - Useful for pure segmentation-based VM
 - Or malloc

can't place!




Free space management




- How to manage the memory well?
 - Find a place for the next process/segment efficiently
 - Avoid fragmentation
- Easy when all requests have same size
 - Also holds for paging-based VM in OS
- More fun when sizes vary
 - Useful for pure segmentation-based VM
 - Or malloc

perfect!
but what if...



Free space management

- How to manage the memory well?
 - Find a place for the next process/segment efficiently
 - Avoid fragmentation
- Easy when all requests have same size
 - Also holds for paging-based VM in OS
- More fun when sizes vary
 - Useful for pure segmentation-based VM
 - Or malloc




Malloc

- What happens then?


```
int* a = malloc(10 * sizeof(int));
```

- malloc does not just go to the OS to ask for memory—syscall/context switch is too slow!
- Instead, malloc will try to fit the request into already allocated memory
- Only if this fails, malloc will ask the OS to extend the heap via e.g. sbrk
- If this fails too, malloc will return NULL—watch out!




Malloc vs OS VM

- Both need memory management logic
 - Where to place the next request?
 - How to keep track of free/occupied memory?




Malloc specifics

- malloc cannot relocate segments
 - There are pointers in the code, and they will be invalidated




Malloc specifics

- malloc cannot relocate segments
 - There are pointers in the code, and they will be invalidated
- malloc can ask for more memory
 - But it should only ask if absolutely necessary
- malloc cannot call malloc!
 - The internal malloc data structure has to be allocated explicitly



Free list




- Let us store all free ranges in the list—**free list**



- Split:** Allocating a segment splits the range




- Coalesce:** Once a segment is freed, merge adjacent free ranges



Header blocks



- When we call `free(10)`, we only supply the address—how to determine the size?
- Store information directly in the memory!
- Each allocated block comes with a header
 - Size—to know what is occupied
 - Magic number—to make sure the header is actually a header

```
int* ptr = malloc(7 * sizeof(int));
```



Embedding free list


- Apply the same trick for the free list
Store free range info at the start of the range
- When a range splits, insert new node



Embedding free list

- Apply the same trick for the free list
Store free range info at the start of the range
- When a range splits, insert new node


```
int* ptr = malloc(4 * sizeof(int));
```



Embedding free list

- Apply the same trick for the free list
Store free range info at the start of the range
- When a range splits, insert new node


```
int* ptr = malloc(4 * sizeof(int));  
int* ptr2 = malloc(2 * sizeof(int));
```



Embedding free list

- Apply the same trick for the free list
 - Store free range info at the start of the range
- When a range splits, insert new node

```
int* ptr = malloc(4 * sizeof(int));
int* ptr2 = malloc(2 * sizeof(int));
free(ptr);
```




Embedding free list

- Apply the same trick for the free list
Store free range info at the start of the range
- When a range splits, insert new node

```
int* ptr = malloc(4 * sizeof(int));  
int* ptr2 = malloc(2 * sizeof(int));  
free(ptr);  
free(ptr2);
```

- Coalesce when a block is freed




Modifying unallocated
memory may break malloc!

Basic strategies

How to pick the free range?

- **Best fit:** Pick the smallest range that fits
- **Worst fit:** Pick the largest range that fits
- **First fit:** Pick the first range that fits
- **Next fit:** Pick the nearest range that fits after the last fit



Summary

- Segmentation: split address space into logical segments
 - The idea is heavily used in OS, but there's more
- Filling memory with arbitrary-length requests is a complicated task with various strategies
- Basic idea behind how malloc manages free/occupied ranges
 - The actual implementation is much more complicated
 - Focus on efficiency and specific fitting strategies

sploitfun.wordpress.com/2015/02/10/understanding-glibc-malloc/

- **Next time:** Paging
- **Homework:** After Chapters 15 and 16

