

INF113: Paging

Kirill Simonov
08.10.2025




Base and bound: Reminder

- Process takes a contiguous chunk of memory, start and end stored in registers
- Translating addresses is easy:

```
if virtual address >= bound  
    physical address = virtual address + base
```

- **Main issue:** Internal fragmentation
 - Each process needs lots of “spare” memory within its address space
- Address space may never be as large as physical memory




Segmentation: Reminder

- Each logical segment takes a contiguous chunk of memory
- For address translation, look up the segment table

	base	bound	dir	rights
code	16KB	1KB	1	r-x
heap	22KB	1KB	1	rw-
stack	32KB	1KB	0	rw-


- **Main issue:** External fragmentation
 - Hard to arrange all process segments without gaps



Paging

- Partition address space into small same-sized **pages**, e.g., 4KB each

process address space




$vp0 \mapsto pf3$
 $vp1 \mapsto pf7$
 $vp2 \mapsto pf5$
 $vp3 \mapsto pf2$

- Each page is stored at some page frame

- Partition physical memory into **page frames** of the same size


physical memory





Page table

- OS stores the **page table**, mapping virtual pages to physical page frames

	VPN	PFN	
virtual page number	0	3	
1	7	physical frame number	
2	5		
3	2		




process address space



Address translation

- When a process requests memory by the virtual address, MMU converts it to the physical address via the page table


assume page
is 16 bytes



Address translation

- When a process requests memory by the virtual address, MMU converts it to the physical address via the page table


assume page
is 16 bytes



Address translation

- When a process requests memory by the virtual address, MMU converts it to the physical address via the page table

assume page
is 16 bytes





Address translation

- If the process performs

```
movl virt_addr, %eax
```


- MMU translates the address as follows

```
offset = virt_addr & 0x3f  
VPN = (virt_addr & 0xc0) >> 6  
phys_addr = (PT[VPN] << 6) | offset
```




Larger address space

- In practice, processes will have huge address spaces
 - Most pages will not be used
 - Need to know whether the page is actually there




VPN	PFN	valid
0	4	1
1	-	0
2	1	1
3	-	0
4	-	0
5	-	0
6	-	0
7	10	1

page table




Page Table Entry (PTE)

- OS stores a PTE record for each page
 - PFN, page frame number
 - validity bit—is page allocated?
 - present bit—is page in RAM?
 - protection bits
 - ...




Benefits of paging

- Great abstraction
 - Each process has huge and (seemingly) contiguous address space
 - But OS can arrange individual pages very flexibly on the RAM




Benefits of paging

- Great abstraction
 - Each process has huge and (seemingly) contiguous address space
 - But OS can arrange individual pages very flexibly on the RAM
- Easy free space management for OS
 - Any frame fits any page, so a page can be placed in any free slot




Benefits of paging

- Great abstraction
 - Each process has huge and (seemingly) contiguous address space
 - But OS can arrange individual pages very flexibly on the RAM
- Easy free space management for OS
 - Any frame fits any page, so a page can be placed in any free slot



Benefits of paging

- Great abstraction
 - Each process has huge and (seemingly) contiguous address space
 - But OS can arrange individual pages very flexibly on the RAM
- Easy free space management for OS
 - Any frame fits any page, so a page can be placed in any free slot
- Almost no fragmentation
 - As soon as a page frame is free, more pages can be allocated
 - Only the remainder of the last page in each logical segment has to be wasted; just a few KBytes per process





Benefits of paging

- Great abstraction
 - Each process has huge and (seemingly) contiguous address space
 - But OS can arrange individual pages very flexibly on the RAM
- Easy free space management for OS
 - Any frame fits any page, so a page can be placed in any free slot
- Almost no fragmentation
 - As soon as a page frame is free, more pages can be allocated
 - Only the remainder of the last page in each logical segment has to be wasted; just a few KBytes per process
- However, there is a price...
 - Memory for storing page tables
 - Each memory access has an overhead of reading the page table

Storing page tables

- Larger address space means larger page table
 - 32-bit addresses—at most 4GB—one million 4KB pages
 - For each page, store at least the PFN—4 bytes
 - At least 4MB for the page table
- Each process has its own page table—400MB for 100 processes!

need to be smarter about
storing the PT—will see
the solution next time



Memory access overhead

- MMU translates the address as follows

```
offset = virt_addr & OFFSET_MASK  
VPN = (virt_addr & PAGE_MASK) >> PAGE_SHIFT  
phys_addr = (PT[VPN] << PAGE_SHIFT) | offset
```




looking up PFN of the given page


- Every memory access now spawns another memory access—to the page table
- Memory access is already 20x–200x slower than referencing caches/registers, now another factor of 2!
- Too many pages to fit into cache/registers
- Maybe at least cache some VPN-PFN pairs?

Translation-lookaside buffer (TLB)

- MMU keeps a small cache of page-table lookups
- To get the PFN from a given VPN:



- If hit, then PFN is found very fast
- If miss, we actually pay the cost of extra memory access
But next operations with the same VPN will be hits



Why TLB helps

- Example: 16-byte pages


```
int i, sum = 0;  
for (i = 0; i < 10; i++) {  
    sum += a[i];  
}
```

i = 0: address 0110100 → page 3 → miss
i = 1: address 0111000 → page 3 → hit
i = 2: address 0111100 → page 3 → hit
i = 3: address 1000000 → page 4 → miss
i = 4: address 1000100 → page 4 → hit
i = 5: address 1001000 → page 4 → hit

...

VPN	PFN
3	7
4	9
5	8

TLB



Typical TLB

- A few thousand entries
 - Each entry covers one 4KB page
 - So millions of addresses are cached at all times
- Extremely fast VPN lookup
 - Hit: 0.5–1 cycles
 - Miss: 10–100 cycles
- Small miss rate
 - 0.01%–1% “on average”
 - 20%–40% for sparse applications
- Better locality in your program—fewer misses!
 - Program only uses few MBs: all its pages numbers will be cached
 - Accessing an array consecutively—always great
 - Array by random indices—also perfect hits if fits in a few MBs


Example:

Assume hit is 1 cycle,
miss is 30 cycles,
miss rate is 1%.

Average TLB overhead is then
 $1 + 30 \cdot 0.01 = 1.30$ cycles per query

TLBs and context switches


- When another process takes over, old TLB is not helpful
 - B accesses a value on page 0
 - page 0 is converted by TLB to frame 3
 - frame 3 is out-of-bounds for B
 - crash!



B	8	4	9	-
	0	1	2	3

TLBs and context switches

- When another process takes over, old TLB is not helpful
- It makes little sense to save TLB and then restore it
 - Just letting it fill naturally will have the same cost
- Simple solution is to clear TLB on every context switch
 - Nanoseconds vs milliseconds, so not too bad
- Many systems support process ID in TLB
 - Then entries for A can be safely left
 - Either they will be eventually overwritten
 - Or process A is restored quickly and most of its TLB survives—win!



Which pages to cache?

When a new entry comes, how to decide which old one to overwrite?

- Has to be a simple strategy
 - Fast
 - No information about the future
 - Very limited space to store past data
- **Example:** Least recently used—kick the entry that stayed unused for too long
- **Example:** Random—kick a random entry
- Similar question for swapping policies
Which memory pages can be safely stowed on the disk?

Summary

- Paging discretizes the address space into small same-sized chunks
- This achieves great flexibility in memory usage
→ and little fragmentation
- Also greatly simplifies free space management for the OS
- The overhead of paging has to be dealt with
 - Extra memory accesses—solved by TLB, caching translations
 - Too much memory for the page table—will see better ways of organizing the page table next time
- Take a look at homework after Chapters 18 and 19
- Mandatory Assignment 2 starts Friday!