

# INF113: Multi-level Page Tables and Swap

Kirill Simonov  
10.10.2025





# Mandatory Assignment 2

- Starts today afternoon
- Deadline in two weeks, Friday Oct 23
- You will experiment with memory management
- In case of issues, ask on Discord or attend group session
- Rule clarification: 10% of final grade, but no pass limit
- New: declaring usage of LLMs required, following the faculty policy

# Paging: Reminder

- Address space is partitioned into 4KB pages, and memory into 4KB frames  
Pages are arbitrarily mapped into frames
- Mapping is stored in the memory as the page table
- Issues:
  - Extra memory accesses—solved by TLB ✓
  - Page tables take lots of memory ?



# Linear page tables in memory


- **Issue 1: size**

4GB memory, 4KB pages, 4B PTE, 100 processes  $\mapsto$  400MB in page tables

- **Issue 2: allocating memory for a page table**

- Single page table takes 4MB, has to be consecutive in memory
- Page table itself is not paged—bad for memory management

32-bit address space




page table




one million page translations in the array

physical memory



# Attempt 1: Larger pages


- Let's make the page 16x larger: 64KB
  - $4\text{MB}/16 = 256\text{KB}$  for a single page table = 4 pages, 25MB for all PTs
- Also increases fragmentation...
  - Now every memory allocation potentially wastes 64KB
  - 4 pages per process, 100 processes = 25MB




# Attempt 2: Data structures

- We know lots of pages may be empty—maybe use a sparser DS?
  - Linked list?
  - Binary search tree?

32-bit address space



page table




linked list only stores existing pages

finding PFN requires cycling through all entries!

# Attempt 2: Data structures

- We know lots of pages may be empty—maybe use a sparser DS?
  - Linked list?
  - Binary search tree?
- Fixes space issues, but overhead for a memory access is terrible



# Multi-level page table

- Let's split the page table into page-sized chunks!

VPN	PFN	valid
0	4	1
1	1	1

PT	PFN	valid
0	3	1
1	-	0
2	-	0
3	6	1

page directory

VPN	PFN	valid
6	-	0
7	10	1

VPN	PFN	valid
0	4	1
1	1	1



VPN	PFN	valid
2	-	0
3	-	0

VPN	PFN	valid
4	-	0
5	-	0

VPN	PFN	valid
6	-	0
7	10	1



# Address translation: MLPT

- Compute page index in PT and fetch its location

```
PD_index = virt_addr & 0x30
```

```
PT = PD[PD_index]
```

- Find target PFN by index within PT page


```
PT_index = virt_addr & 0x08
```

```
PFN = PT[PT_index]
```

- Add offset to obtain final address

```
offset = virt_addr & 0x07
```

```
phys_addr = (PFN << 3) | offset
```



	PFN	valid
0	3	1
1	-	0
2	-	0
3	6	1

page directory

	PFN	valid
0	4	1
1	1	1


PT page 0

6	-	0
7	10	1

PT page 3


physical address  
is 1010001

physical memory



# Benefits of MLPT


- Size of page table scales with used memory pages
  - Page table size is 4KB (PD) + 4KB for each used PT page
  - Each PT page handles a thousand virtual pages
- Page table is also stored in pages
  - PD and individual PT pages can be stored at any frame
  - PD keeps references to where each individual PT page is stored
- Issue: more memory reads
  - Each memory access turns into three
  - However, TLB still helps
- Issue: page directory size
  - PD might not always fit in a single page
  - Example: 32-bit addresses, 4KB page, 4B per PTE/PDE  
20-bit page index, 1024 PTE in a PT page, 1024 PT pages described in PD



# MLPT as a tree



- Two-level page table is a search tree of depth 2

physical memory



# More levels


- MLPT may also have, e.g., 3 levels



- Handles larger memory space and/or smaller page size
- Downside: one extra memory access per level

# Beyond physical memory

- It may be too generous to keep all pages of all processes in memory at all times
  - Some processes may be sleeping for a while
  - Some pages may be almost never in use
- **Swap space:** portion of the disk where memory pages can be stored
  - With paging, we can swap any page out, without affecting the rest




# Present bit

- **Present bit** in the page table marks whether the page is in memory
- Access to an invalid page—crash
- Access to a not present page—call OS to load the page from disk


reading from disk is slow

enough to allow context switch!




	VPN	PFN	valid	pr.	prot.
code	0	4	1	1	r-x
heap	1	-	0	0	---
heap	2	1	1	1	rw-
heap	3	-	1	0	rw-
heap	4	-	0	0	---
heap	5	-	0	0	---
heap	6	-	0	0	---
stack	7	10	1	1	rw-

page table




# Between memory and swap

- A page is moved from swap upon access
  - Sometimes also in advance, e.g., the next code page



# Between memory and swap

- A page is moved from swap upon access
  - Sometimes also in advance, e.g., the next code page
- A frame must be free for the page to be brought back
- OS will try to keep some memory free proactively
  - If less than **low** frames are free, OS removes pages to reach **high**
  - Moving pages to disk in bulk saves time—buffering



# Swapping policies

How to decide which page to swap?

- With good policy OS multiplies memory size for free
- With bad policy every memory access can turn into a disk access...

memory 100 ns

SSD 0.1 ms

HDD 10 ms

- Can see memory as a **cache** of pages
  - Fast but small memory
  - Large but slow disk
  - Want to maximize the number of hits
- Same considerations as for other caches: TLB, L1/L2, ...
- Swap specialty: misses are very costly!

## Example:

Memory vs HDD, 10% miss rate

Average access =  $100\text{ns} + 0.1 \cdot 10\text{ms} = 1.0001\text{ms}$

Memory vs HDD, 0.1% miss rate

Average access =  $100\text{ns} + 0.001 \cdot 10\text{ms} = 10.1\mu\text{s}$

# Example policies

- **Optimal:** swap page that is furthest in the future
- **Random:** swap a random page
- **FIFO:** swap the earliest page
- **LRU:** swap the least recently used page

query	cache	result	
0	empty	miss	
1	0	miss	optimal
2	0, 1	miss	
0	0, 1, 2	hit	4/9 ~ 44.4% miss rate
1	0, 1, 2	hit	
3	0, 1, 2	miss	
0	0, 1, 3	hit	
3	0, 1, 3	hit	
1	0, 1, 3	hit	


FIFO

5/9 ~ 55.6% miss rate

query	cache	result
0	empty	miss
1	0	miss
2	0, 1	miss
0	0, 1, 2	hit
1	0, 1, 2	hit
3	0, 1, 2	miss
0	1, 2, 3	miss
3	1, 2, 3	hit
1	1, 2, 3	hit


# Policies vs scenarios

- **No-locality:** 100 pages, full random
  - Nothing helps, except knowing the future




# Policies vs scenarios

- **No-locality:** 100 pages, full random
  - Nothing helps, except knowing the future
- **80-20:** 20% of pages get 80% of queries
  - LRU predicts “heavy hitters”




# Policies vs scenarios

- **No-locality:** 100 pages, full random
  - Nothing helps, except knowing the future
- **80-20:** 20% of pages get 80% of queries
  - LRU predicts “heavy hitters”
- **Looping-sequential:** 50 pages accessed in the same sequence, over and over
  - Worst-case for LRU



# Policies vs scenarios

- **No-locality:** 100 pages, full random
  - Nothing helps, except knowing the future
- **80-20:** 20% of pages get 80% of queries
  - LRU predicts “heavy hitters”
- **Looping-sequential:** 50 pages accessed in the same sequence, over and over
  - Worst-case for LRU
- **Clock policy:** approximates LRU by only keeping a single “accessed” bit
  - More realistic



# Memory usage tools

- pmap shows the memory map of a process

```
pmap -x <PID>
```

- free displays memory usage

```
free -h
```

- vmstat allows to keep track of memory, swap, cpu, also in real-time

```
vmstat 1
```

- swapon -s to show swap devices

```
swapon -s
```

- Use mem.c from Chapter 21 Homework to experiment with these

# Summary

- Multi-level page tables solve the space problem of page tables
- Levels provide trade-off between the size of table and extra memory accesses
- Paging makes it easy to swap unpopular pages to disk, saving space
- Good policy is needed to keep miss rates very low
- Modern OS would use some approximations to LRU
- Spread of SSDs shifts balance:  
misses are less costly → more headspace for policies
- Look at Chapters 20–22 for more details