

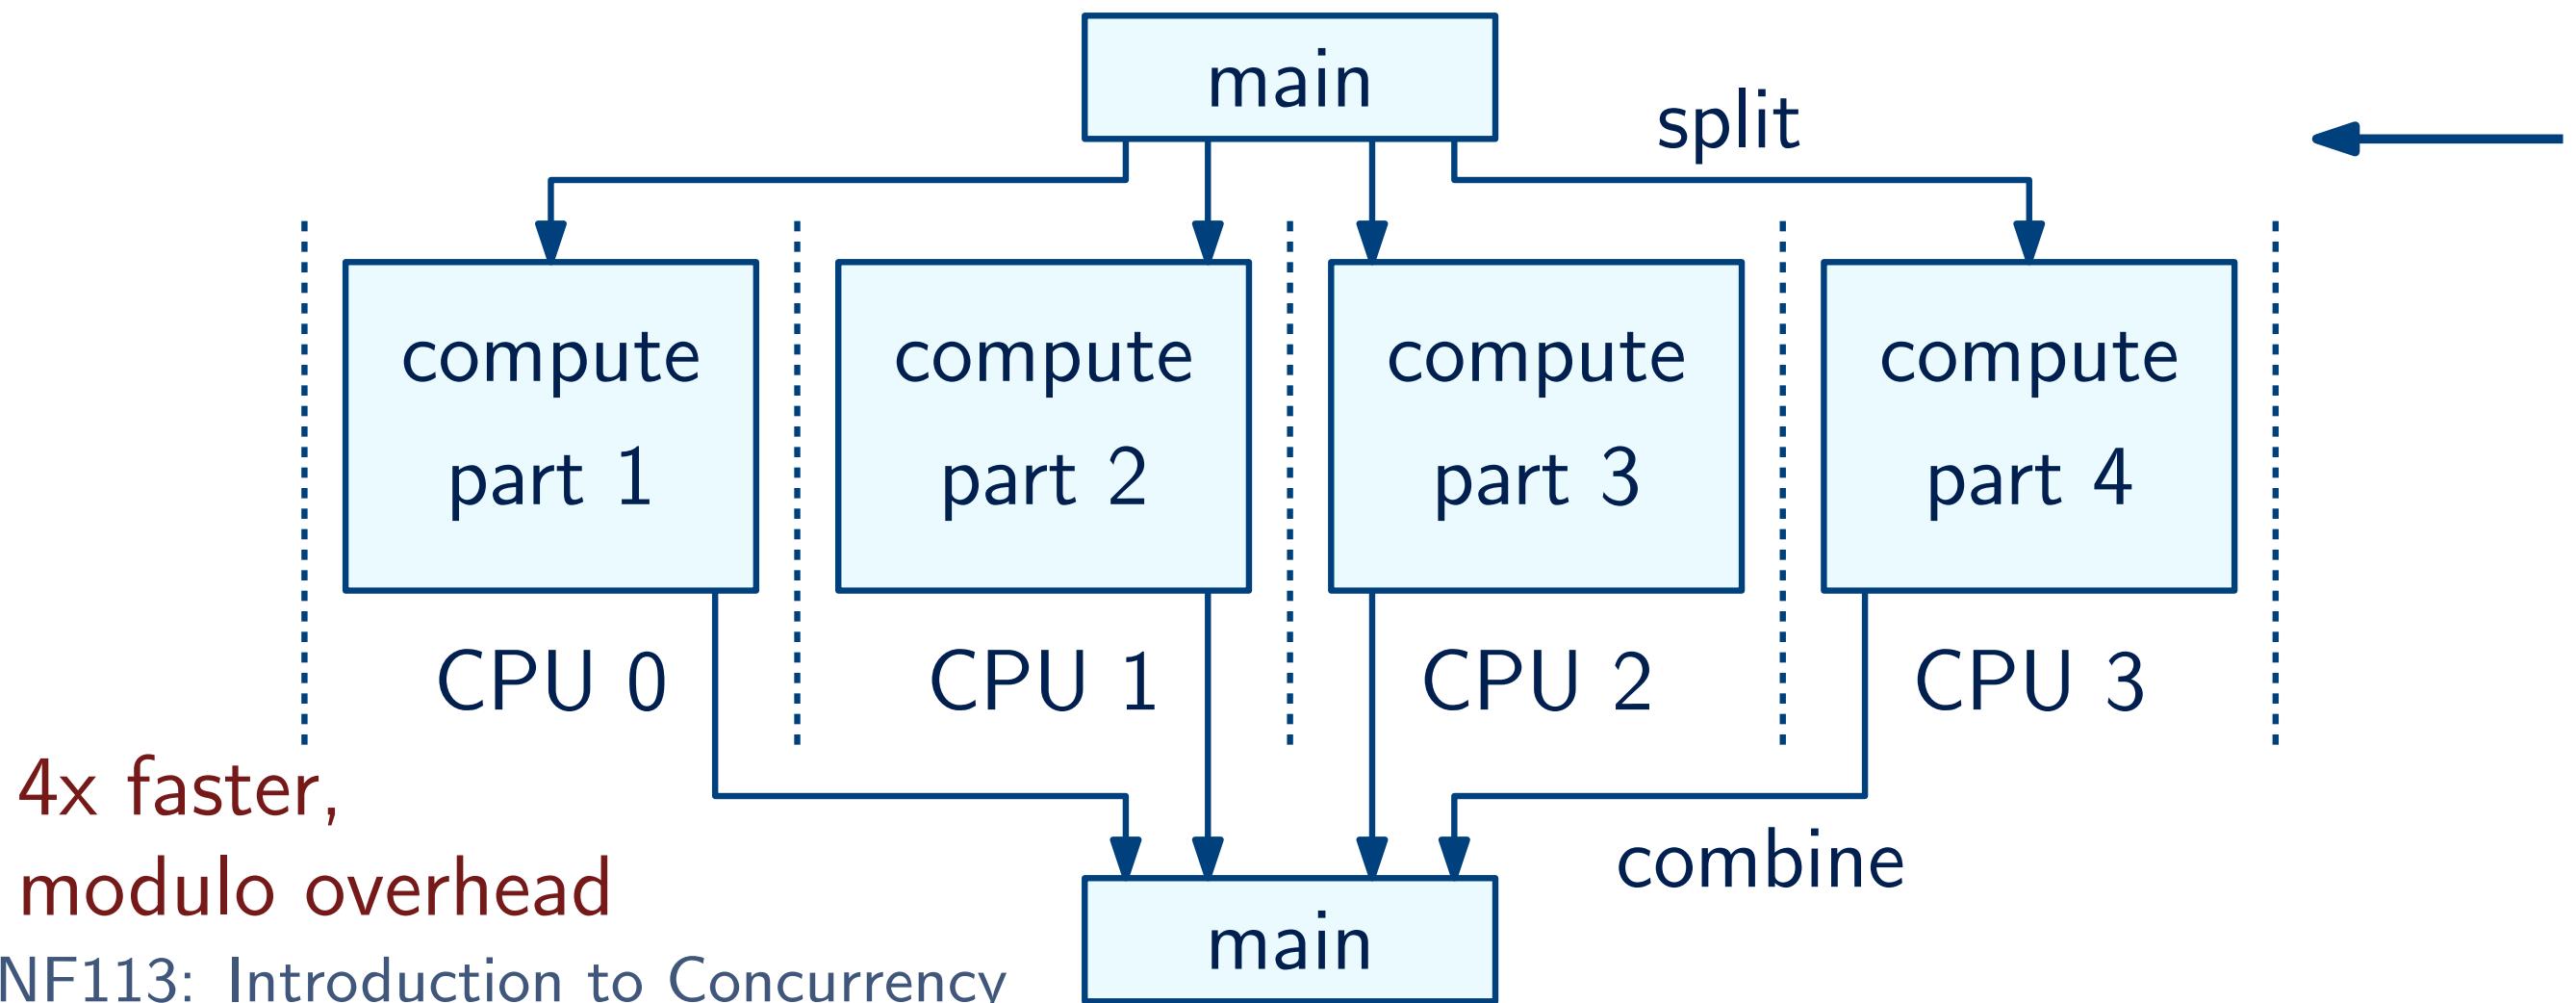
INF113: Introduction to Concurrency

Kirill Simonov
15.10.2025



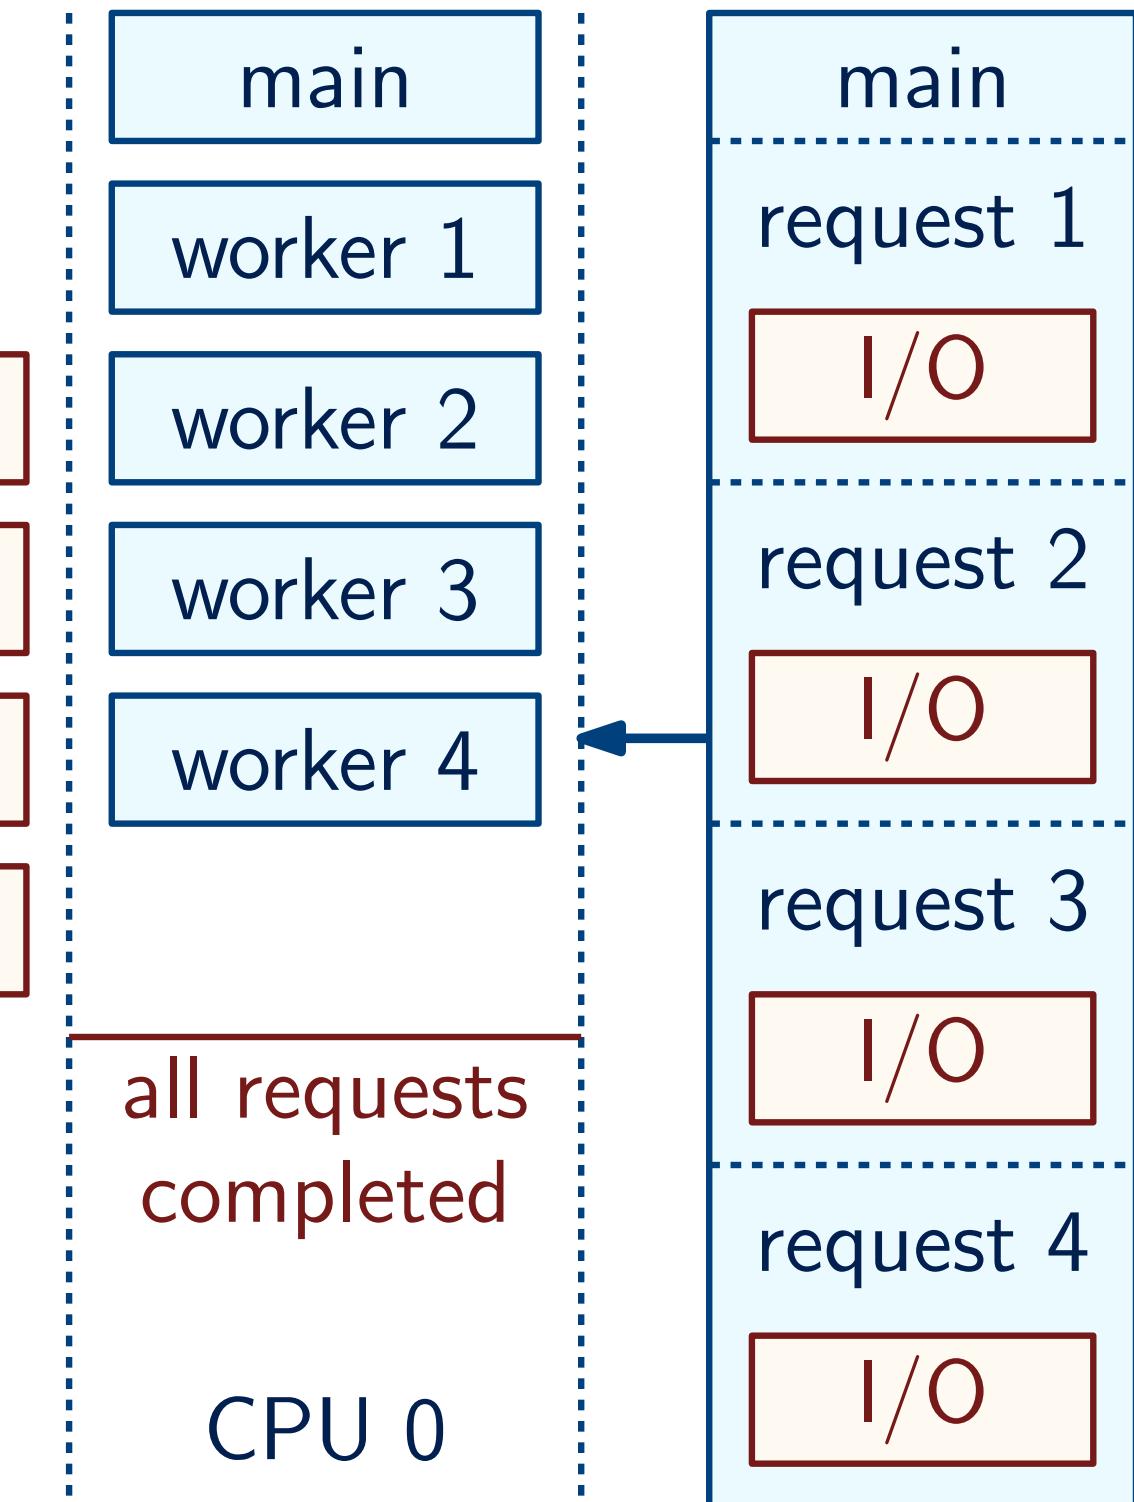
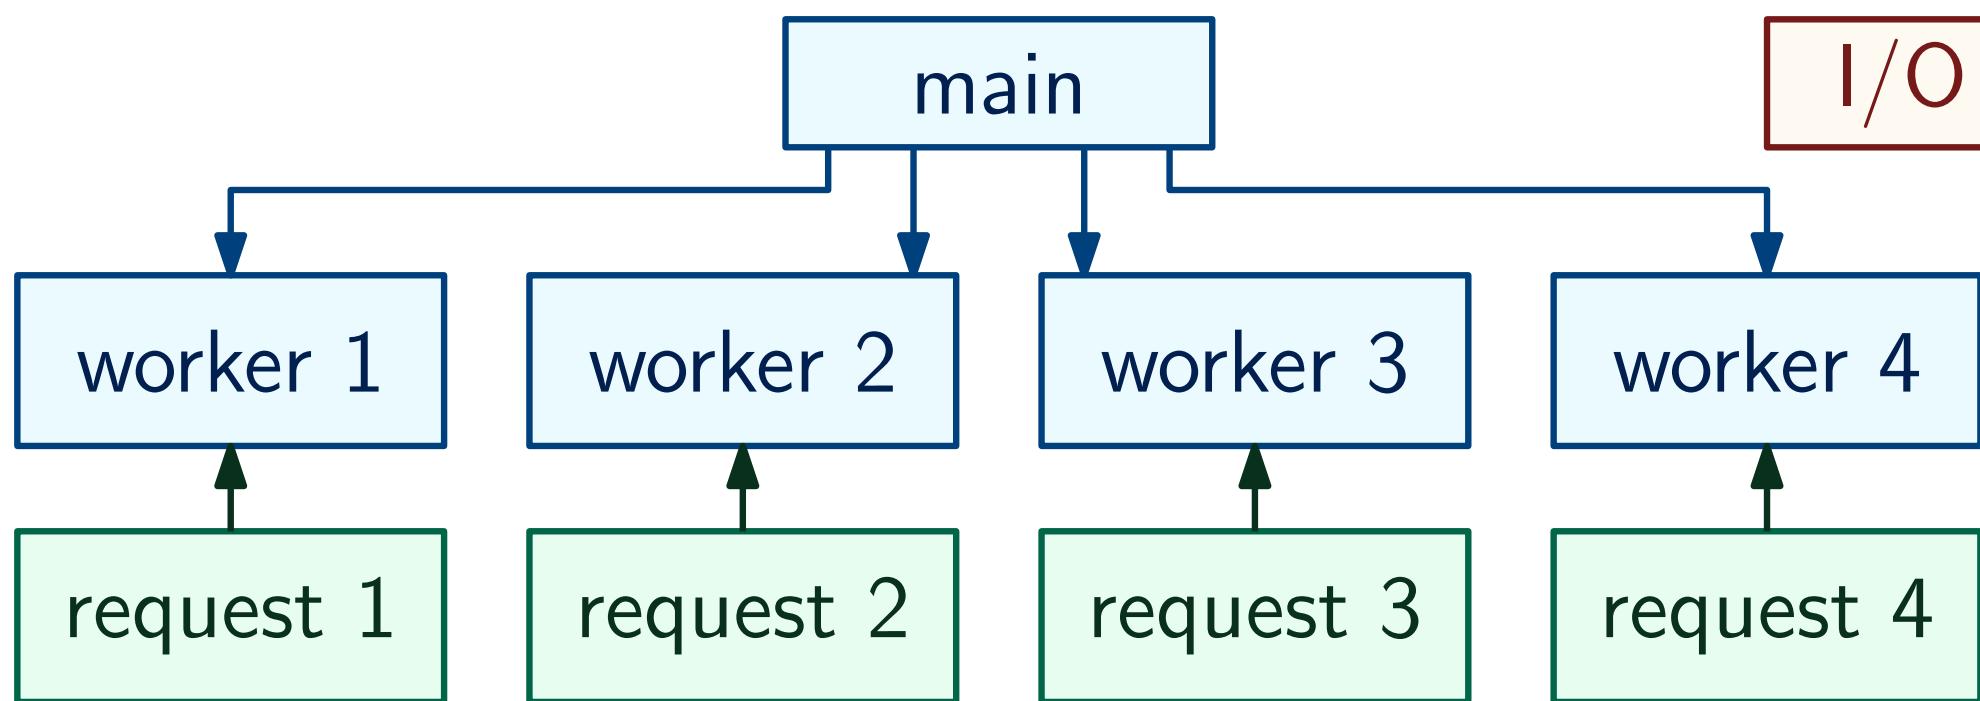
Parallel computation

- A program may split computation into smaller pieces to run in parallel
- Each “subprocess” gets a piece
- Examples: encoding a video, reverting a hash, solving TSP...



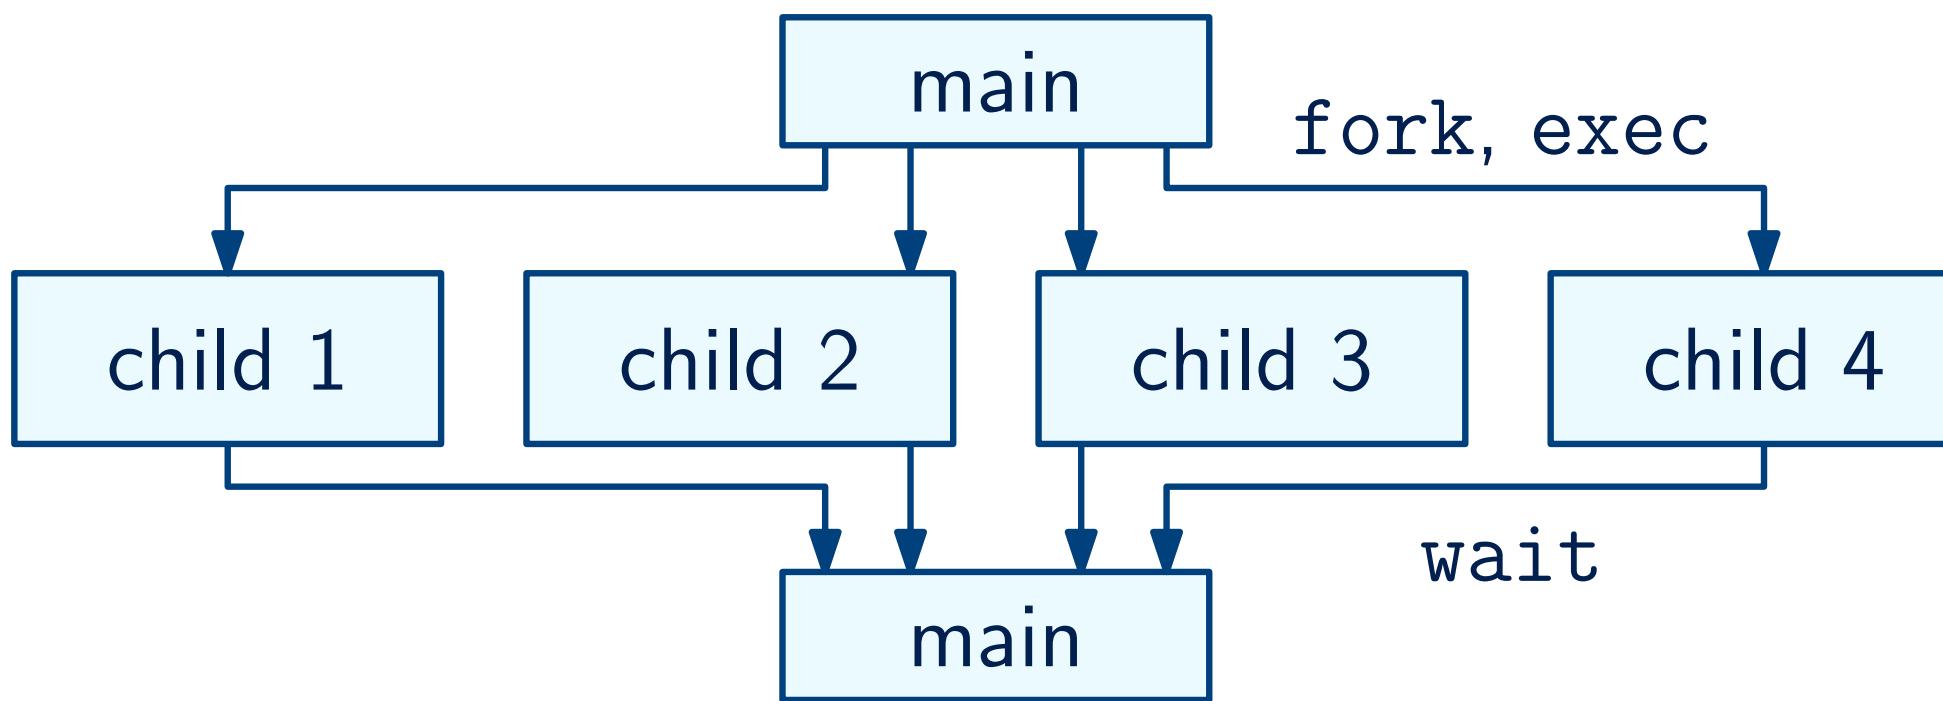
Servicing requests

- A program split into “subprocesses” may progress while waiting on I/O
- Saves time even on a single CPU core
- Example: database



Child processes

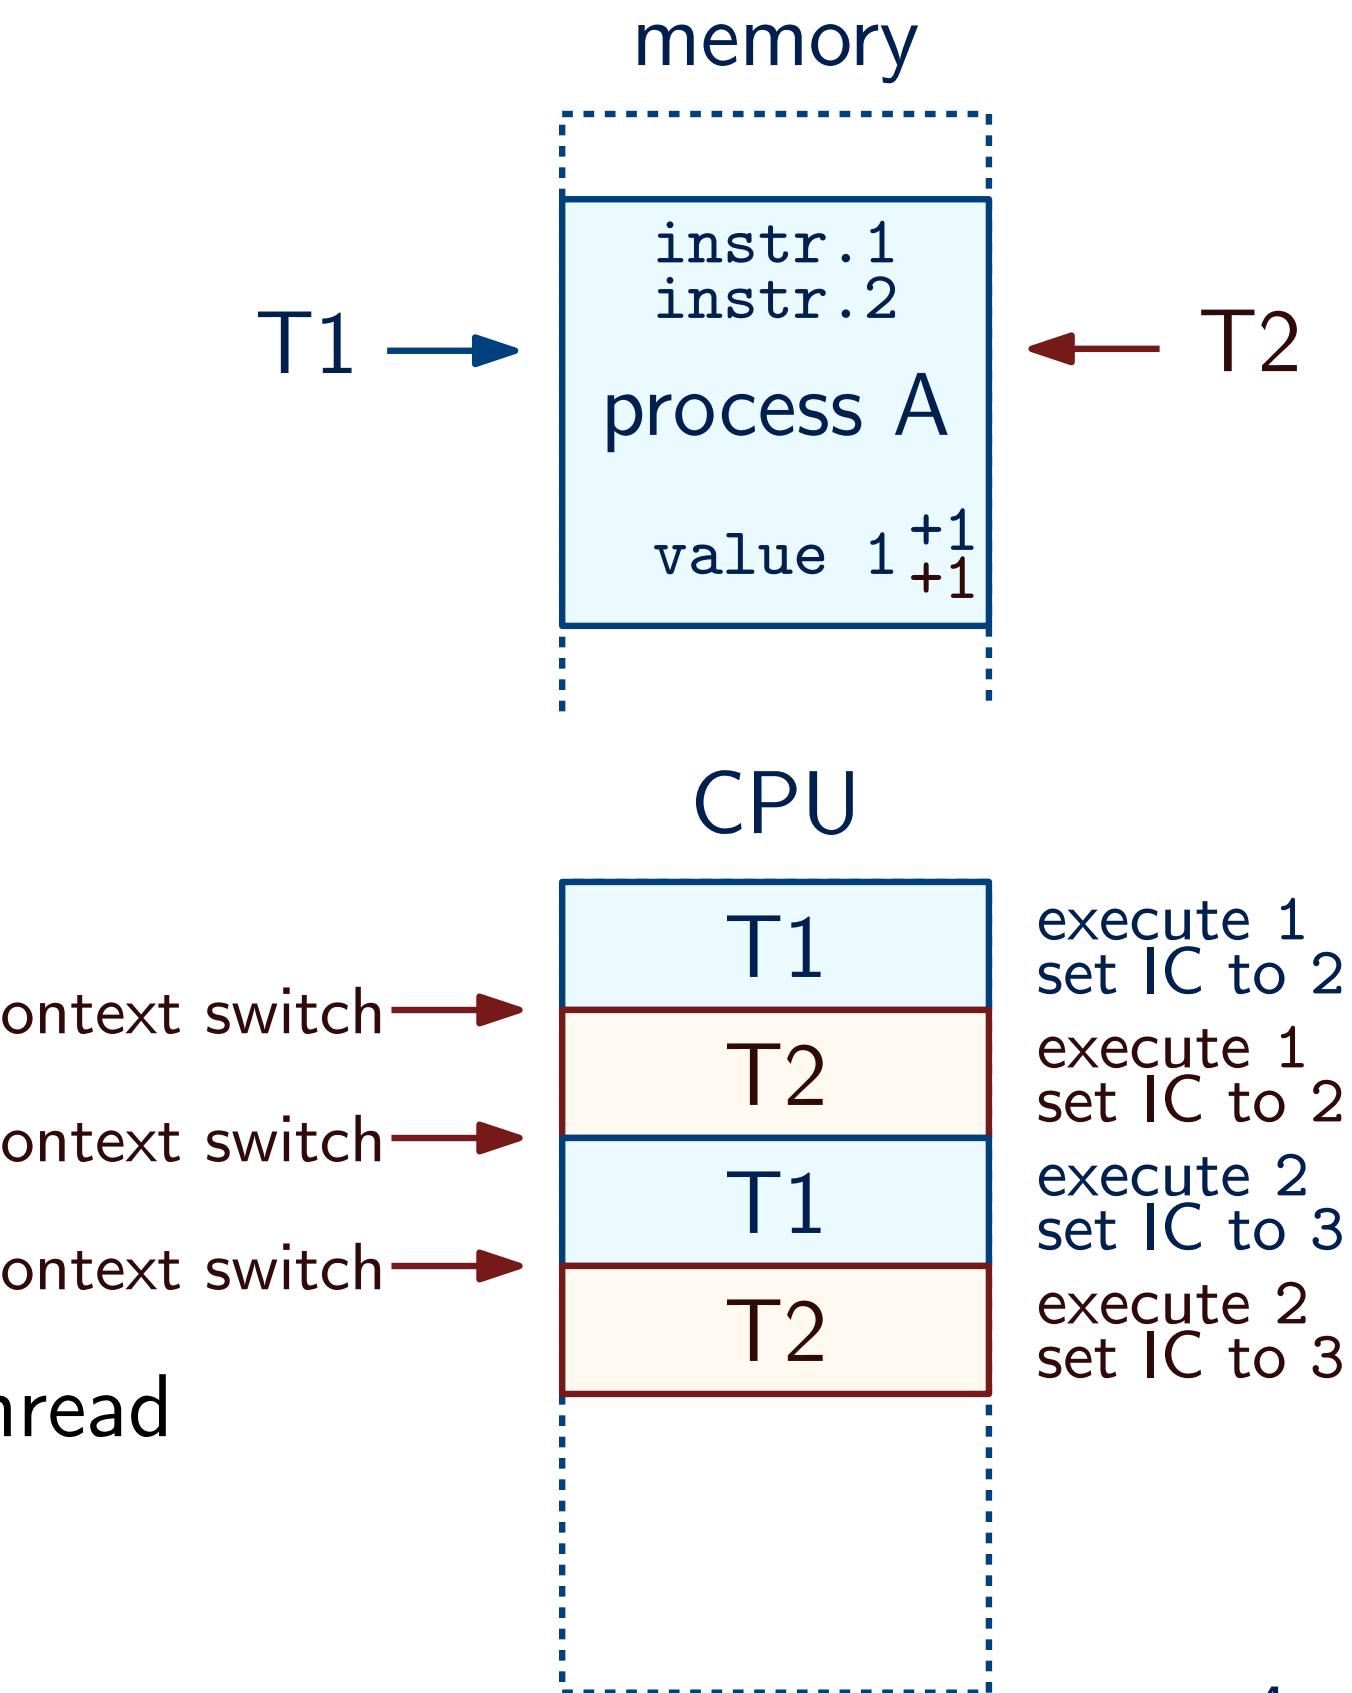
- We could implement parallelization by spawning child processes
 - “Prime” each child with any data
 - Run independently, also in parallel
 - Blocked independently, e.g., with I/O
 - Wait for completion if needed
- Issue: almost no interaction between the processes
 - Possible via files, but very slow



```
for (int i = 0; i < 4; ++i) {  
    int rc = fork();  
    if (rc == 0) {  
        //preparation  
        //code  
        exec();  
    }  
}  
  
while (wait() > 0) {};  
//parent  
//continues
```

Threads

- “Subprocesses” need to have shared memory
- **Thread** is a point of execution within a process
 - Its own value of Instruction Counter
 - And other registers
 - But no own memory
- Each process may have multiple threads
- OS schedules threads independently
- On a context switch, OS saves the state of the thread
 - No need to flush TLB if still the same process



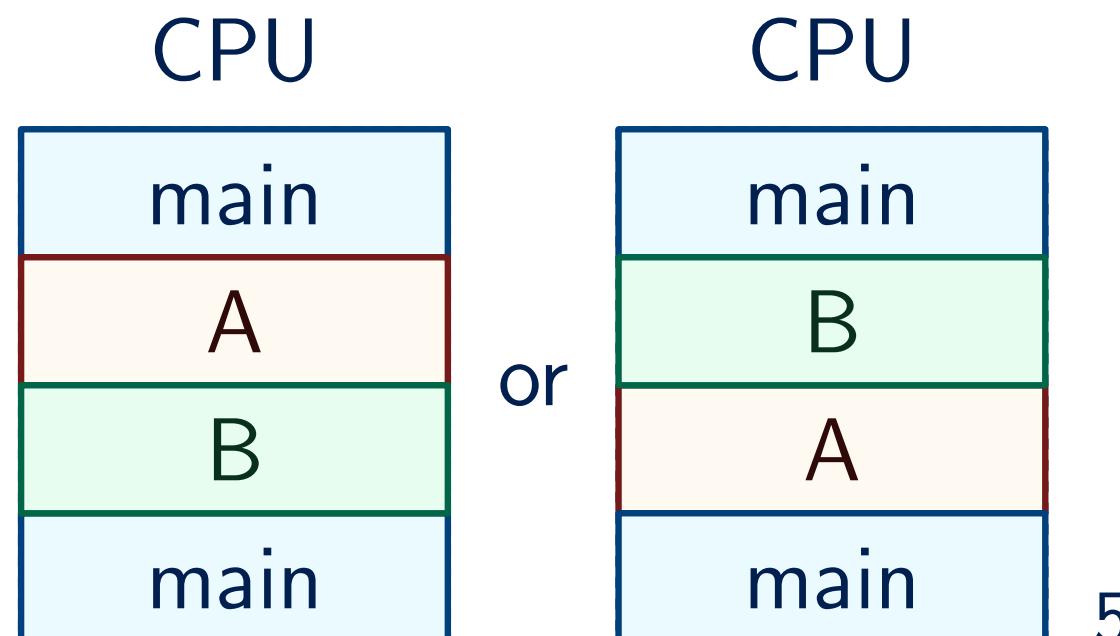
Example in C

```
#include <stdio.h>
#include <stdlib.h>
#include <pthread.h>

void *mythread(void *arg) {
    printf("%s\n", (char *) arg);
    return NULL;
}

int main() {
    pthread_t p1, p2;
    printf("main: begin\n");
    pthread_create(&p1, NULL, mythread, "A");
    pthread_create(&p2, NULL, mythread, "B");
    pthread_join(p1, NULL);
    pthread_join(p2, NULL);
    printf("main: end\n");
    return 0;
}
```

- `pthread_create` starts a new thread with a given entry point
- `pthread_join` waits for a given thread
- `void*` allows to pass/return any data
- order is not fixed:



C tricks: Wrapping arguments

```
...  
typedef struct {  
    int arg1, arg2;  
} args_t;  
  
void *mythread(void *arg) {  
    args_t *a = (args_t *) arg;  
    printf("%d %d\n", a->arg1, a->arg2);  
    return NULL;  
}  
  
int main() {  
    ...  
    args_t a = {3, 2};  
    pthread_create(&p, NULL, mythread, &a);  
    ...  
}
```

- Any collection of arguments can be arranged in a `struct` and passed as a `void*` pointer

C tricks: Wrapping arguments

```
void *mythread(void *arg) {  
    args_t *a = (args_t *) arg;  
    args_t *ret = malloc(sizeof(args_t));  
    ret->arg1 = a->arg1 + a->arg2;  
    ret->arg2 = a->arg1 - a->arg2;  
    return ret;  
}  
  
int main() {  
    ...  
    void *ret;  
    pthread_join(p, &ret);  
    args_t *aret = ret;  
    printf("main: %d %d\n",  
        aret->arg1, aret->arg2);  
    ...  
}
```

- Any collection of arguments can be arranged in a struct and passed as a `void*` pointer
- Same for return values

C tricks: Wrapping arguments

```
...  
void *mythread(void *arg) {  
    int a = *(int *)arg;  
    int b = a + 1; pointer to b will not be valid!  
    return &b;  
}  
  
int main() {  
    pthread_t p;  
    printf("main: begin\n");  
    int a = 3;  
    pthread_create(&p, NULL, mythread, &a);  
    void *ret;  
    pthread_join(p, &ret);  
    printf("main: %d\n", *(int *)ret);  
    printf("main: end\n");  
    return 0;  
}
```

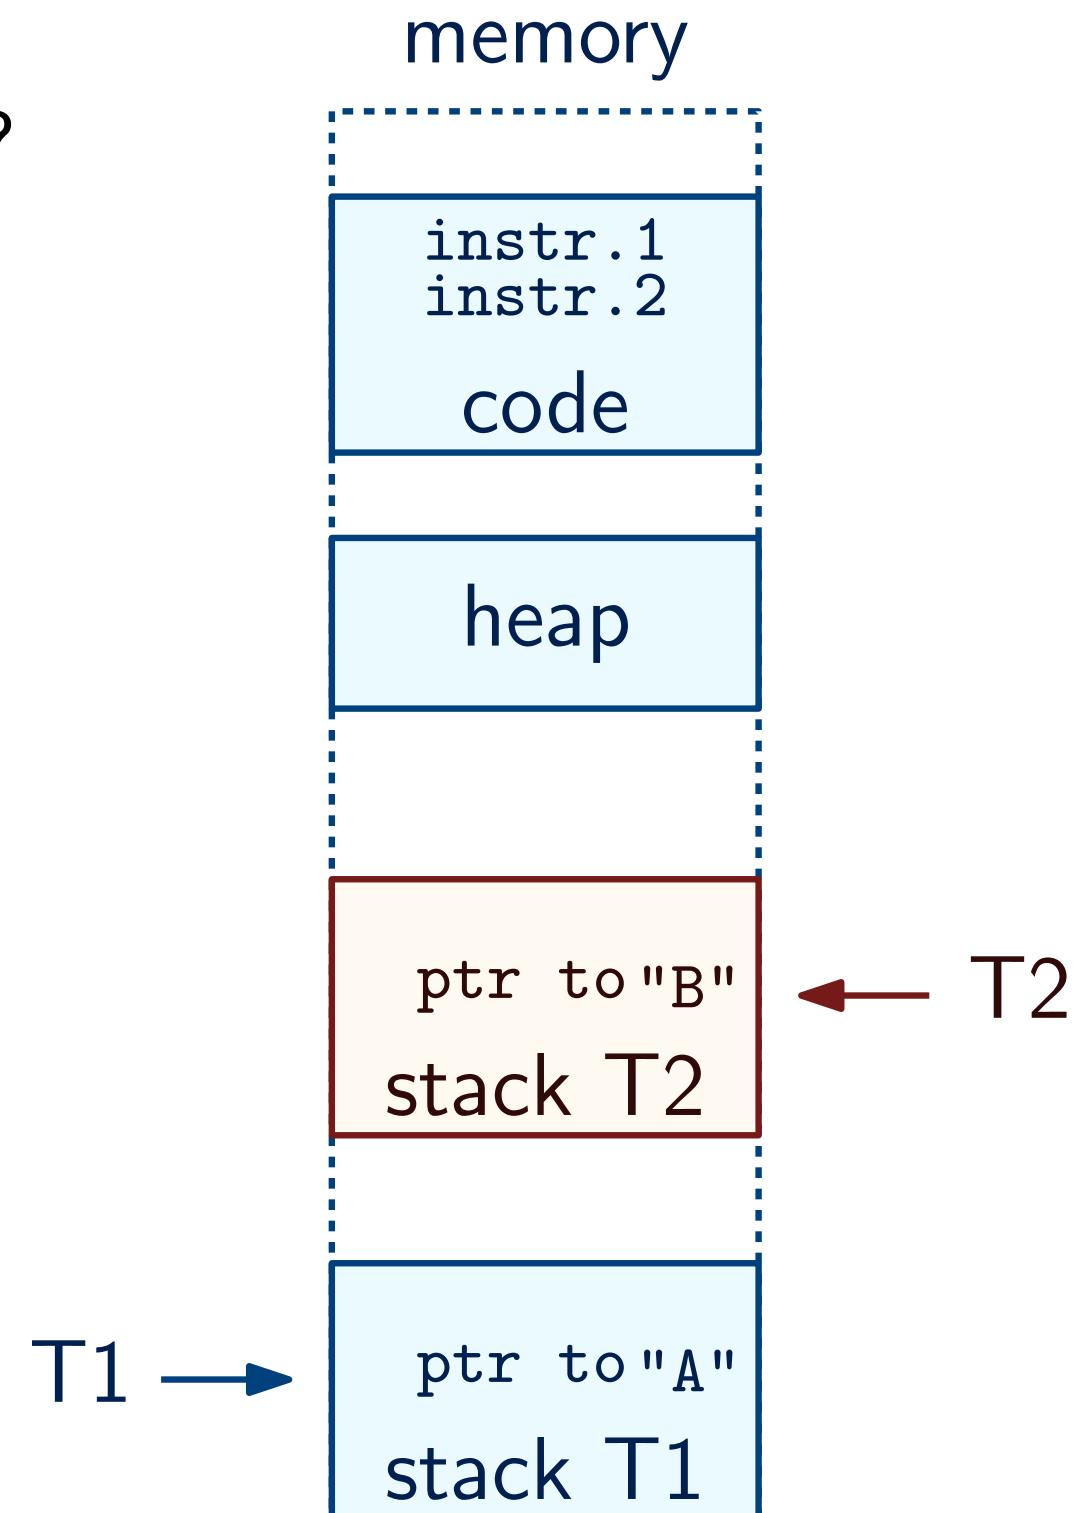
- Any collection of arguments can be arranged in a struct and passed as a void* pointer
- Same for return values
- **Beware:** Variables allocated on stack get invalidated!

Threads and stack

- If threads possess no memory, how the args are different?

```
...
void *mythread(void *arg) {
    printf("%s\n", (char *) arg);
    return NULL;
}
...
pthread_create(&p1, NULL, mythread, "A");
pthread_create(&p2, NULL, mythread, "B");
...
```

- Actually, each thread gets its own place for stack
 - Hence, local variables are “private” for the thread
 - But not protected, since threads share the same address space

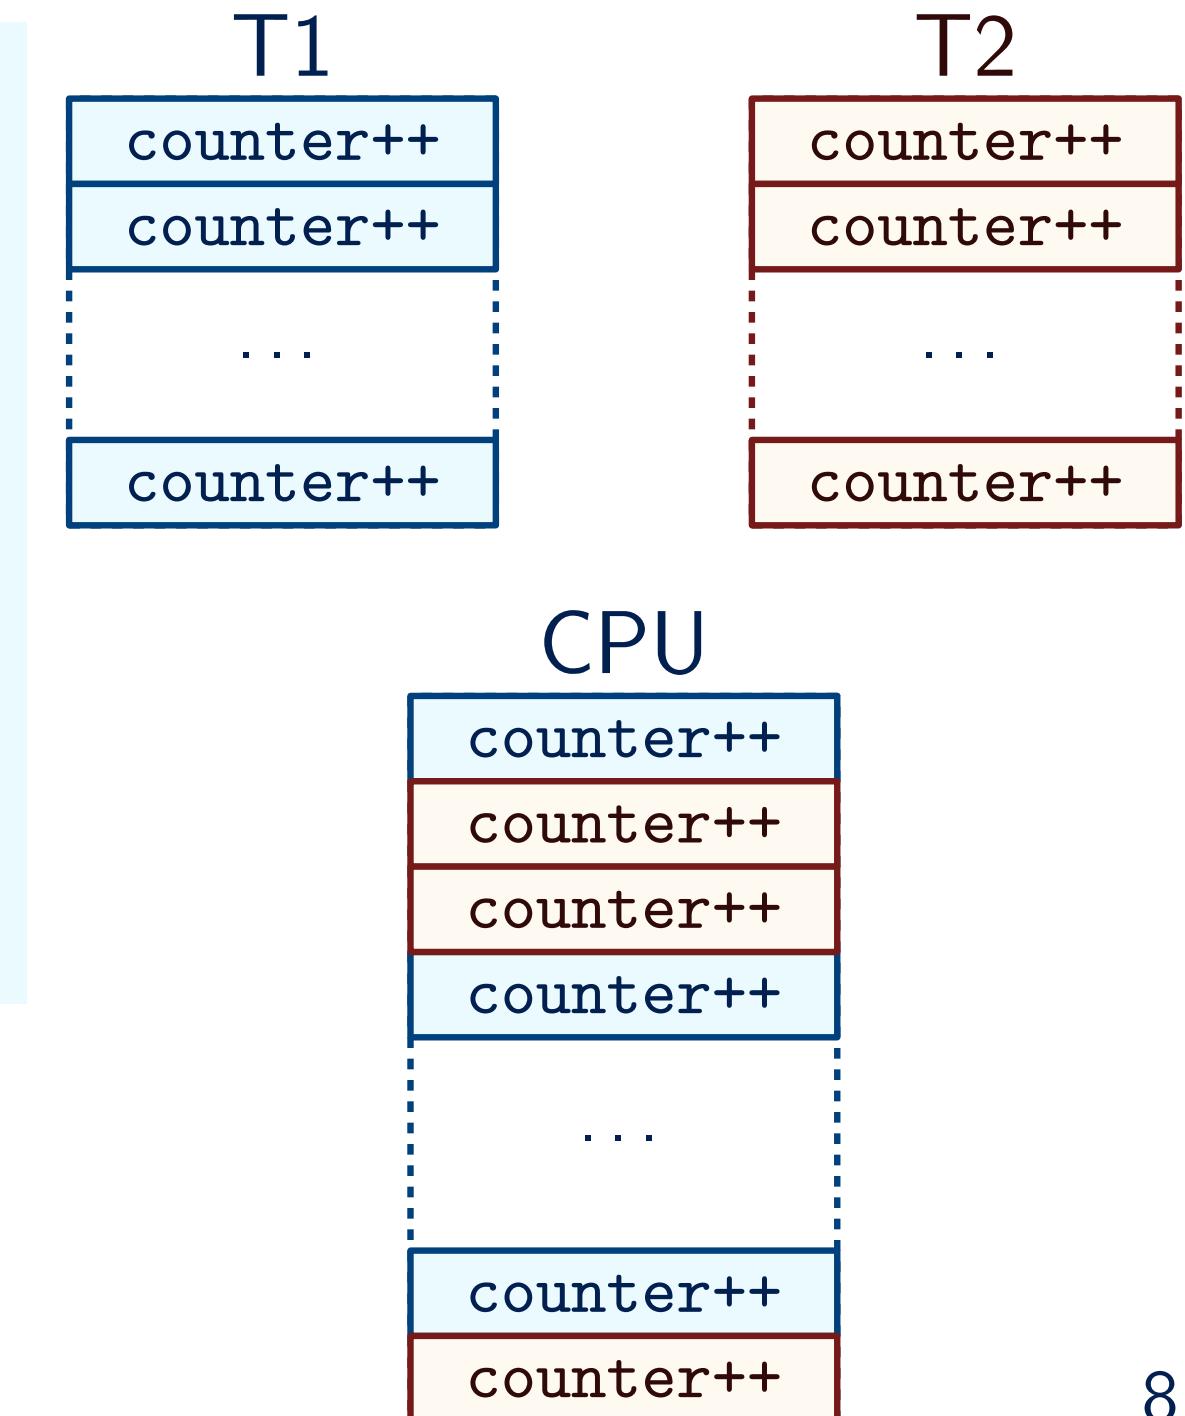


Example 2

- If two threads modify the same counter, the result is very unpredictable—why?

```
...
volatile int counter = 0;

void *mythread(void *arg) {
    printf("%s: begin\n", (char *) arg);
    for (int i = 0; i < 1e7; i++) {
        counter = counter + 1;
    }
    printf("%s: end\n", (char *) arg);
    return NULL;
}
...
```



- Seemingly, the order in which the threads increment the counter should not matter

Example 2: Explanation

- counter = counter + 1 is not **atomic**!
- Context switch may happen **within** the operation

```
counter = counter + 1;  
mov 0x8049a1c, %eax  
add $0x1, %eax  
mov %eax, 0x8049a1c
```

T1 → T2 ←

CPU
T1: mov 0x8049a1c, %eax
T2: mov 0x8049a1c, %eax
T2: add \$0x1, %eax
T2: mov %eax, 0x8049a1c
T2: mov 0x8049a1c, %eax
T2: add \$0x1, %eax
T2: mov %eax, 0x8049a1c
T1: add \$0x1, %eax
T1: mov %eax, 0x8049a1c
T2: mov 0x8049a1c, %eax
T2: add \$0x1, %eax
T2: mov %eax, 0x8049a1c

counter	%eax in T1	%eax in T2
0	0	
0	0	0
0	0	1
1	0	1
1	0	1
1	0	2
2	0	2
2	1	2
1	1	2
1	1	1
1	1	2
2	1	2

after 1 iteration of T1
and 3 iterations of T2,
counter is at 2!

Some terminology

- **Critical section** is a part of code accessing a **shared resource**
- **Race condition** occurs when multiple threads enter a critical section simultaneously
- A program that admits race conditions is **indeterminate**
- To avoid this, there must be **mutual exclusion**: if one thread enters the critical section, the others cannot

```
...
volatile int counter = 0;
    shared resource
void *mythread(void *arg) {
    printf("%s: begin\n", (char *) arg);
    for (int i = 0; i < 1e7; i++) {
        counter = counter + 1;
    } critical section
    printf("%s: end\n", (char *) arg);
    return NULL;
}
...
```



The diagram illustrates a shared resource, `counter`, which is modified within a critical section. The `counter` variable is highlighted with a red box and labeled "shared resource". Inside the `mythread` function, there is a loop where the `counter` is incremented by 1. This increment operation is also highlighted with a red box and labeled "critical section". Two arrows point from the text "shared resource" and "critical section" to their respective highlighted code segments.

Atomic instructions

- We could hope that hardware implements common modifications **atomically**
- I.e., there is a single instruction instead of several
- While some exist, it is not sufficient in general
- Example: Concurrently modify a BST
 - Hardware should support atomics on our specific node structure?

```
T1 → int search(node *v, int key) {  
    if (v->key < key) {  
        search(v->left, key);  
    }  
    ...  
}
```

```
T2 → int remove(node *v, int key) {  
    if (v->key == key) {  
        tree* sub = merge(v->left, v->right);  
        free(v);  
        return sub;  
    }  
    ...  
}
```

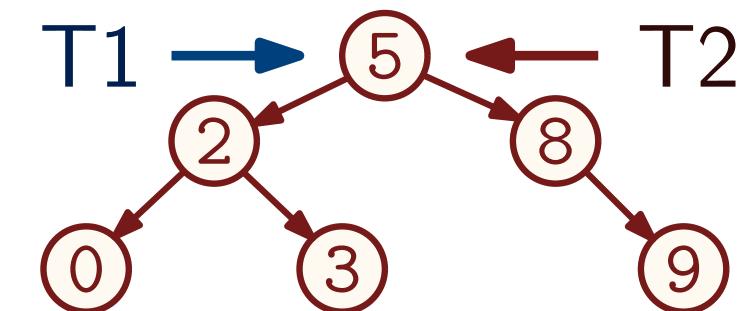
```
counter = counter + 1;
```

```
mov 0x8049a1c, %eax
```

```
add $0x1, %eax
```

```
mov %eax, 0x8049a1c
```

```
mem-add 0x8049a1c, $0x1
```



Summary

- We aim to design a set of general **synchronization primitives** based on limited hardware instructions
- Allow to safely “block” parts of code when in use by a thread
- Allow to “wake” threads when the wait condition is fulfilled

```
void *mythread(void *arg) {  
    printf("%s: begin\n", (char *) arg);  
    ↓  
    for (int i = 0; i < 1e7; i++) {  
        counter = counter + 1;  
    }  
    ↓  
    printf("%s: end\n", (char *) arg);  
    return NULL;  
}
```

Summary

- We aim to design a set of general **synchronization primitives** based on limited hardware instructions
- Allow to safely “block” parts of code when in use by a thread
- Allow to “wake” threads when the wait condition is fulfilled
- OS viewpoint: what support from hardware and the OS is needed to make concurrency in programs easy?
- OS viewpoint: OS is the original concurrent process
 - Write to disk, start/finish process, . . . , all in presence of interrupts