

INF113: Locks

Kirill Simonov
17.10.2025



Reminder

- We aim to design a set of general **synchronization primitives** based on limited hardware instructions
- Allow to safely “block” parts of code when in use by a thread
- In other words, when entering a critical section T1 “locks” it from T2
- When T1 is done, the section should be unlocked again



```
void *mythread(void *arg) {
    printf("%s: begin\n", (char *) arg);
    for (int i = 0; i < 1e7; i++) {
        counter = counter + 1;
    }
    printf("%s: end\n", (char *) arg);
    return NULL;
}
```

Locks: the interface

- A datatype that stores the locked/unlocked state + possibly additional data

```
lock_t mutex;
```

- A function that sets the lock into “locked”, called **before** the critical section

```
lock(&mutex);
```

- A function that sets the lock into “unlocked”, called **after** the critical section

```
unlock(&mutex);
```

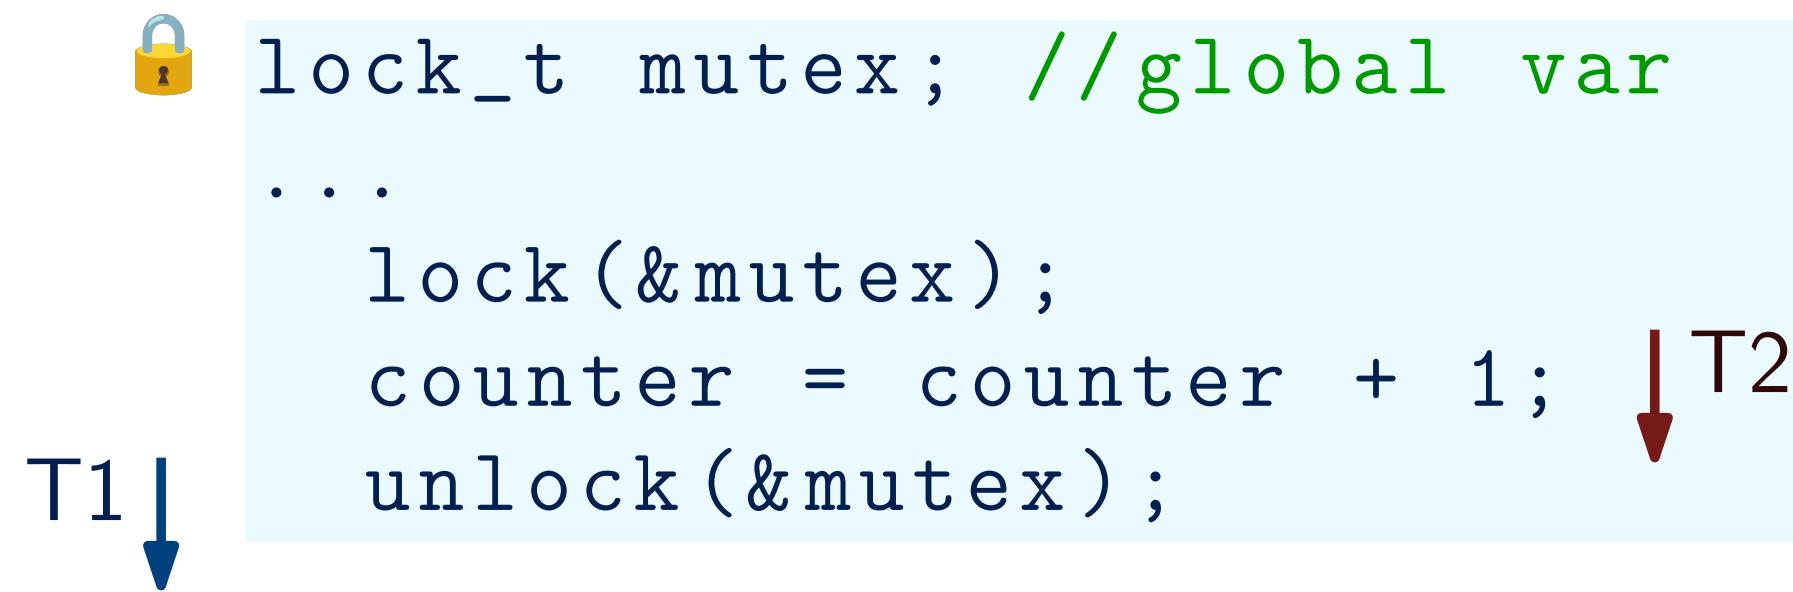


altogether:

```
lock_t mutex; // global var  
...  
lock(&mutex);  
counter = counter + 1;  
unlock(&mutex);
```

Lock function behavior

- mutex is initially “unlocked”
- T1 calls lock: sets mutex to “locked” and returns
- T2 calls lock: while mutex is “locked”, T2 is stuck within lock
- T1 calls unlock: sets mutex to “unlocked”
- T2 sets mutex to “locked” within lock and exits lock



```
lock_t mutex; // global var  
...  
lock(&mutex);  
counter = counter + 1; ↓ T2  
unlock(&mutex);
```

POSIX locks

- `pthread.h` provides implementation for a lock, called a **mutex**

```
#include <pthread.h>
pthread_mutex_t lock = PTHREAD_MUTEX_INITIALIZER;
...
pthread_mutex_lock(&lock);
counter = counter + 1;
pthread_mutex_unlock(&lock);
```

- A mutex can be dynamically initialized and destroyed

```
pthread_mutex_init(&lock, NULL);
...
pthread_mutex_destroy(&lock);
```

Locking tips

- Locking scope trade-offs
 - Larger scope: less flexibility, more waiting
 - Smaller scope: less waiting, more context switching

```
pthread_mutex_lock(&lock);
for (int i = 0; i < 1e7; ++i) {
    counter = counter + 1;
}
pthread_mutex_unlock(&lock);
```

```
for (int i = 0; i < 1e7; ++i) {
    pthread_mutex_lock(&lock);
    counter = counter + 1;
    pthread_mutex_unlock(&lock);
}
```

Locking tips

- Locking scope trade-offs
 - Larger scope: less flexibility, more waiting
 - Smaller scope: less waiting, more context switching
- Use different locks for different shared resources

```
T1 ↓
pthread_mutex_lock(&lock1);
counter1 = counter1 + 1;
pthread_mutex_unlock(&lock1);
...
pthread_mutex_lock(&lock2);
counter2 = counter2 + 1;
pthread_mutex_unlock(&lock2);
```

while T1 modifies
counter1, T2 is safe
to access counter2

↓ T2

Locking tips

- Locking scope trade-offs
 - Larger scope: less flexibility, more waiting
 - Smaller scope: less waiting, more context switching
- Use different locks for different shared resources
- Less time in critical sections, less locking

```
int local_cnt = 0;
for (int i = 0; i < 1e7; ++i) {
    local_cnt = local_cnt + 1;
    if (local_cnt > 1000) {
        pthread_mutex_lock(&lock);
        counter += local_cnt;
        pthread_mutex_unlock(&lock);
        local_cnt = 0;
    }
}
```

Example:
approximate/delayed
counter

Building locks

- How to build an efficient lock?
 - Hardware needs?
 - OS needs?

Evaluating the lock design:

- **Mutual exclusion:** Does it fulfill the purpose?
- **Fairness:** Does every thread get a chance of taking the lock?
Alternatively, could a thread **starve**, waiting for a lock forever?
- **Performance:** What is the extra overhead of locking/unlocking?

```
void lock() {  
    if (T1) {  
        return;  
    } else {  
        while (true) {};  
    }  
}
```

Example: only T1
gets through the lock

Naïve implementation

```
typedef struct __lock_t { int flag; } lock_t;

void init(lock_t *mutex) {
    // 0 -> lock is available, 1 -> held
    mutex->flag = 0;
}

void lock(lock_t *mutex) {
    while (mutex->flag == 1) // TEST the flag
        ; // spin-wait (do nothing)
    mutex->flag = 1; // now SET it!
}

void unlock(lock_t *mutex) {
    mutex->flag = 0;
}
```

- Let's just store the state in a 0/1 variable
- Locking sets the flag to 1
Unlocking sets to 0
- If the flag is taken, the thread will wait in a while-loop
- See the problem?

CPU

T1: TEST
T2: TEST
T1: SET
T2: SET

Another thread
can sneak
in-between!

Test-and-set instruction

- It seems there should be a single atomic instruction that performs test and set at the same time:

```
int test_and_set(int *old_ptr, int new) {  
    int old = *old_ptr; // fetch old value at old_ptr  
    *old_ptr = new; // store 'new' into old_ptr  
    return old; // return the old value  
}
```

but actually implemented in hardware

- Hardware support exists indeed: e.g., called xchg on x86
- We can now redesign the lock implementation with test_and_set in mind

Test-and-set lock

```
typedef struct __lock_t { int flag; } lock_t;  
  
void init(lock_t *mutex) {  
    // 0 -> lock is available, 1 -> held  
    mutex->flag = 0;  
}  
  
void lock(lock_t *mutex) {  
    while (test_and_set(&mutex->flag, 1) == 1)  
        ; // spin-wait (do nothing)  
}  
  
void unlock(lock_t *mutex) {  
    mutex->flag = 0;  
}
```

- If flag is 0, instruction sets it to 1 and reports 0 —thread proceeds
- If flag is 1, instruction sets it to 1 and reports 1 —thread waits

```
int test_and_set(  
    int *old_ptr,  
    int new)  
{  
    int old = *old_ptr;  
    *old_ptr = new;  
    return old;  
}
```

Compare-and-swap

- Similar to test-and-set, but also compares old value with an expected value

```
int compare_and_swap(int *old_ptr, int expected, int new) {  
    int old = *old_ptr; // fetch old value at old_ptr  
    if (old == expected) // update only if value matches  
        *old_ptr = new; // store 'new' into old_ptr  
    return old; // return the old value  
}
```

- Can be used in the lock

```
void lock(lock_t *mutex) {  
    while (compare_and_swap(&mutex->flag, 0, 1) == 1)  
        ; // spin-wait (do nothing)  
}
```

Compare-and-swap

- Similar to test-and-set, but also compares old value with an expected value

```
int compare_and_swap(int *old_ptr, int expected, int new) {  
    int old = *old_ptr; // fetch old value at old_ptr  
    if (old == expected) // update only if value matches  
        *old_ptr = new; // store 'new' into old_ptr  
    return old; // return the old value  
}
```

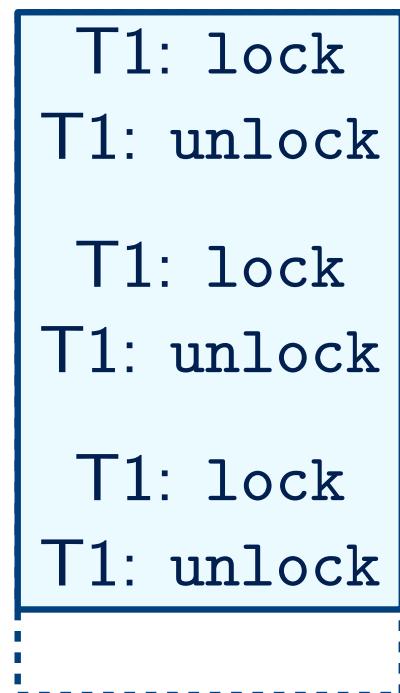
- Can be used in the lock
- Also to implement other operations atomically, for example increment:

```
void inc(int *ptr) {  
    while (true) {  
        int old = *ptr;  
        if (compare_and_swap(ptr, old, old + 1) == old) {  
            break;  
        }  
    }  
}
```

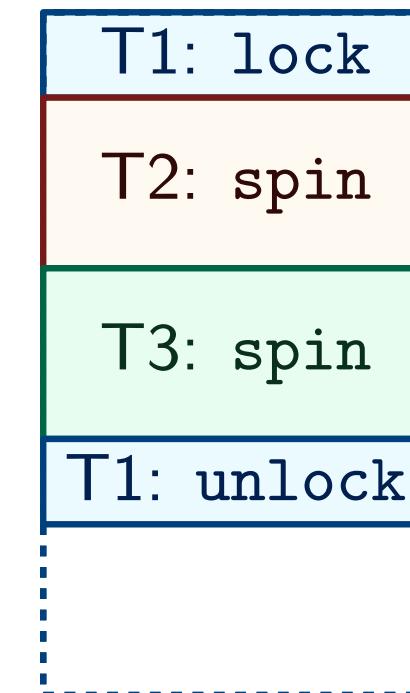
Evaluating spin locks

- **Mutual exclusion:** Done! ✓
- **Fairness:** A thread can be spinning indefinitely, if it's unlucky ?
- **Performance:** Depends on scenario, but spinning could be costly ?

Single thread



Single core



nobody else takes the lock so no waiting



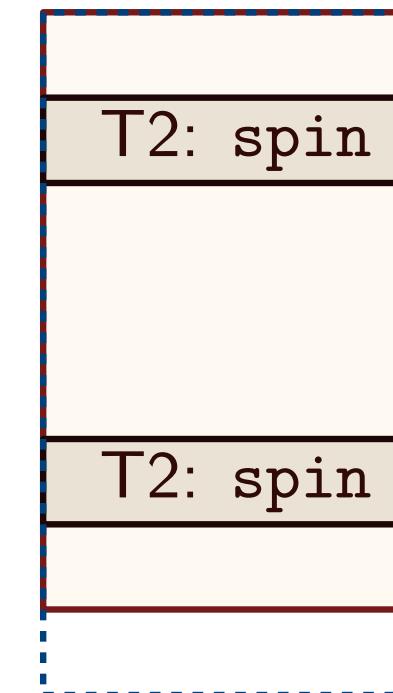
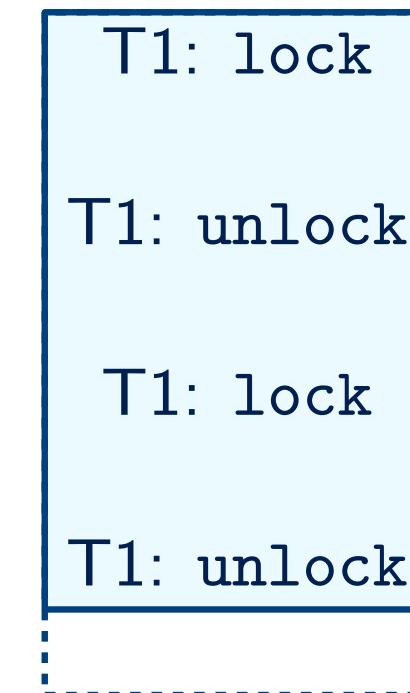
CPU

each thread spins for a time slice before T1 unlocks



CPU

Threads = cores



only spins while the other is in a critical section



CPU 0

CPU 1

Fetch-and-add

```
typedef struct __lock_t {  
    int ticket, turn;  
} lock_t;  
  
void init(lock_t *mutex) {  
    lock->ticket = 0;  
    lock->turn = 0;  
}  
  
void lock(lock_t *mutex) {  
    int my_turn = fetch_and_add(&lock->ticket);  
    while (lock->turn != myturn)  
        ; // spin  
}  
  
void unlock(lock_t *mutex) {  
    lock->turn = lock->turn + 1;  
}
```

- Atomically increment and return the value of the counter:

```
int fetch_and_add(int  
                  *ptr) {  
    int old = *ptr;  
    *ptr = old + 1;  
    return old;  
}
```

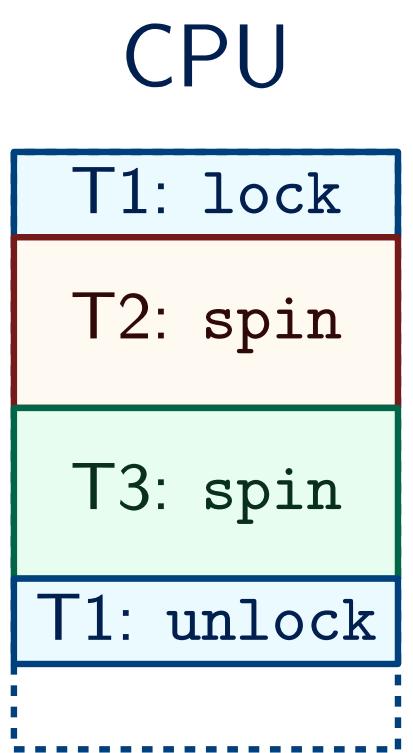
- Ticket-based lock: once a thread gets a ticket, it is guaranteed to run!

Fairness: ✓

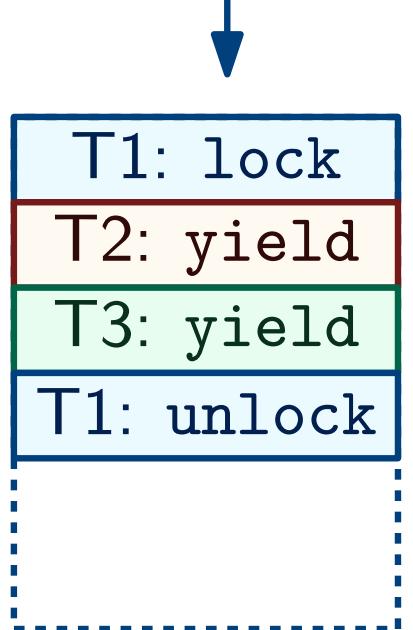
Resolving spins

- If a thread without the lock is run, it wastes CPU for its whole timeslice
- Attempt 1: Let the thread voluntarily give up control instead of spinning

```
void lock(lock_t *mutex) {  
    while (test_and_set(&mutex->flag, 1) == 1)  
        yield(); // pass control back  
}
```



- Still, lots of context switches before we get to unlock
- Threads can still starve



Queue-based lock

- We will put waiting threads to sleep in a queue
- OS support, Solaris example:
 - Syscall park(); to yield control and be put to sleep,
 - Syscall unpark(thread); to wake thread and pass control
- flag keeps track of the lock
- guard is an internal lock to access the queue
- queue stores the waiting threads

```
typedef struct __lock_t {  
    int flag;  
    int guard;  
    queue_t *q;  
} lock_t;  
  
void lock_init(lock_t *m) {  
    m->flag = 0;  
    m->guard = 0;  
    queue_init(m->q);  
}
```

Queue-based lock: lock

```
void lock(lock_t *m) {
    while (test_and_set(&m->guard, 1) == 1)
        ; // acquire guard lock by spinning
    if (m->flag == 0) {
        m->flag = 1; // lock is acquired
        m->guard = 0;
    } else {
        queue_add(m->q, gettid());
        m->guard = 0;
        park();
    }
}
```

- Then either get the flag and pass the main lock, or put yourself into the queue and yield

- First, try to pass the guard lock by spinning

```
typedef struct __lock_t {
    int flag;
    int guard;
    queue_t *q;
} lock_t;

void lock_init(lock_t *m) {
    m->flag = 0;
    m->guard = 0;
    queue_init(m->q);
}
```

Queue-based lock: unlock

```
void unlock(lock_t *m) {  
    while (test_and_set(&m->guard, 1) == 1)  
        ; // acquire guard lock by spinning  
    if (queue_empty(m->q))  
        m->flag = 0; // let go of lock; no  
        one wants it  
    else  
        unpark(queue_remove(m->q)); // hold  
        lock (for next thread!)  
    m->guard = 0;  
}
```

- Then either clear the flag, or put the next thread in the queue to work

- First same, try to pass the guard lock by spinning

```
typedef struct __lock_t {  
    int flag;  
    int guard;  
    queue_t *q;  
} lock_t;  
  
void lock_init(lock_t *m) {  
    m->flag = 0;  
    m->guard = 0;  
    queue_init(m->q);  
}
```

Queue-based lock: benefits

- Threads only spin to access the internal queue
 - Only the small and controlled queue operation is a critical section for that, not arbitrary user code
- **Fairness:** Once a thread is put in the queue, it will clear the lock eventually 
- **Performance:** Less guessing by the scheduler
 - When a thread is blocked, it yields and never randomly runs again
 - When a thread is done, it wakes the next waiting one

Summary

- A simple atomic instruction like test-and-set can be used to build safe locks
- Performance of a locking mechanism may depend a lot on the scenario: no. of threads, cores, time spent in critical sections, . . .
- To avoid wasting too much time on spinning, a good lock design introduces some order into thread scheduling—e.g., with the queue
 - **futex**, used in Linux, follows a similar queue-based approach
- Take a look at Chapter 28 for more details and a simulator
- **Next week:** more sync primitives, condition variables and semaphores