

# INF113: Semaphores

Kirill Simonov  
24.10.2025



# Locks and conditional variables

- Lock a.k.a mutex: “blocks” a part of code when in use by a thread
- Conditional variable: can either wait for a specific condition or signal the condition
  - Always used together with a lock
- **Next:** A simpler-to-use primitive that can model both\*

```
pthread_mutex_t mutex;

void *mythread(void *arg) {
    ...
    pthread_mutex_lock(&mutex);
    counter = counter + 1;
    pthread_mutex_unlock(&mutex);
    ...
}
```



```
void thr_exit() {
    pthread_mutex_lock(&m);
    done = 1;
    pthread_cond_signal(&c);
    pthread_mutex_unlock(&m);
}

void thr_join() {
    pthread_mutex_lock(&m);
    while (done == 0)
        pthread_cond_wait(&c, &m);
    pthread_mutex_unlock(&m);
}
```

# Semaphore: Interface

- Semaphore holds an integer **value**

```
#include <semaphore.h>
sem_t s;
```

- Has to be initialized, the value is then set to the argument


```
sem_init(&s, 0, value);
```

- sem\_wait decrements value and waits if negative

```
int sem_wait(sem_t *s) {
    decrement the value of semaphore s
    if value of semaphore s is negative, wait
}
```

- sem\_post increments value and sends a signal


```
int sem_post(sem_t *s) {
    increment the value of semaphore s
    if there are threads waiting, wake one
}
```



assume these are  
done atomically

# Semaphore as a lock

- Simply wrap the critical section in `sem_wait` and `sem_post`



- Binary semaphore = lock

```
sem_t m;  
sem_init(&m, 0, 1);  
...  
sem_wait(&m);  
// critical section here  
sem_post(&m);
```

```
int sem_wait(sem_t *s) {  
    decrement value  
    if negative, wait  
}  
  
int sem_post(sem_t *s) {  
    increment value  
    wake a waiting thread  
}
```


# Semaphores for ordering

- Example: Parent thread waits for the child thread to finish

```
sem_t s;

void *child(void *arg) {
    sem_post(&s);
    return NULL;
}
```

```
int main() {
    sem_init(&s, 0, 0);
    pthread_t c;
    pthread_create(&c, NULL,
                  child, NULL);
    sem_wait(&s);
    return 0;
}
```



```
int sem_wait(sem_t *s) {
    decrement value
    if negative, wait
}

int sem_post(sem_t *s) {
    increment value
    wake a waiting thread
}
```


# Semaphores for ordering

- Example: Parent thread waits for the child thread to finish

```
sem_t s;

void *child(void *arg) {
    sem_post(&s);
    return NULL;
}
```

```
int main() {
    sem_init(&s, 0, 0);
    pthread_t c;
    pthread_create(&c, NULL,
                  child, NULL);
    sem_wait(&s);
    return 0;
}
```




```
int sem_wait(sem_t *s) {
    decrement value
    if negative, wait
}

int sem_post(sem_t *s) {
    increment value
    wake a waiting thread
}
```

# Reminder: producer/consumer

- **Producer** threads generate data items and place them in the buffer
- **Consumers** pick the items from the buffer and process them
- Examples:
  - Multi-threaded web server
  - `grep foo file.txt | wc -l`
- Now with semaphores!



# Producer/consumer: Template

- Need to add waiting/signaling upon put/get

```
void *producer(void *arg) {  
    for (int i = 0; i < loops; i++) {  
        put(i); ← wait on empty  
    } → signal full  
}  
  
void *consumer(void *arg) {  
    while (1) {  
        int tmp = get(); ← wait on full  
        printf("%d\n", tmp); → signal empty  
    }  
}
```

template producer and consumer

```
int buffer[MAX];  
int fill_ptr = 0;  
int use_ptr = 0;  
  
void put(int value) {  
    buffer[fill_ptr] = value;  
    fill_ptr = (fill_ptr + 1) % MAX;  
}  
  
int get() {  
    int tmp = buffer[use_ptr];  
    use_ptr = (use_ptr + 1) % MAX;  
    return tmp;  
}
```

helper functions

# Two semaphores

- Add one semaphore for each of the two conditions:
- Initially, MAX “empty” resources, and 0 “full” resources
- Every producer cycle moves one from “empty” to “full”

```
sem_t empty;
sem_t full;


void *producer(void *arg) {
    for (int i = 0; i < loops; i++) {
        sem_wait(&empty);           empty--
        put(i);
        sem_post(&full);          full++
    }
}
```

```
void *consumer(void *arg) {
    while (1) {
        sem_wait(&full);      full--
        int tmp = get();
        sem_post(&empty);   empty++
        printf("%d\n", tmp);
    }
}

int main() {
    ...
    sem_init(&empty, 0, MAX);
    sem_init(&full, 0, 0);
    ...
}
```

# Two semaphores: Example 1

- One producer, one consumer, MAX = 1



```
void *producer(void *arg) {  
    for (int i = 0; i < loops; i++) {  
        sem_wait(&empty);  
        put(i);  
        sem_post(&full);  
    }  
}
```

```
void *consumer(void *arg) {  
    while (1) {  
        sem_wait(&full);  
        int tmp = get();  
        sem_post(&empty);  
        printf("%d\n", tmp);  
    }  
}
```

# Two semaphores: Example 2

- Two producers, one consumer, MAX = 1

consumer

```
sem_wait full  
└ value: 0 → -1  
└ wait on full
```

producer 1

```
sem_wait empty  
└ value: 1 → 0  
put(0)  
interrupt
```

producer 2

```
sem_wait empty  
└ value: 0 → -1  
└ wait on empty
```

producer 1

```
sem_post full  
└ value: -1 → 0  
wake full
```

consumer

```
exit sem_wait  
get()  
sem_post empty  
└ value: -1 → 0  
└ wake empty
```

- Producers may run in parallel, but only one gets through `sem_wait` since the value is one

```
void *producer(void *arg) {  
    for (int i = 0; i < loops; i++) {  
        sem_wait(&empty);  
        put(i);  
        sem_post(&full);  
    }  
}
```

```
void *consumer(void *arg) {  
    while (1) {  
        sem_wait(&full);  
        int tmp = get();  
        sem_post(&empty);  
        printf("%d\n", tmp);  
    }  
}
```

# Two semaphores: Example 3

- Two producers, one consumer, MAX = 2

consumer

```
sem_wait full  
└value: 0 → -1  
└wait on full
```

producer 1

```
sem_wait empty  
└value: 2 → 1  
put(0)  
interrupt
```

producer 2

```
sem_wait empty  
└value: 1 → 0  
put(0)  
...
```

possible interleaving

```
p1: buffer[0] = 0;  
p2: buffer[0] = 0;  
p1: fill_ptr = 1;  
p2: fill_ptr = 1;
```

one data item lost

- The producers both enter put without a lock!


```
void *producer(void *arg) {  
    for (int i = 0; i < loops; i++) {  
        sem_wait(&empty);  
        put(i);  
        sem_post(&full);  
    }  
}
```

```
void *consumer(void *arg) {  
    while (1) {  
        sem_wait(&full);  
        int tmp = get();  
        sem_post(&empty);  
        printf("%d\n", tmp);  
    }  
}
```

# Two semaphores and a mutex

- Let us add a lock—another binary semaphore—to guard the critical section
- Only one thread at a time can access the buffer
- But now we can have a deadlock

```
void *producer(void *arg) {
    for (int i = 0; i < loops; i++) {
        sem_wait(&mutex);
        sem_wait(&empty);
        put(i);
        sem_post(&full);
        sem_post(&mutex);
    }
}
```



```
void *consumer(void *arg) {
    while (1) {
        sem_wait(&mutex);
        sem_wait(&full);
        int tmp = get();
        sem_post(&empty);
        sem_post(&mutex);
        printf("%d\n", tmp);
    }
}
```

# Full solution



- Put lock on the critical section only
- Now, a thread cannot go into sleep while holding mutex
- Finally, a working solution for producer/consumer with semaphores only

```
void *producer(void *arg) {
    for (int i = 0; i < loops; i++) {
        sem_wait(&empty);
        [ sem_wait(&mutex);
          put(i);
          sem_post(&mutex);
          sem_post(&full);
        }
}
```

```
void *consumer(void *arg) {
    while (1) {
        sem_wait(&full);
        [ sem_wait(&mutex);
          int tmp = get();
          sem_post(&mutex);
          sem_post(&empty);
          printf("%d\n", tmp);
        }
}
```

# Reader-writer lock

- Assume a data structure that many threads read from, but only a few write to
- Reader threads can go in simultaneously
- Writer thread must block all else



# Reader-writer lock: Implementation

- Keep two locks and a counter

```
void init() {  
    readers = 0;  
    sem_init(&lock, 0, 1);  
    sem_init(&writelock, 0, 1);  
}
```

writer

```
void acquire_writelock() {  
    sem_wait(&writelock);  
}  
  
void release_writelock() {  
    sem_post(&writelock);  
}
```


reader

```
void acquire_readlock() {  
    sem_wait(&lock);  
    readers++;  
    if (readers == 1)  
        sem_wait(&writelock);  
    sem_post(&lock);  
}
```

```
void release_readlock() {  
    sem_wait(&lock);  
    readers--;  
    if (readers == 0)  
        sem_post(&writelock);  
    sem_post(&lock);  
}
```

# Dining philosophers

- Five philosophers seat around the table
- Each has a plate, and there is a fork between each pair of plates
- Philosophers alternate between thinking and eating
- A philosopher can only eat if both adjacent forks are free
- **Task:** Design an algorithm so that no philosopher starves




# Dining philosophers: Attempt

- Eat whenever both forks are available

```
while (1) {  
    think();  
    get_forks(p);  
    eat();  
    put_forks(p);  
}
```

```
void get_forks(int p) {  
    sem_wait(&forks[left(p)]);  
    sem_wait(&forks[right(p)]);  
}
```

```
void put_forks(int p) {  
    sem_post(&forks[left(p)]);  
    sem_post(&forks[right(p)]);  
}
```




cc wiki

deadlock—each has one fork, no one lets go

# Dining philosophers: Solution

- We need to break the symmetry

```
void get_forks(int p) {  
    if (p == 4) {  
        sem_wait(&forks[right(p)]);  
        sem_wait(&forks[left(p)]);  
    } else {  
        sem_wait(&forks[left(p)]);  
        sem_wait(&forks[right(p)]);  
    }  
  
}  
  
void put_forks(int p) {  
    sem_post(&forks[left(p)]);  
    sem_post(&forks[right(p)]);  
}
```



cc wiki

# Throttling

- We might want to limit access even without a race condition
- Example: memory-intensive section
  - 100 threads allocate 1GB each—out of memory, everything goes to swap
  - Only 5 threads allowed at a time—smooth performance
- Solvable by semaphores:

```
sem_t s;
sem_init(&s, 0, 5);
...
sem_wait(&m);
// at most 5 threads at a time
sem_post(&m);
```

# Implementing semaphores

- Store value, a lock and a conditional variable:

```
int value;  
pthread_cond_t cond;  
pthread_mutex_t lock;  
  
init(int v) {  
    value = v;  
    pthread_cond_init(&cond);  
    pthread_mutex_init(&lock);  
}
```

```
void wait() {  
    pthread_mutex_lock(&lock);  
    while (value <= 0)  
        pthread_cond_wait(&cond, &lock);  
    s->value--;  
    pthread_mutex_unlock(&lock);  
}  
  
void post() {  
    pthread_mutex_lock(&lock);  
    value++;  
    pthread_cond_signal(&cond);  
    pthread_mutex_unlock(&lock);  
}
```

# Summary

- Often, semaphore is the only primitive you need to use
  - Implements both “locking” and “ordering” mechanics
  - But, it is challenging to derive conditional variables in full generality
- Locks and conditional variables enable semaphore’s implementation
- Sometimes more general synchronization primitives are more costly
  - A general semaphore vs a lock
  - A read-write lock vs a lock
- See Chapter 31 for more details, and Chapters 32–33 for extra background
- Next two weeks: Storage