

INF113: Long-Term Storage

Kirill Simonov
29.10.2025




Persistence

- We covered many aspects of the OS:
 - How to let the processes “time-share” the CPU
 - How to share memory access
 - How to run multiple parallel threads over the same address space
- All of this vanishes once the power goes out
- **Persistence:** How to manage long-term storage?
- Challenges:
 - Slow devices
 - Reliability
 - Abstraction


System hierarchy

- Slower rates—farther from CPU
- The slower the device, the more “custom” it can be
- Modern systems would have specialized chips and faster point-to-point interconnects
 - e.g., to the GPU



A canonical device

- How to work with wildly different devices?
- Every device will have registers for control




- The canonical interaction protocol:


```
while (STATUS == BUSY)
; // wait until device is not busy
write data to DATA register
write command to COMMAND register
(starts the device and executes the command)
while (STATUS == BUSY)
; // wait until device is done with your request
```

Optimization 1: Interrupts

- Polling wastes CPU cycles




- Interrupts instead of active polling:
 1. Request for a status of a device
 2. Process is put to sleep
 3. Interrupt raised when the device is done
 4. Operation is completed




Optimization 2: DMA

- Write to disk:
 1. Copy data to disk interface
 2. Disk performs the write




- Direct Memory Access (DMA)
 1. DMA engine receives instructions
 2. Performs copy from memory in parallel
 3. Raises interrupt when it's done



Device drivers

- Certain privileged instructions are used to interact with the device registers
 - e.g., in and out on x86
 - need to specify **port**, i.e., which device, and the register
- **Driver** provides an abstract interface to the device to the rest of the OS
 - Hides the specific register locations, signal values, interaction protocol



Hard disk drive (HDD)


- Has a spinning magnetic **platter** and a **head** that shifts to the desired sector
- Cons:
 - Slower than SSDs
 - Bulky and noisy
 - Sensitive to mechanical action
- Pros:
 - Cheaper per 1TB
 - Reliable long-term
 - Faster with continuous writing
- Still most of the datacenter storage



cc wiki

Understanding HDDs

- The platter consists of 512-byte **sectors/blocks**, indexed from 0 to $n - 1$
- The head reads/writes one sector at a time, atomically
- The head needs to mechanically reach the sector
- **Rotational delay:** waiting for a sector on the same track
 - For 10 000 RPM, 6ms for the whole rotation
- **Seek:** moving the head to the right track
 - Settling time 0.5-2ms
- Once the head is in place, transfer is fast—e.g., $30\mu\text{s}$
- Much faster to access sectors sequentially
 - Random access: 5–8 ms, next sector: $30\ \mu\text{s}$



Solid-state drive (SSD)

- Stores individual bit values in transistors
- Pros:
 - Fast
 - No mechanical parts
- Cons:
 - More expensive
 - Wears out on rewrites
 - Writing is complicated





cc wiki

Device	Random Reads (MB/s)	Random Writes (MB/s)	Sequential Reads (MB/s)	Sequential Writes (MB/s)
Samsung 840 Pro SSD	103	287	421	384
Seagate 600 SSD	84	252	424	374
Intel SSD 335 SSD	39	222	344	354
Seagate Savvio 15K.3 HDD	2	2	223	223

Understanding SSDs


- Bits are grouped in 4KB **pages**, pages in 256KB **blocks** and blocks into **banks**



- **Read:** Get contents of a page by its index, $\sim 10\mu\text{s}$
- Cannot write to a page at will, instead:
 - **Erase a block:** sets the entire block to 1, $\sim 3\text{ms}$
 - **Program a page:** sets an erased page to the desired contents, $\sim 100\mu\text{s}$
- **Wear out:** each block can only do a finite number of erase cycles
- **Solution: logging**—append changes, instead of rewriting directly
- **Solution: wear leveling**—spread writes across all pages

Files and directories

- **File** is an array of bytes
 - OS keeps a low-level identifier, **inode number**
- **Directory** stores a list of pairs:
(user-readable name, inode number)
 - Is really a special kind of file
 - Also has an inode number
- Files and directories form a tree-like hierarchy
 - Starts from the root directory /
 - Each file can be referenced by the **absolute pathname**
`/bar/foo/bar.txt`
- In UNIX systems, everything is accessible within a single tree
 - Also devices, pipes, processes



File system interface

- Creating files: possible by open syscall

```
int fd = open("foo", O_CREAT|O_WRONLY|O_TRUNC, S_IRUSR|S_IWUSR);
```

- Return value: **file descriptor**, a local index of an opened file
 - OS stores the list of files opened by a process

```
struct proc {  
    ...  
    struct file *ofile[NOFILE]; // Open files  
    ...  
};
```

- Reading files

```
$ echo hello > foo  
$ cat foo  
hello
```

0: stdin
1: stdout
2: stderr
3: first file
...
...

```
$ strace cat foo  
...  
open("foo", O_RDONLY|O_LARGEFILE) = 3  
read(3, "hello\n", 4096) = 6  
write(1, "hello\n", 6) = 6  
hello  
read(3, "", 4096) = 0  
close(3) = 0  
...
```

Non-sequential reading and writing

- A process need not to go through the whole file byte-by-byte
- lseek syscall shifts the “next character” by the desired offset

```
off_t lseek(int fd, off_t offset, int whence);
```

- For each file opened by the process, OS tracks the current offset
 - updates implicitly by read/write
 - updates explicitly by lseek

```
struct file {  
    int ref;  
    char readable;  
    char writable;  
    struct inode *ip;  
    uint off;  
};
```

Renaming files

```
$ mv foo bar
```

- rename syscall “atomically” renames a file

```
int rename(char *old, char *new);
```

- Can be used to safely edit files:

```
int fd = open("foo.txt.tmp", O_WRONLY|O_CREAT|O_TRUNC,  
              S_IRUSR|S_IWUSR);  
write(fd, buffer, size); // write out new version  
fsync(fd);  
close(fd);  
rename("foo.txt.tmp", "foo.txt");
```

- fsync syscall forces the write to disk as soon as possible, instead of buffering

File metadata

- stat syscall gives information about the file

```
$ stat foo
  File: foo
  Size: 6          Blocks: 8          IO Block: 4096   regular file
Device: 259,4 Inode: 23598029      Links: 1
Access: (0664/-rw-rw-r--)  Uid: ( 1000/ seemann)    Gid: ( 1000/ seemann)
Access: 2025-10-27 19:58:53.527389303 +0100
Modify: 2025-10-27 19:58:50.937389495 +0100
Change: 2025-10-27 19:58:50.937389495 +0100
 Birth: 2025-10-27 19:58:50.937389495 +0100
```

- OS stores this information about files in a structure called **inode**

Directories

- Directory files are stored in a special format so should not be read directly
- Create empty directory: `mkdir`

```
$ strace mkdir foo  
...  
mkdir("foo", 0777) = 0  
...
```

- Cycle over directory entries: `opendir`, `readdir`, `closedir`

```
struct dirent {  
    char d_name[256]; // filename  
    ino_t d_ino; // inode number  
    off_t d_off; // offset to the next dirent  
    unsigned short d_reclen; // length of this record  
    unsigned char d_type; // type of file  
};
```

```
$ ls
```

Links

- Removing a file simply calls `unlink()`

```
$ strace rm foo  
...  
unlink("foo") = 0  
...
```

- `link()` makes a **hard link**—a copy of the directory entry

```
$ ln file file2  
$ ls -i file file2  
67158084 file  
67158084 file2
```

- OS keeps track of the number of hard links to the actual file (i.e., inode)
 - When created, the file is linked to its human-readable name
 - When linked, counter increases
 - When unlinked, counter decreases (also in `rm foo`)
 - When it reaches zero, the whole entry at inode is removed

Soft links

- **Soft link** is a special type of file that contains the pathname of another file
 - Can point to a directory
 - Can point to a different device

```
$ ln -s foo foo2
```

- We can see that soft link is a special file by running ls -l and stat

```
$ ls -l
-rw-rw-r-- 1 ... foo
lrwxrwxrwx 1 ... foo2 -> foo
```

```
$ stat foo2
  File: foo2 -> foo
  Size: 3          Blocks: 0          IO Block: 4096   symbolic link
Device: 259,4  Inode: 23593299      Links: 1
```

Owners and permission bits

- Each file has an owner (user) and a group (collection of users)
- A file has also permission settings for each of the three access types: by own user, by own group, or by anyone else

```
$ ls -l foo  
-rw-rw-r-- 1 user group 0 Oct 28 19:13 foo
```

- Permission could be to read the file (r), to write to the file (w), or execute it (x)
- Permission set has a corresponding bitmask, e.g., rwx is 111 or 7

```
$ chmod 640 foo
```

6—to owner	4—to group	0—to others
110 = read and write	100 = read only	000 = nothing


Making and mounting a file system

- `mkfs` creates an empty file system of given `fs` type on a given device
don't run on your devices that already have a file system!
- To be part of the directory tree, the file system has to be **mounted**
- Mount attaches the directory tree of the file system under the given **mount point**

```
$ mount -t ext3 /dev/sda1 /home/users
$ cd /home/users/a
```

- List all mounted fs by calling `mount` without arguments:

```
$ mount
/dev/sda1 on / type ext3 (rw)
...
```



- With `mount`, all different file systems and devices are in the same directory tree