


INF113: FFS, Crash Consistency

Kirill Simonov
05.11.2025





Reminder

- We described a very simple file system
- VSFS stores:
 - Superblock
 - Allocation bitmaps
 - Inode blocks
 - Data region
- Any other file system can be implemented by replacing the internal representation and open/read/write/... syscalls




Locality

- Storage benefits from “locality”—storing related content close
 - Example: hard disk drives
 - If two consecutive reads are far apart, HDD has a long seek to perform
 - VSFS is not great for locality
 - Reading from a file: first inode contents, then data itself
 - But inode and its data blocks may be very far apart
 - Even more jumps from determining the file’s inode number
`/foo/bar: inode / → data / → inode foo → data foo → inode bar` —




Fragmentation

- Data blocks of a file will not always be consecutive
create A → create B → remove A → create C — C is split!
- “Allocate first available block” strategy will produce fragmented files
 - The more fragmented a file is, the slower it is to read (from HDD)
- **Defragmentation:** rearranging data blocks on the whole disk to make files continuous



Cylinder groups

- HDD motivation: blocks in the same “cylinder” can be read in a sequence quickly
- Cylinder groups: batch a few close cylinders together
- Each group has the same internal structure as the whole VSFS:
 - Superblock
 - Allocation blocks
 - Inodes
 - Data blocks




4 cylinders

2 cylinder
groups


Block groups

- Same idea without relying on physical structure of the disk
- Block group: a sequence of consecutive blocks on the disk
- Each group has its own
 - Superblock
 - Allocation blocks
 - Inodes
 - Data blocks
- Creating a file may happen within a single block group
 - inode bitmap → new inode → data bitmap → new data block




Fast File System (FFS)

- Designed for 4.2BSD (1983)
- Introduced block groups and other features
 - Long file names, symbolic links, atomic rename, ...
- Policies for choosing the group:
 - Single file: Inode and data blocks in the same group
 - Directory and its files in the same group
 - Different directories far away
 - Large file: first few blocks in a single group
- Block groups are still present in modern file systems



Crash consistency


- A single write request changes multiple on-disk structures
 - Write a new data block
 - Add its number to inode
 - Mark it allocated in the data bitmap
 - Potentially more: directory data, new indirect node, ...
- In case of a crash, only some of these changes will be applied
- This might not just lose user data, but also put the file system into an **inconsistent** state



Crash behaviors

Example: Append a data block to a file

1. Write new data block D2
2. Add reference to D2 to the inode: $\text{lv1} \rightarrow \text{lv2}$
3. Mark D2 allocated: $\text{db v1} \rightarrow \text{db v2}$



Crash behaviors

Example: Append a data block to a file


1. Write new data block D2
 2. Add reference to D2 to the inode: $\text{lv1} \rightarrow \text{lv2}$
 3. Mark D2 allocated: $\text{db v1} \rightarrow \text{db v2}$
- Only step 1: D2 is written, but nothing points to it—same as no write



Crash behaviors

Example: Append a data block to a file


1. Write new data block D2
 2. Add reference to D2 to the inode: $\text{lv1} \rightarrow \text{lv2}$
 3. Mark D2 allocated: $\text{db v1} \rightarrow \text{db v2}$
- Only step 1: D2 is written, but nothing points to it—same as no write
 - Only step 2: Inode points to the block, but the data was not written!
 - Inconsistency: Inode thinks D2 is part of the file, data bitmap thinks it's free



Crash behaviors

Example: Append a data block to a file


1. Write new data block D2
 2. Add reference to D2 to the inode: $\text{lv1} \rightarrow \text{lv2}$
 3. Mark D2 allocated: $\text{db v1} \rightarrow \text{db v2}$
- Only step 1: D2 is written, but nothing points to it—same as no write
 - Only step 2: Inode points to the block, but the data was not written!
 - Inconsistency: Inode thinks D2 is part of the file, data bitmap thinks it's free
 - Only step 3: Data bitmap marks D2 as allocated, but no file uses the block
 - Inconsistency: Data bitmap marks more blocks than in use \rightarrow space leak



Crash behaviors—part 2


Example: Append a data block to a file

1. Write new data block D2
 2. Add reference to D2 to the inode: $\text{lv1} \rightarrow \text{lv2}$
 3. Mark D2 allocated: $\text{db v1} \rightarrow \text{db v2}$
- Only 2 and 3: Inode and data bitmap agree on the block—but no data there!
 - Only 1 and 2: Inode points to the block and the data is correct.
 - Inconsistency: Inode thinks D2 is part of the file, data bitmap thinks it's free
 - Only 1 and 3: Data bitmap marks D2 as allocated, but no file uses the block
 - Inconsistency: Data bitmap marks more blocks than in use \rightarrow space leak



File system checker


- Assume a crash that causes inconsistency happens
- Fix inconsistencies after reboot, before mounting the device again
 - UNIX utility `fsck`
- No additional information—scan the whole filesystem
- Detects and fixes inconsistencies in file system structures
 - Cannot detect data blocks that weren't written



Stages of fsck

- **Superblock:**


- Sanity checks: file system size, etc.
- See if some of the superblock copies are corrupted
- Replace by another copy



Stages of fsck

- **Superblock:**
 - Sanity checks: file system size, etc.
 - See if some of the superblock copies are corrupted
 - Replace by another copy
- **Allocated data blocks:**
 - List all data blocks in use, as marked in inodes
 - Trust inodes if inconsistent with data bitmap


file
8192 bytes
2 blocks: 9, 12
inode



Stages of fsck

- **Superblock:**
 - Sanity checks: file system size, etc.
 - See if some of the superblock copies are corrupted
 - Replace by another copy
- **Allocated data blocks:**
 - List all data blocks in use, as marked in inodes
 - Trust inodes if inconsistent with data bitmap
- **Allocated inodes:**
 - Scan through inode locations and directory contents


file
8192 bytes
2 blocks: 9, 12
inode



Stages of fsck—part 2

- **Inode state:**
 - If inode contents look corrupted, it is cleared

file
8192 bytes
2 blocks: 9, 12
inode



Stages of fsck—part 2

- **Inode state:**
 - If inode contents look corrupted, it is cleared
- **Inode links:**
 - Scan directories and count links to the particular inode
 - Update link count
 - No directory point to the inode: move to lost+found

file
8192 bytes
2 blocks: 9, 12
inode



Stages of fsck—part 2

- **Inode state:**
 - If inode contents look corrupted, it is cleared
- **Inode links:**
 - Scan directories and count links to the particular inode
 - Update link count
 - No directory point to the inode: move to lost+found
- **Duplicates:**
 - If two inodes point to the same data block, copy the block

file
8192 bytes
2 blocks: 9, 12
inode



Stages of fsck—part 3

- **Bad blocks:**
 - If inode references an out-of-range block, remove the reference



Stages of fsck—part 3

- **Bad blocks:**
 - If inode references an out-of-range block, remove the reference
- **Directories:**
 - Aren't arbitrary user data, so can be checked too
 - . and .. are the first two entries
 - Each referenced inode is allocated
 - Each directory is linked once

inum	name
2	.
2	..
4	bar
5	bin
...	

directory



Summary: fsck

- Can restore the file system to a consistent state with minimal modifications
- Does not catch errors that leave no inconsistencies
- **Main issue:** Scans the whole file system, so is extremely slow
- **Next time:** Journaling
 - Leave a note before each write
 - Know exactly what to look for after the crash
- Chapter 41 for FFS and Chapter 42 for crash consistency and journaling