

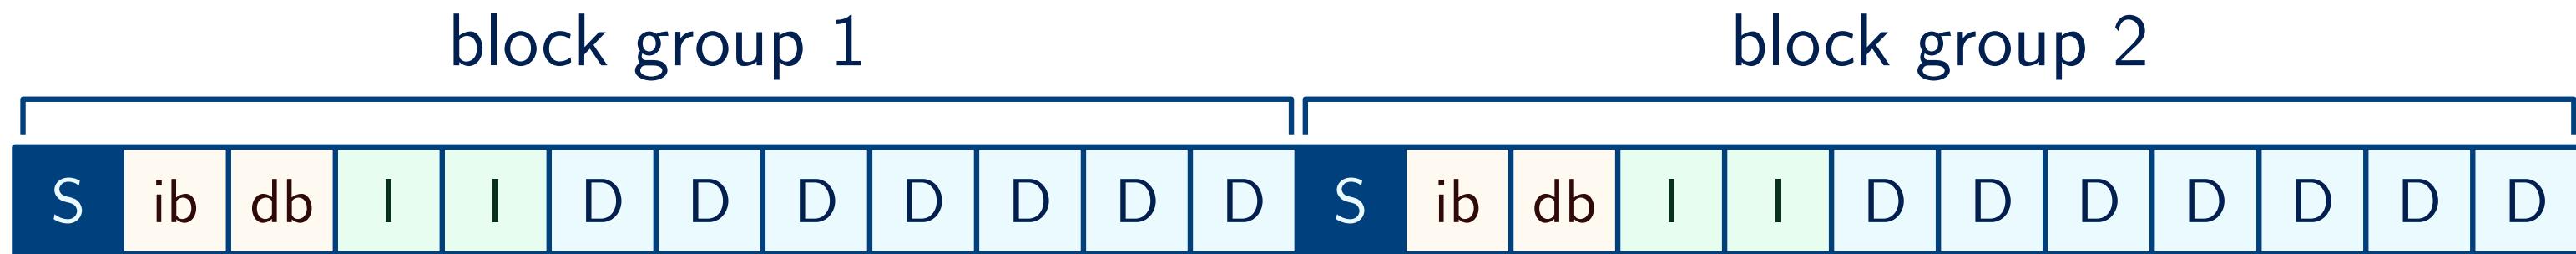
INF113: Journaling

Kirill Simonov
07.11.2025



Reminder

- We enhanced the very simple file system with block groups



- We discussed crash scenarios where the file system becomes **inconsistent**



- **Solution 1:** fsck, scan the whole file system after crash
 - Slow, does not detect errors that leave fs consistent
- **Today:** what if we leave a note to help us find which write failed?

Journaling

- Keep a portion of the disk as the **journal**
- Before any write operation, log the write into the journal first
 - Another term: “write-ahead logging”
- Little extra work for each write → Much easier recovery

Linux ext2 (no journaling)

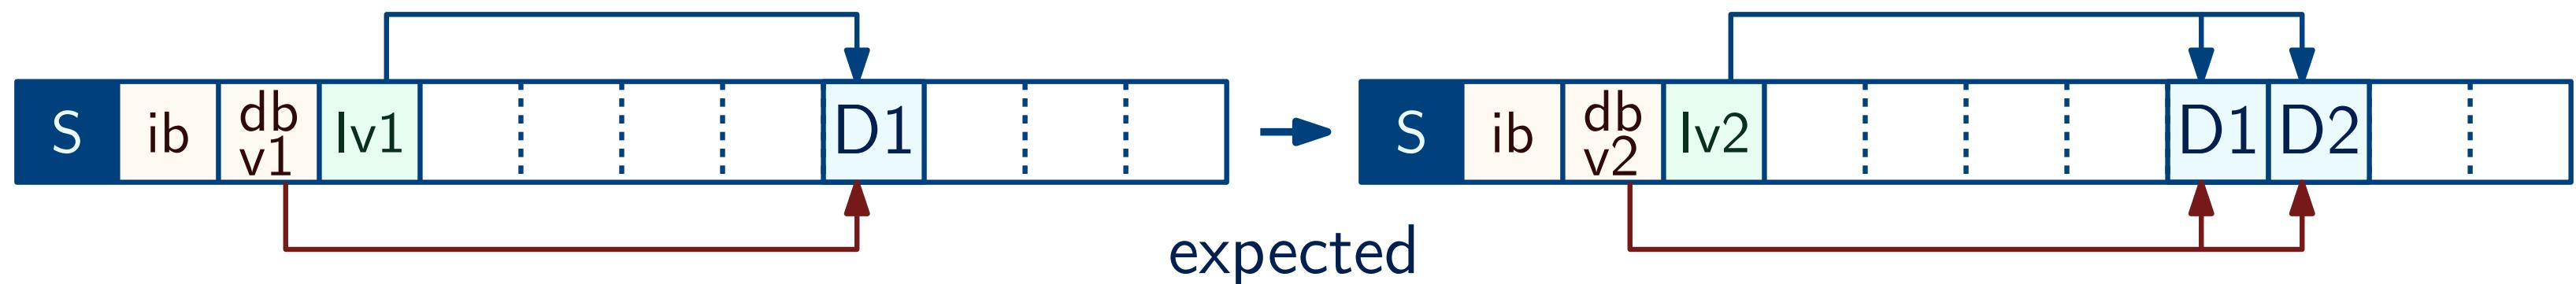


Linux ext3 (with journaling)



Data journaling

Append a data block to a file: write data block (D2), new inode block (lv2), new data bitmap (db v2)



Writes to the journal:

1. Transaction begin block (TxB): addresses and transaction id
2. Copies of blocks to be written (lv2, db v2, D2)
3. Transaction end block (TxE): transaction id

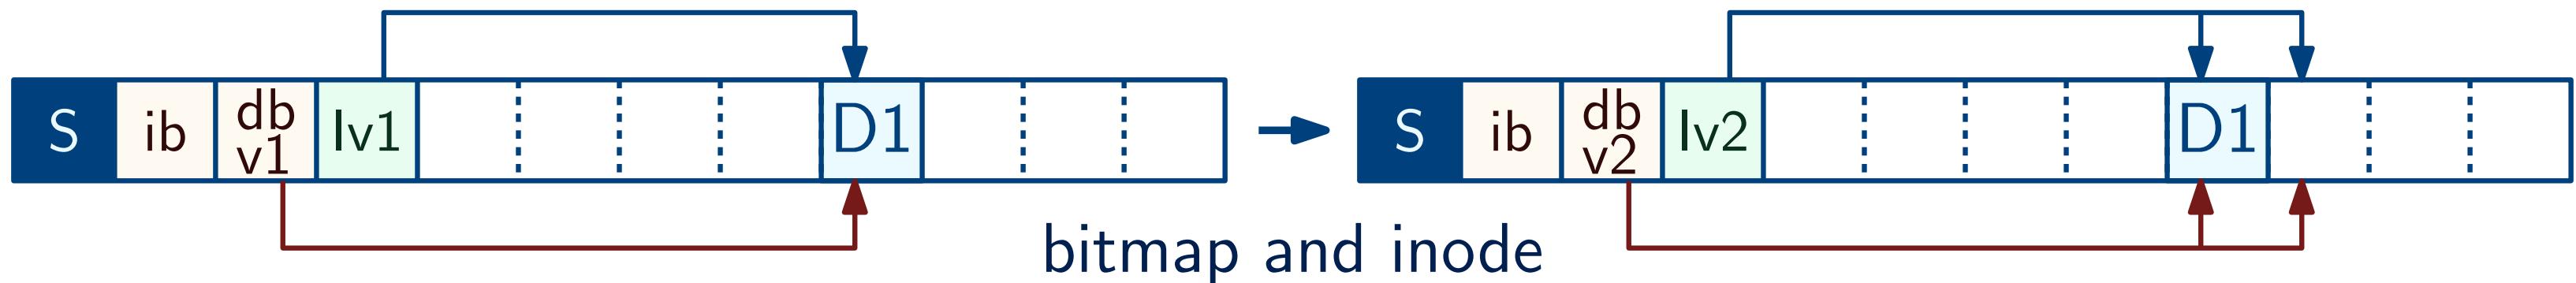


Checkpointing

1. **Journal write:** Write the transaction (begin+blocks+end), wait for it to complete



2. **Checkpoint:** Write the updates to their final locations in the file system



- If a crash happens during checkpointing, repeat from the journal
 - Need to know that a crash happened
 - And which transactions in the journal are not completed

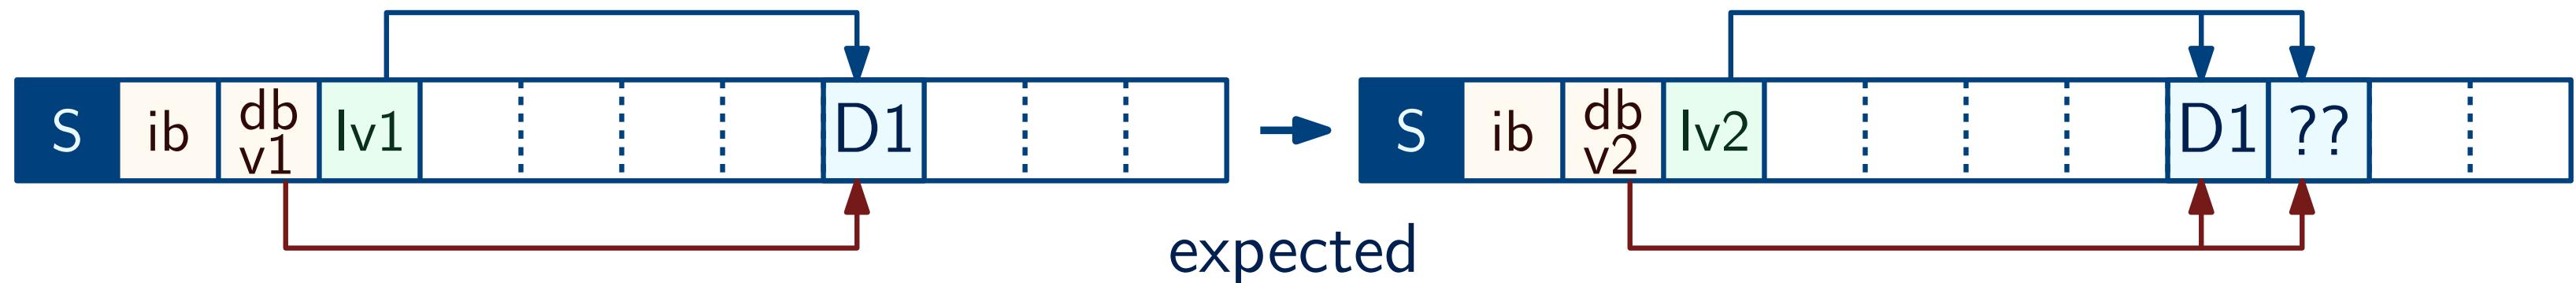
even if
consistent!

Crashes during logging

- Assume journal blocks are written in arbitrary order; crash before D2 is written



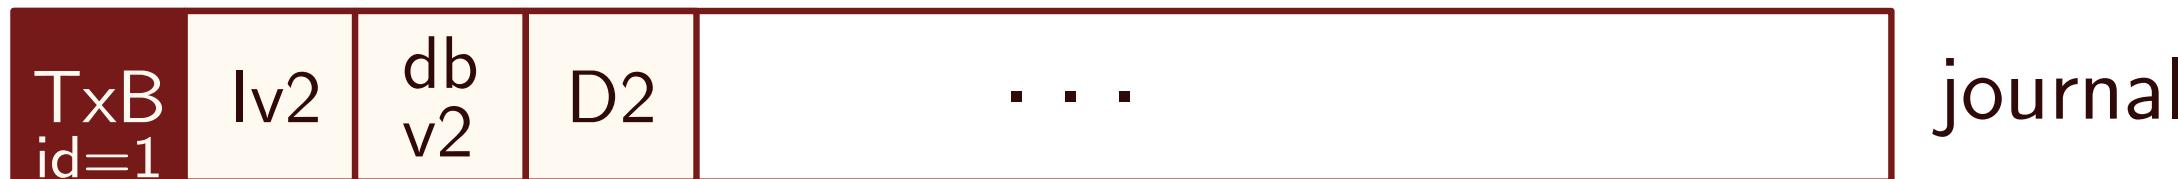
- System attempts to reapply the transaction, which looks valid



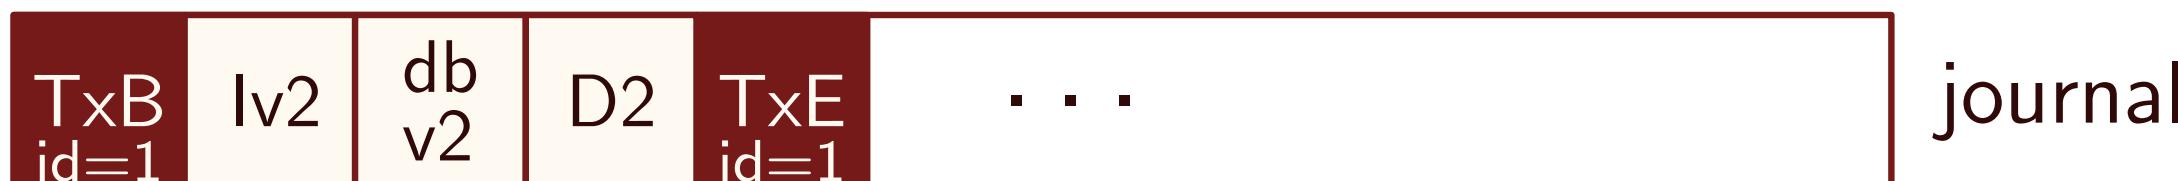
- The file is now corrupted!
- We need to make sure the transaction was fully written to the journal

Writing the transaction

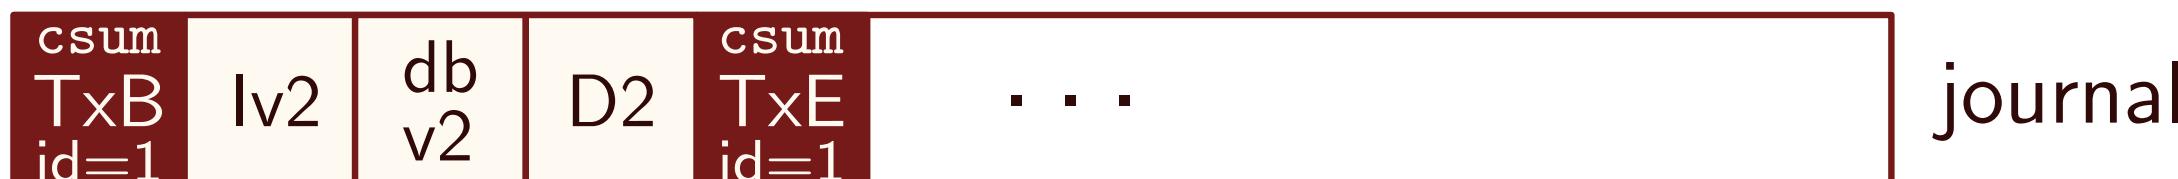
- Writing all blocks one after the other is too slow
- First write the start ($\text{Tx}B + \text{blocks}$) in any order, wait to complete



- Then write the end block ($\text{Tx}E$), wait to complete

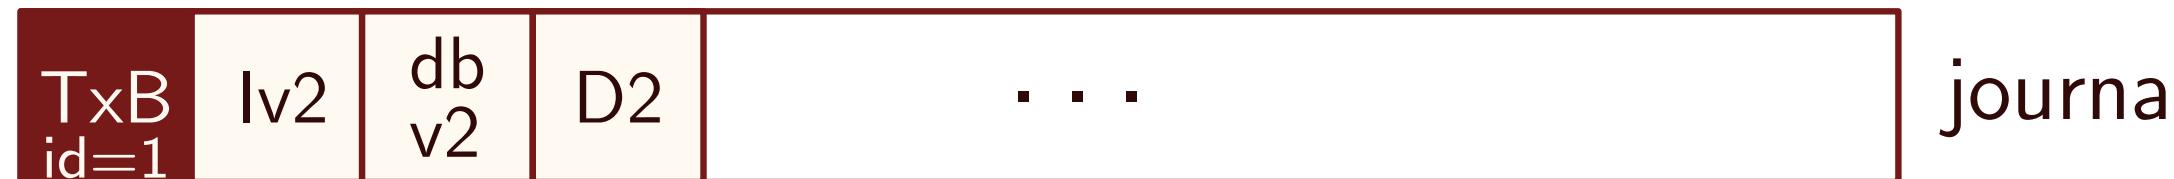


- Alternative solution (ext4): any order, but write a checksum to $\text{Tx}B$ and $\text{Tx}E$

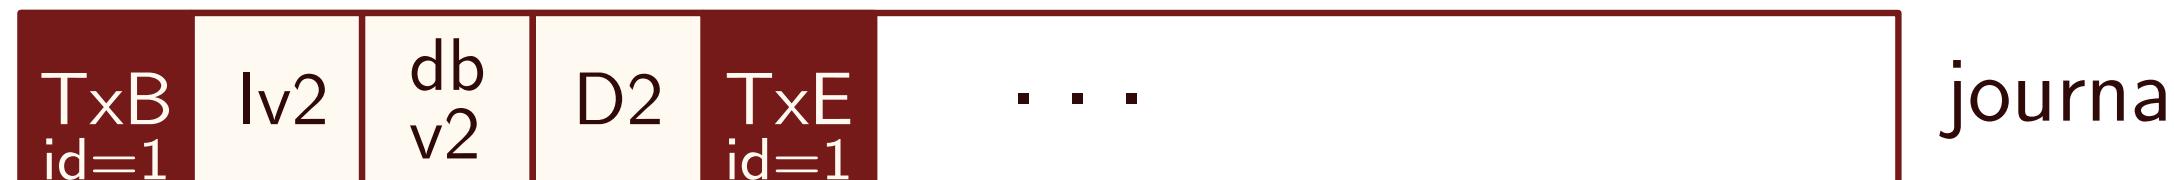


Final protocol

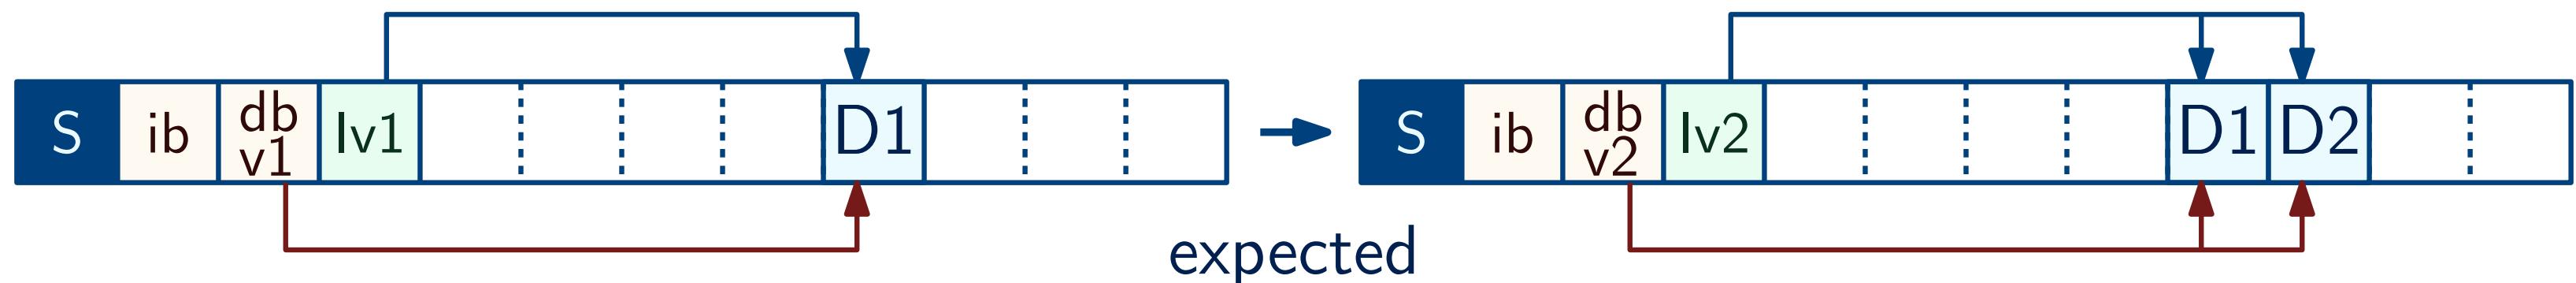
1. **Journal write:** Write the transaction (TxB+blocks), wait to complete



2. **Journal commit:** Write the transaction commit block (TxE), wait to complete

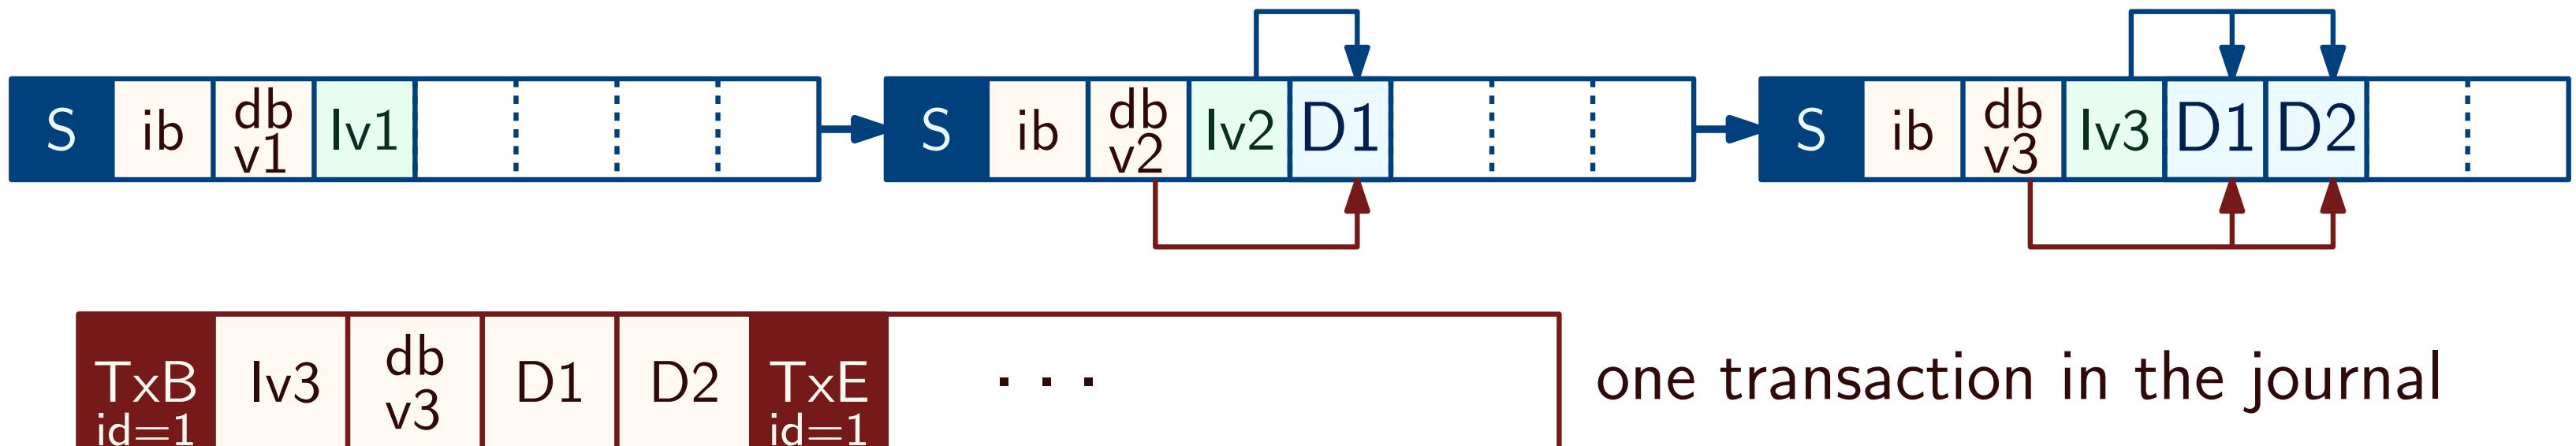


3. **Checkpoint:** Write the updates to their final locations in the file system



Batching updates

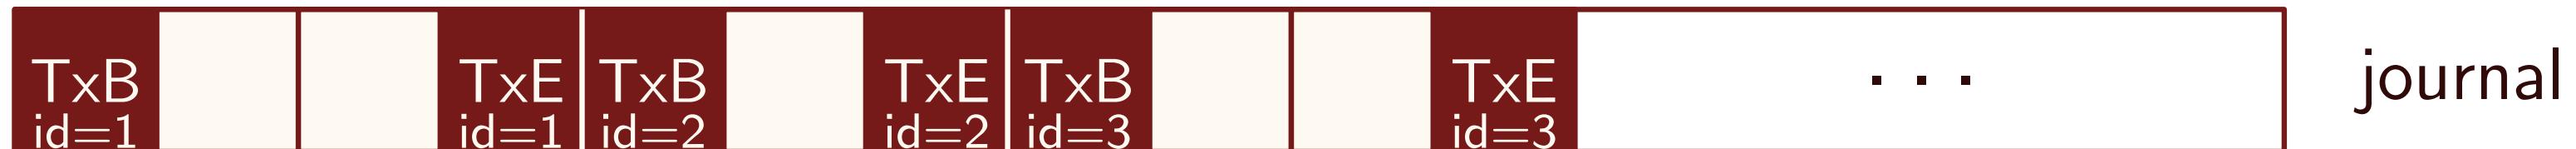
- Every write is now doubled: first in the log, then in the destination
- We can improve performance by **batching** write requests:
 - Gather writes for 5 seconds, merge and send to the disk
 - Journal records a single transaction
- Rewriting same data block multiple times → one write sent to disk, one log entry
- Appending multiple blocks to the same file → one inode rewrite, one bitmap rewrite



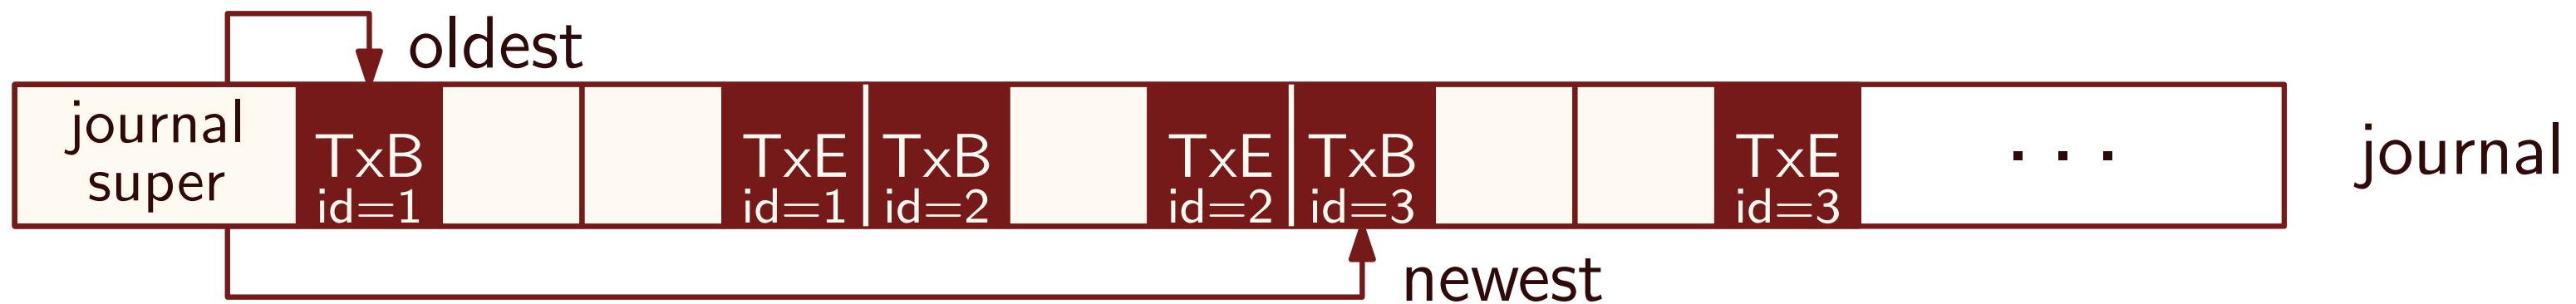
- Saves rewrites in other cases too, e.g., for creating two files within the same directory

Clearing the log

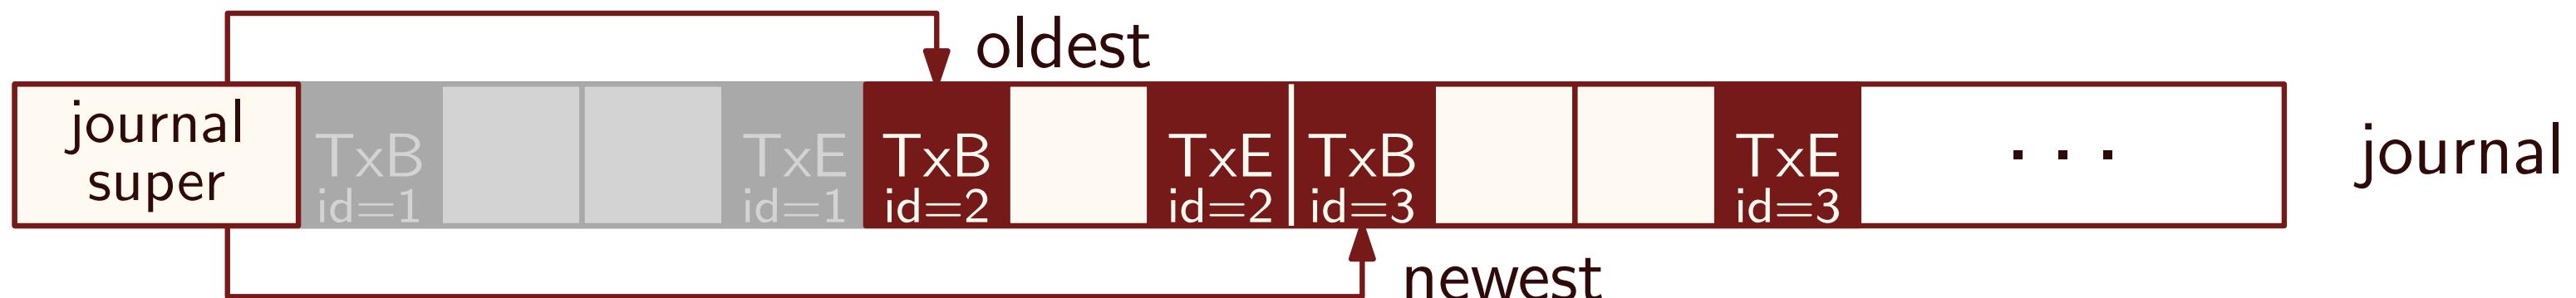
- Eventually the journal runs out of space



- **Circular log:** keep the oldest and the newest unapplied transaction

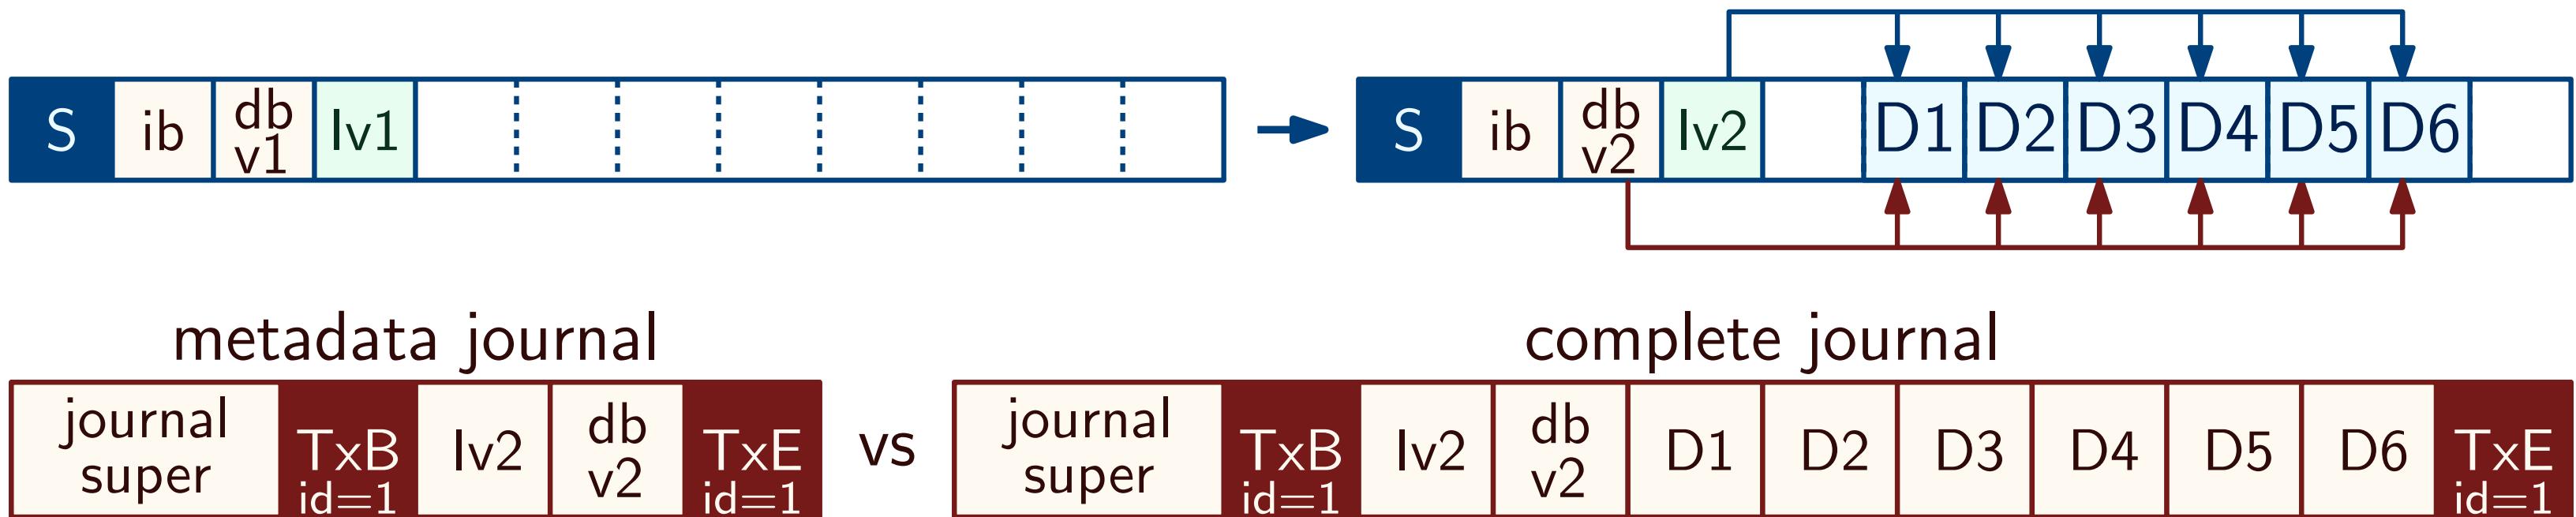


- After the transaction is checkpointed, mark it free—shift the “oldest” pointer



Metadata journaling

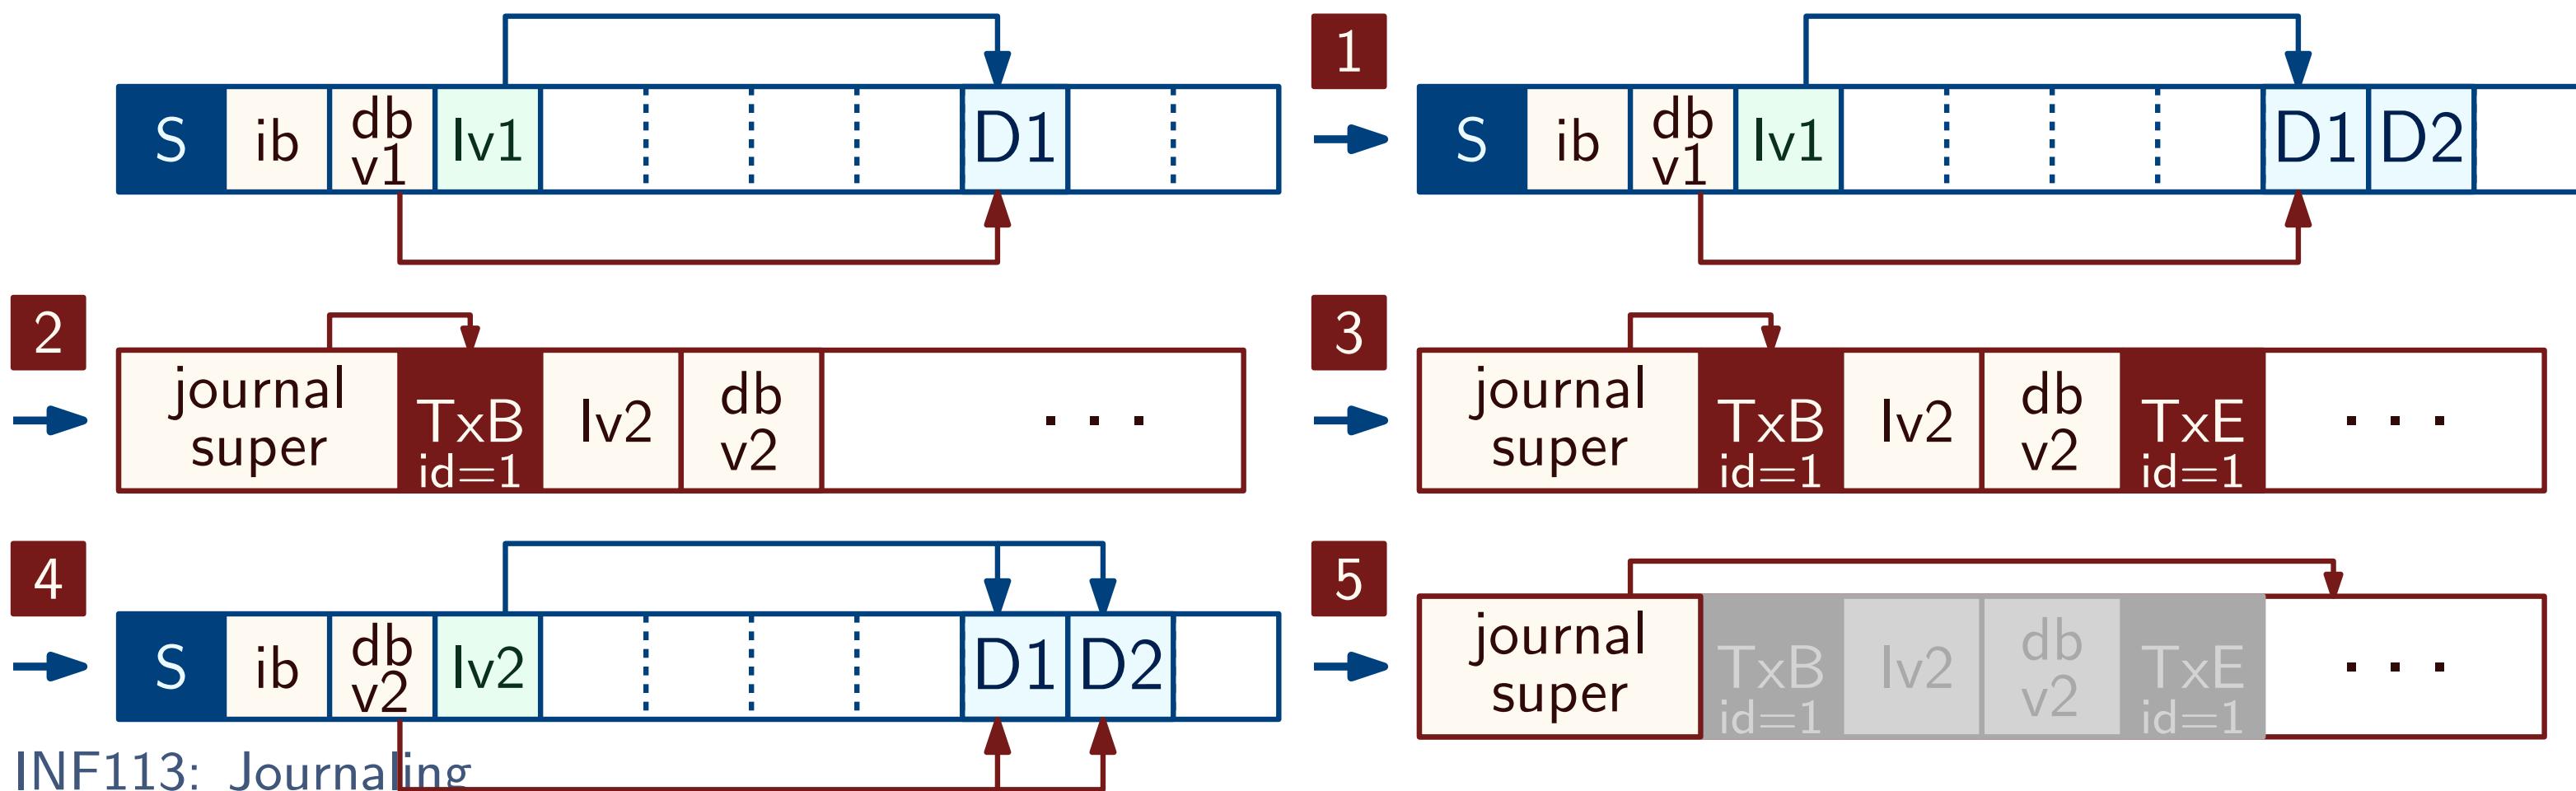
- Every block to be written is also saved in the journal → twice the number of writes
- **Metadata journaling:** Only save file system metadata to the log



- User data is saved directly to the destination
- Most data is user data → much smaller overhead

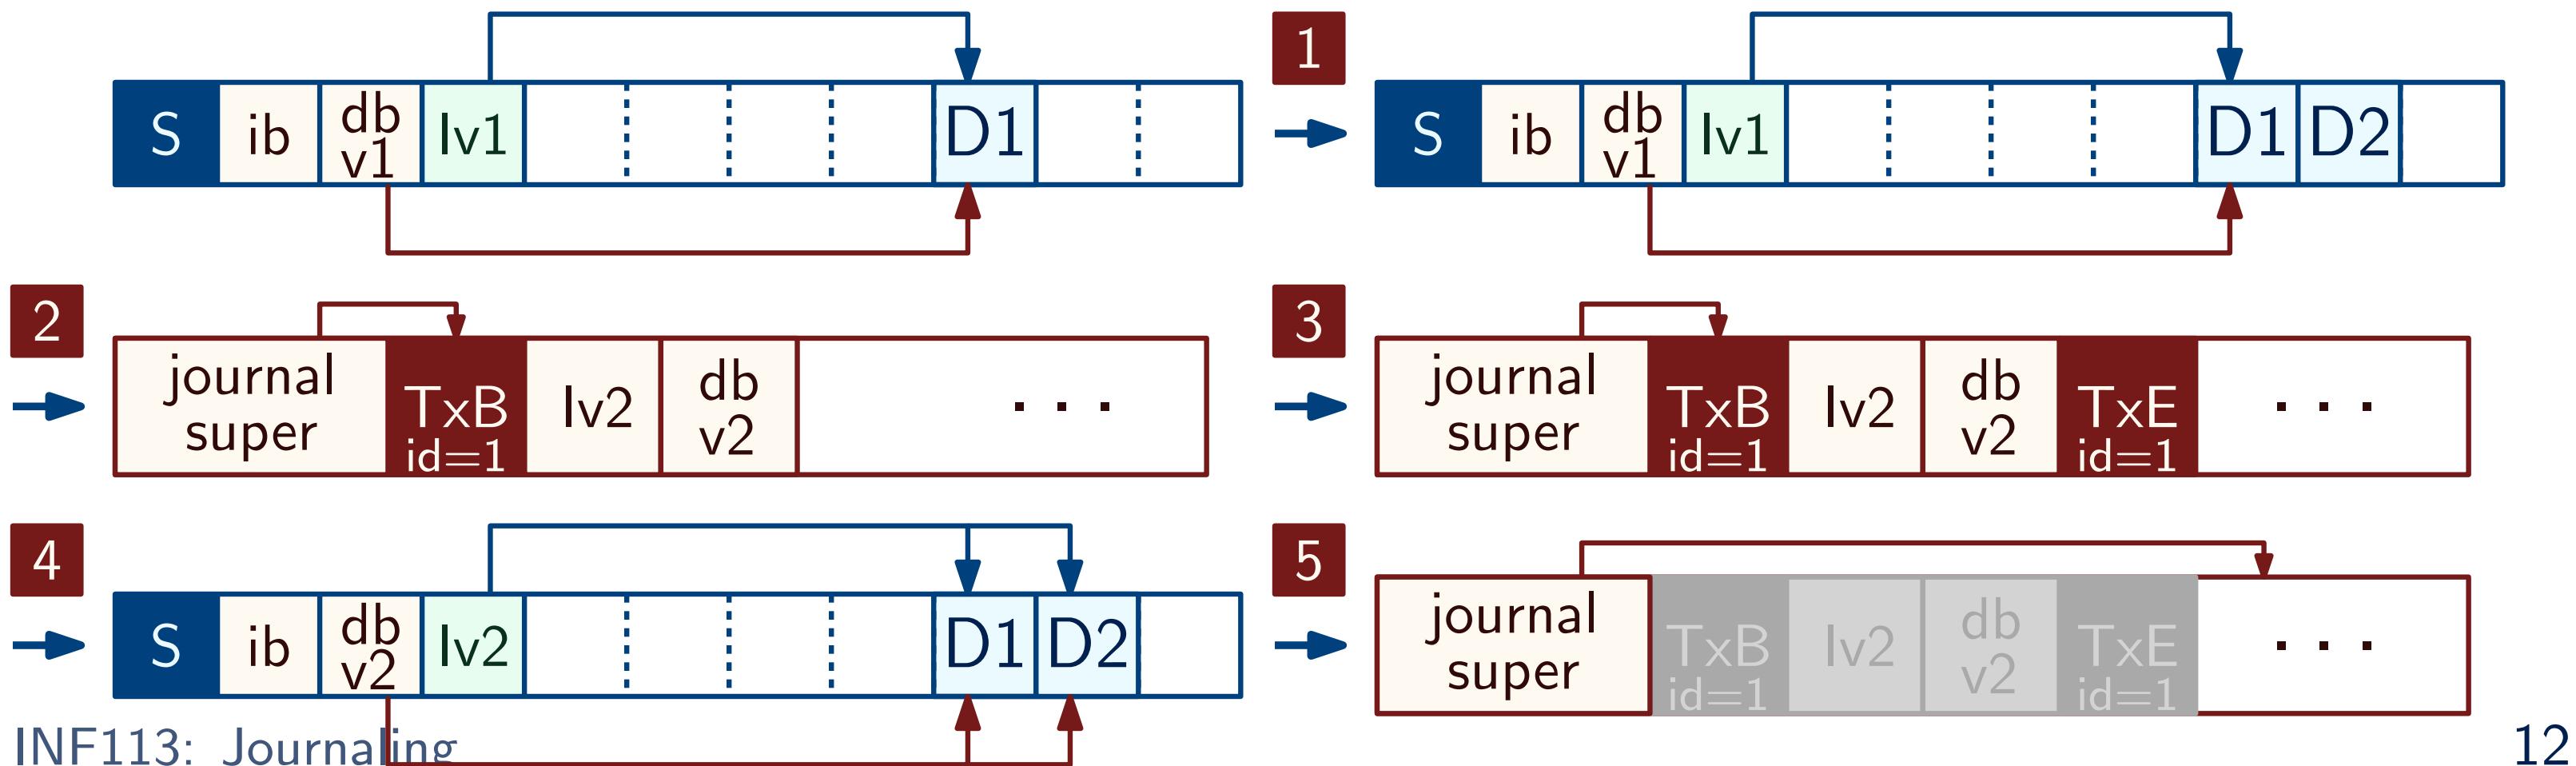
Metadata journaling protocol

1. **Data write:** Write the data to the final location, wait to complete
2. **Journal metadata write:** Write the begin block and metadata to the log, wait
3. **Journal commit:** Write the transaction commit block to the log, wait to complete
4. **Checkpoint metadata:** Write the metadata updates to their final locations
5. **Free:** Later, mark the transaction free in journal superblock



Crash scenarios

1. Crash after data write: no metadata/journal changed
 2. Crash after journal write: incomplete transaction, discarded on restore
 3. Crash after journal commit: steps 4 and 5 reattempted upon restore
 4. Crash after checkpoint: 4 and 5 reattempted upon restore
 5. Crash after free: All changes are saved to disk
-] write lost
] write saved



Block reuse

Metadata journaling may have an issue with reusing data blocks:

1. Add a file to directory foo, data in block 10
2. Remove directory foo
3. Create file bar, its data gets written to block 10
4. System crashes, and the journal is replayed
 - The outdated data block of foo is written to block 10
 - The data of bar is not in the journal, and is lost

Solution: Do not reuse same data blocks until checkpoint is completed

Or **revoke** outdated journal entries (ext3)—special entry in the journal



Summary

- Journaling allows to quickly and reliably repair the file system in case of a crash
- Only adding metadata to the journal greatly reduces the overhead
- (Some) file systems that use journaling:
 - ext3/ext4: user can choose between data/metadata journaling
 - ntfs: some form of metadata journaling
- Alternative solutions:
 - **Soft updates**: writes are ordered so that inconsistency is impossible, or is a leak
 - **Log-structured fs**: journal is the whole file system, data is read from the journal
 - **Copy-on-write**: no in-place writes, always to a new block (e.g., APFS)

Next week & exam

- Today is the final topic of the course
- Wednesday next week—Q&A session for course topics and Assignment 3
- Friday next week—no lecture
- Group sessions next week as usual
- Assignment 3 deadline next Friday (Nov 14)
- Exam preparation:
 - Try past exams—“Exam preparation” module on Mitt
 - Revisit lecture slides/recordings
 - Revisit book chapters—see the list on course homepage